

CS 6476: Computer Vision

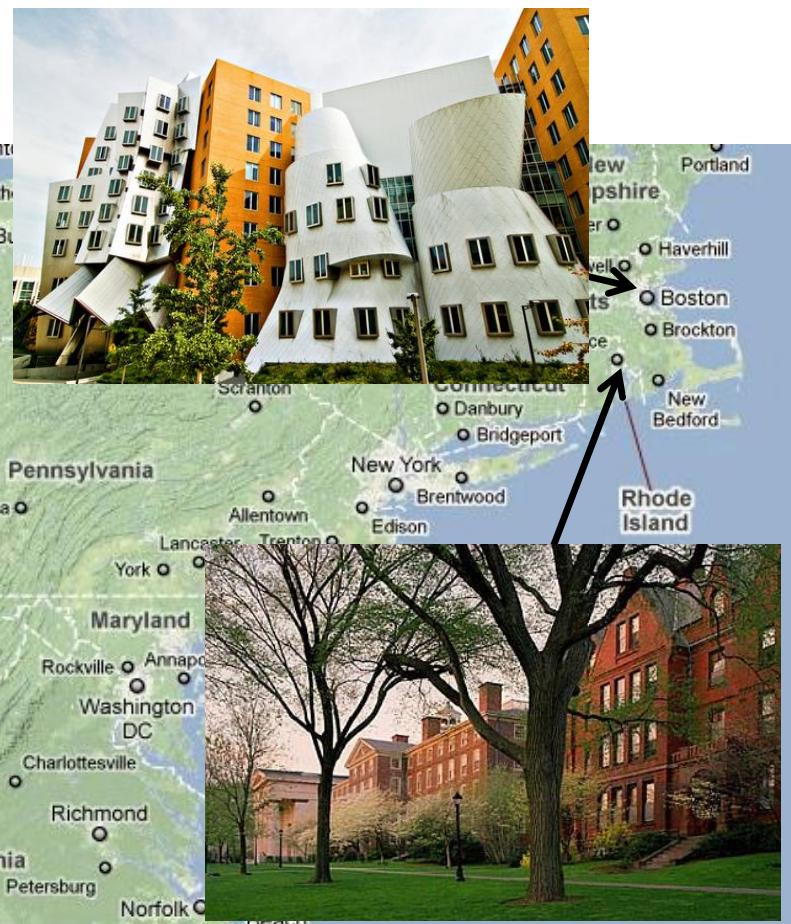
Instructor: James Hays

Tas: Cusuh Ham (head TA), Sean Foley, Jianan Gao,
John Lambert, (more to come)

Today's Class

- Course enrollment
- Who am I?
- What is Computer Vision?
- Specifics of this course
- Geometry of Image Formation
- Questions

A bit about me

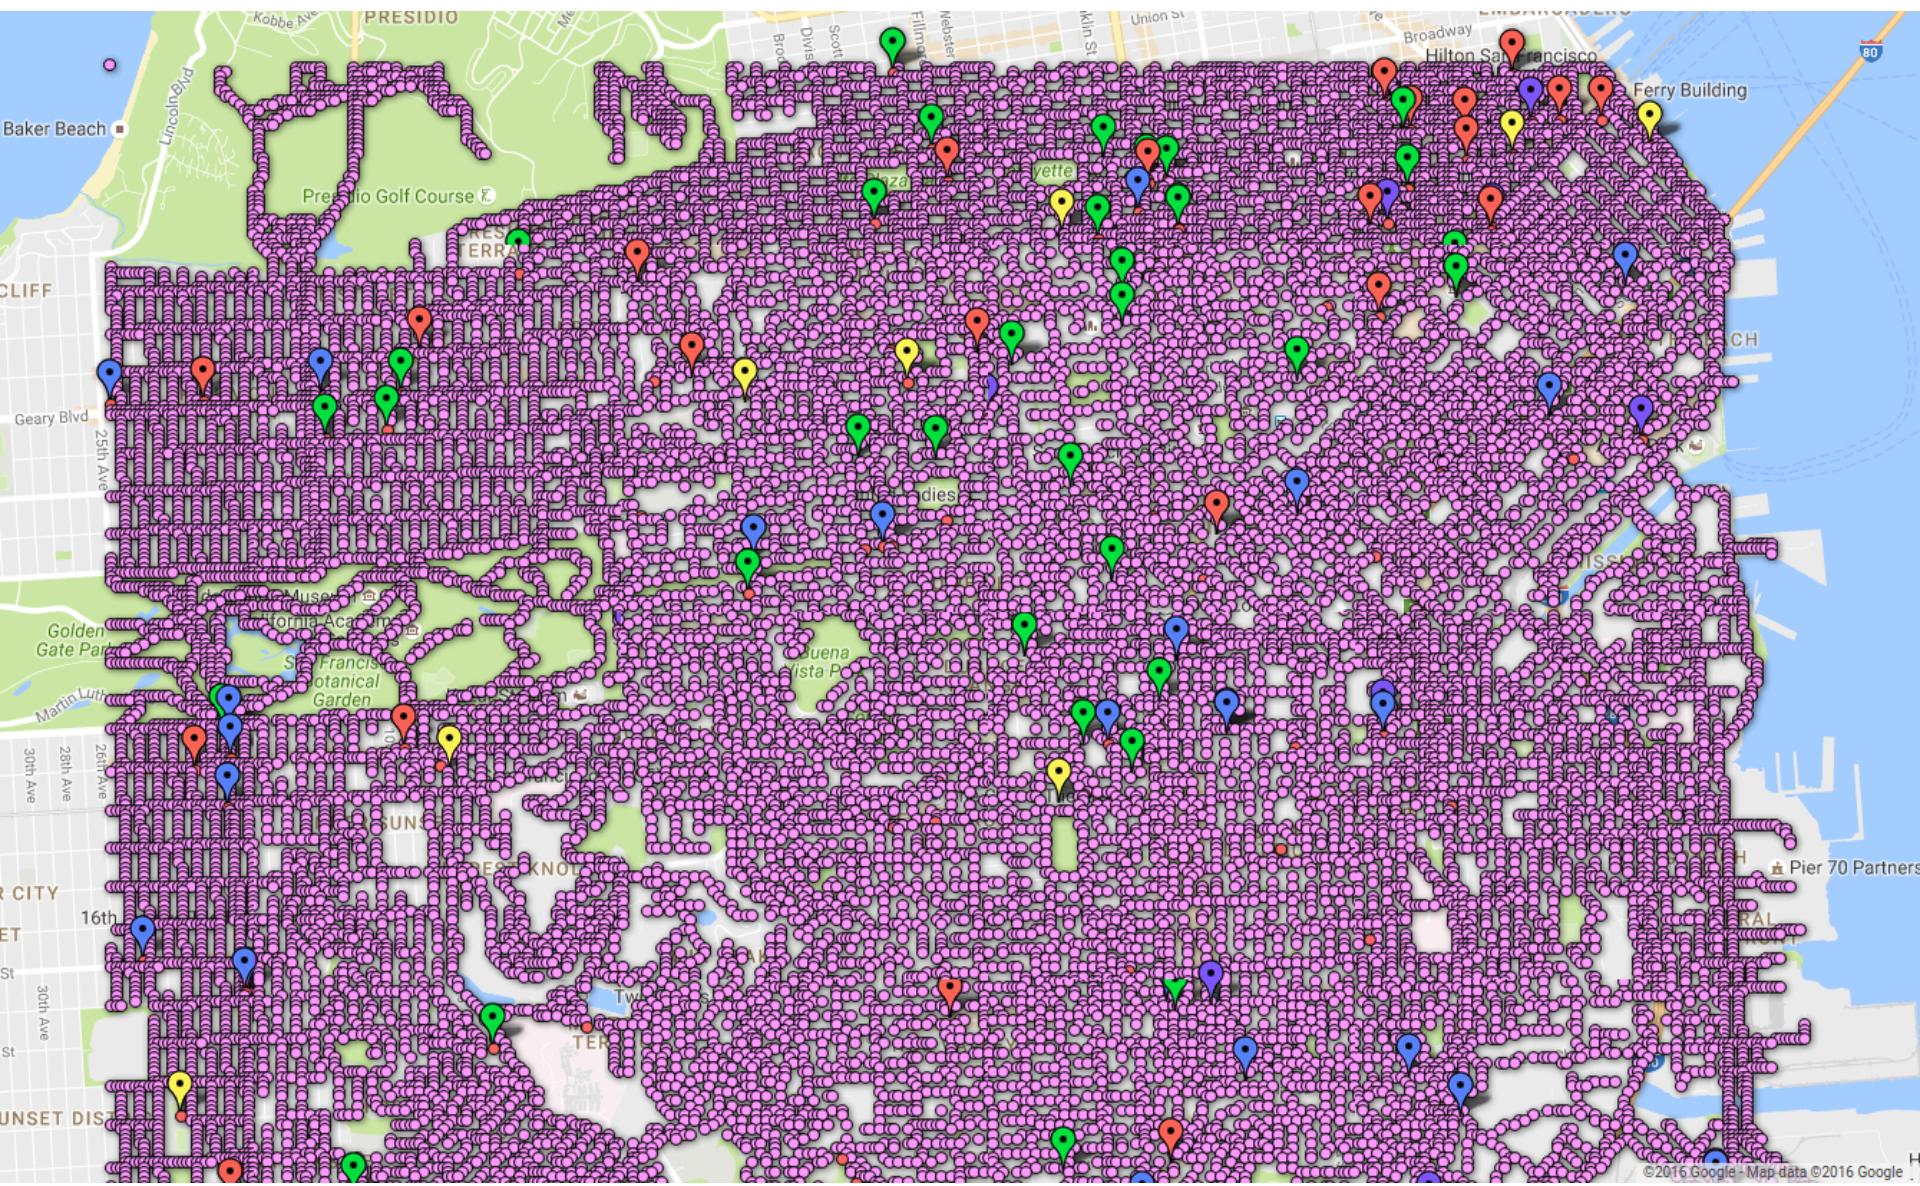




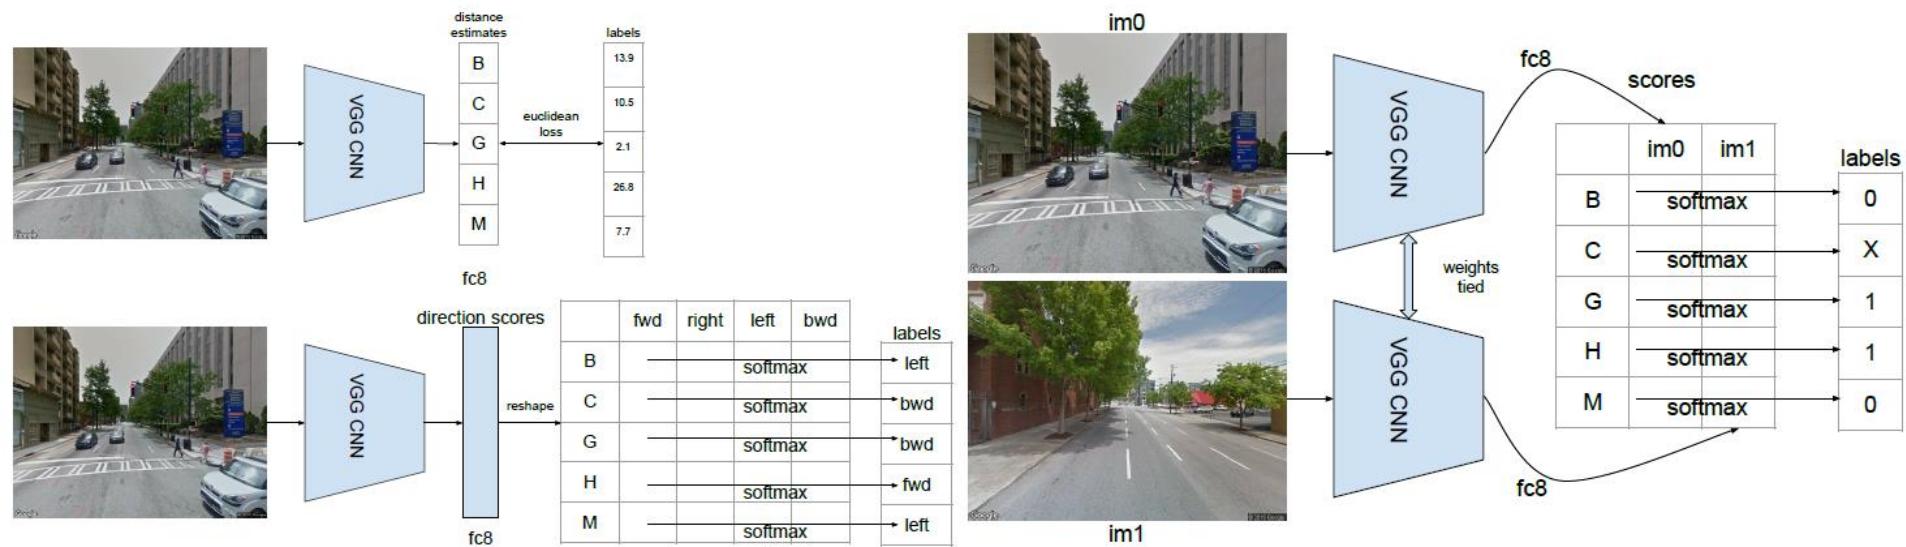
DeepNav



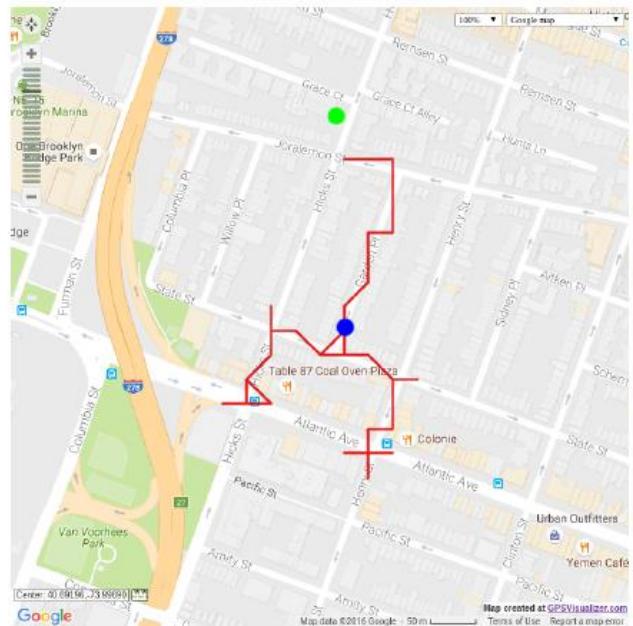
DeepNav: Learning to Navigate Large Cities
Samarth Brahmbhatt and James Hays. CVPR 2017



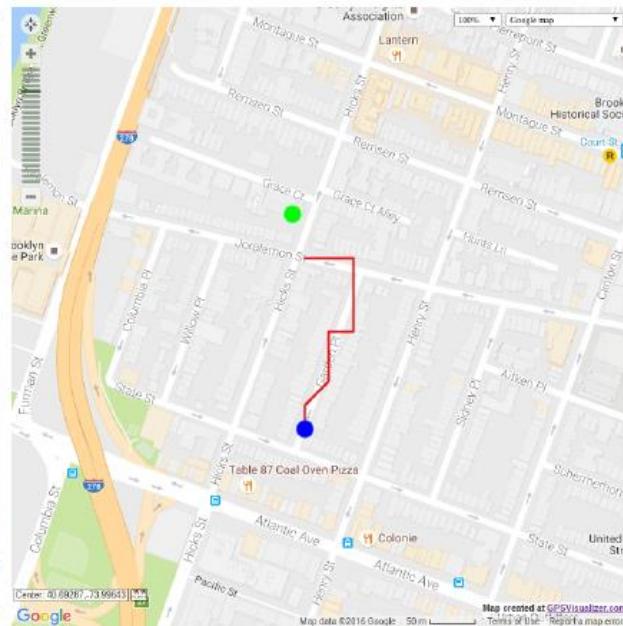
Network Architectures



Qualitative Results



Found, length = 60



Found, length = 10



Found, length = 520

Revisiting IM2GPS in the Deep Learning Era

Nam Vo, Nathan Jacobs, James Hays. ICCV 2017

The problem set up:

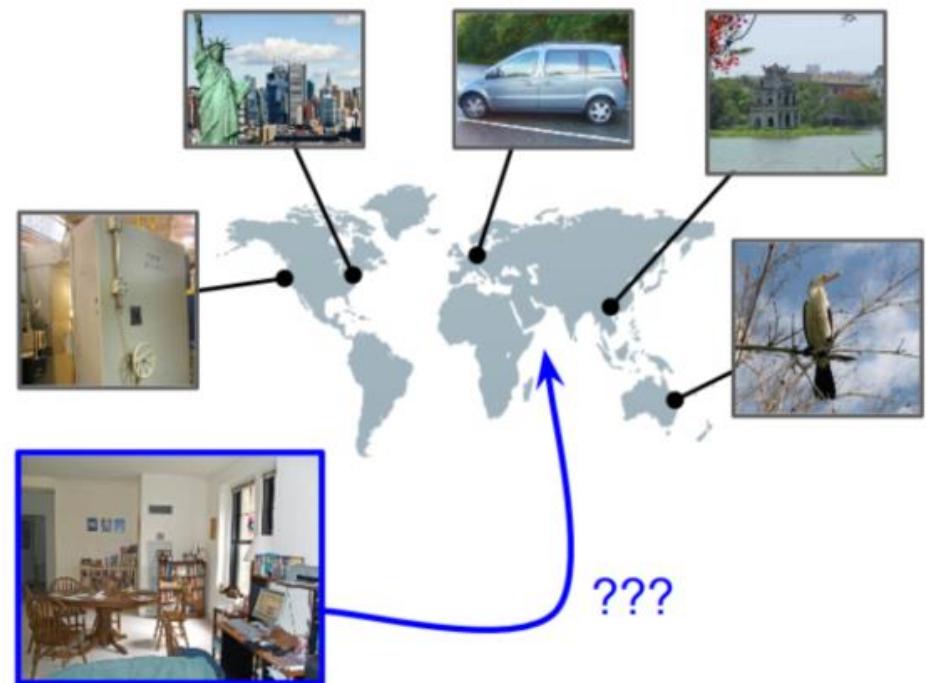
Give a large set of GPS-tagged images.

Learn to infer GPS coordinate of query images with unknown location.

Approaches:

Image retrieval

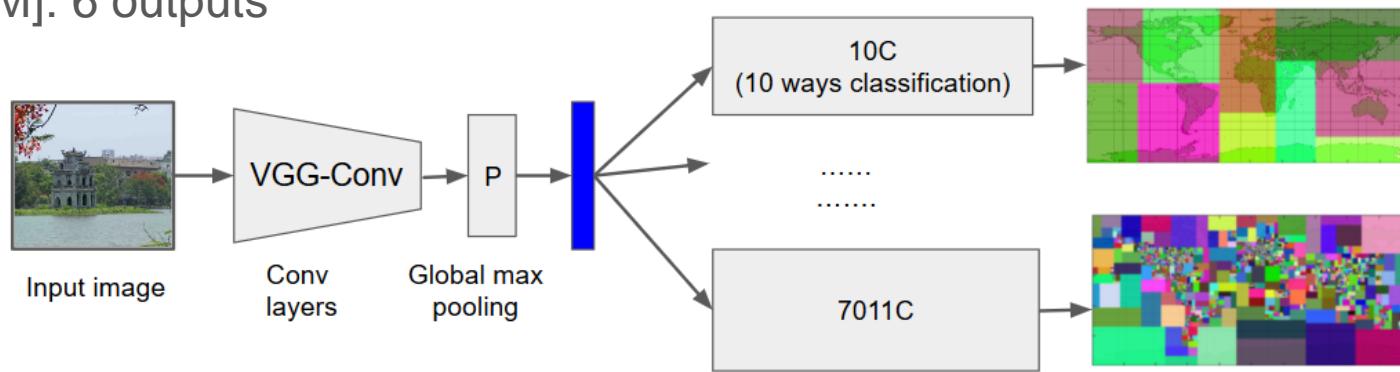
Image classification



Geolocation at planet scale

Caffe library, Vgg-16 imagenet initialization, training data: Im2GPS (~6m images)

Model [M]: 6 outputs



Model [L]: 7011C only

Model [L2]: 359C only

Model [R]: finetuned from [M] with ranking loss

Geolocalization at planet scale, Quantitative result

Table 1. Retrieval performance on Im2GPS test set. (Human* performance is average from 30 mturk workers over 940 trials, so it might not be directly comparable)

Threshold (km)	Street	City	Region	Country	Cont.
	1	25	200	750	2500
Human*			3.8	13.9	39.3
Im2GPS [9]		12.0	15.0	23.0	47.0
Im2GPS [10]	02.5	21.9	32.1	35.4	51.9
PlaNet [35]	08.4	24.5	37.6	53.6	71.3
[L] 7011C	06.8	21.9	34.6	49.4	63.7
[L] kNN, $\sigma=4$	12.2	33.3	44.3	57.4	71.3
... 28m database	14.4	33.3	47.7	61.6	73.4

Scribbler: Controlling Deep Image Synthesis with Sketch and Color



Patsorn Sangkloy, Jingwan Lu, Chen Fang , Fisher Yu, and James Hays. CVPR 2017

Training Data – (Mostly) Synthetic Sketches

Style Transfer

[Ulyanov et al. 2016]



XDoG

[Holger et al. 2012]



Photoshop

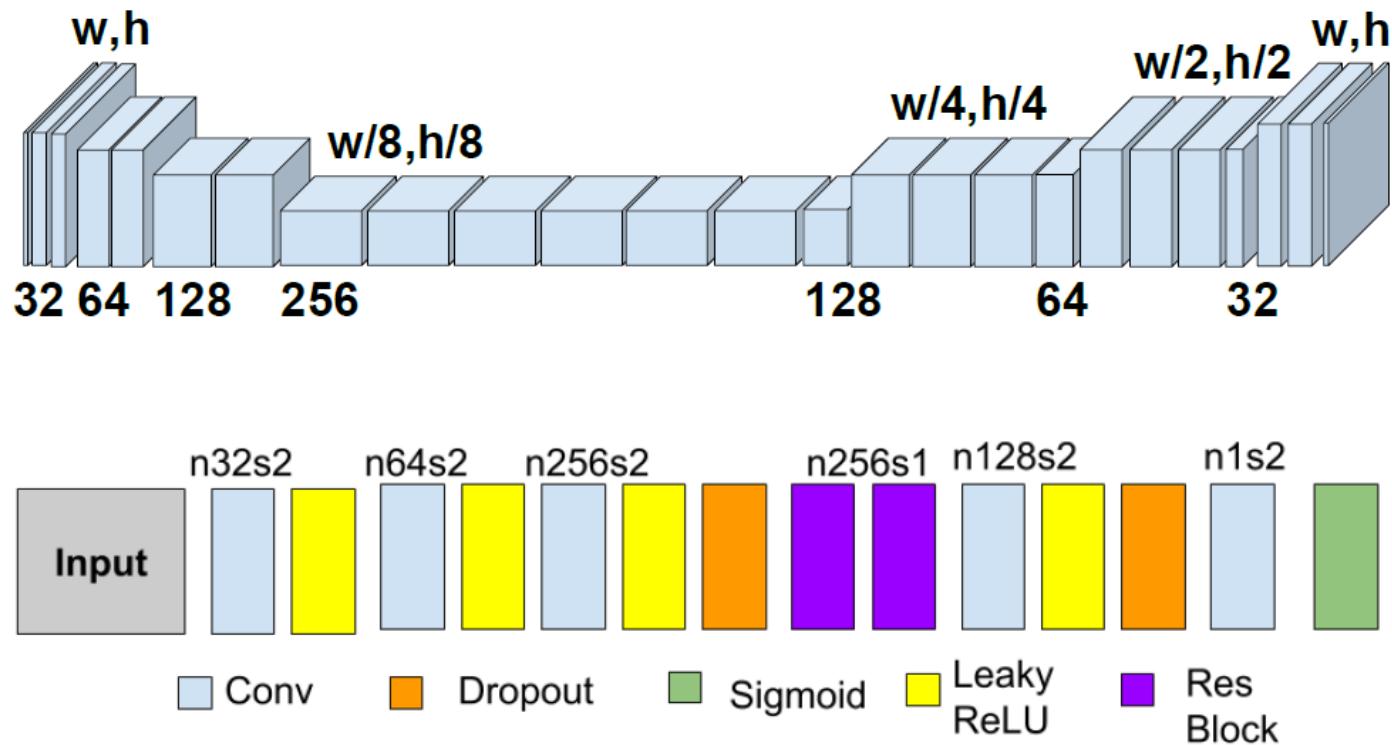


CUHK

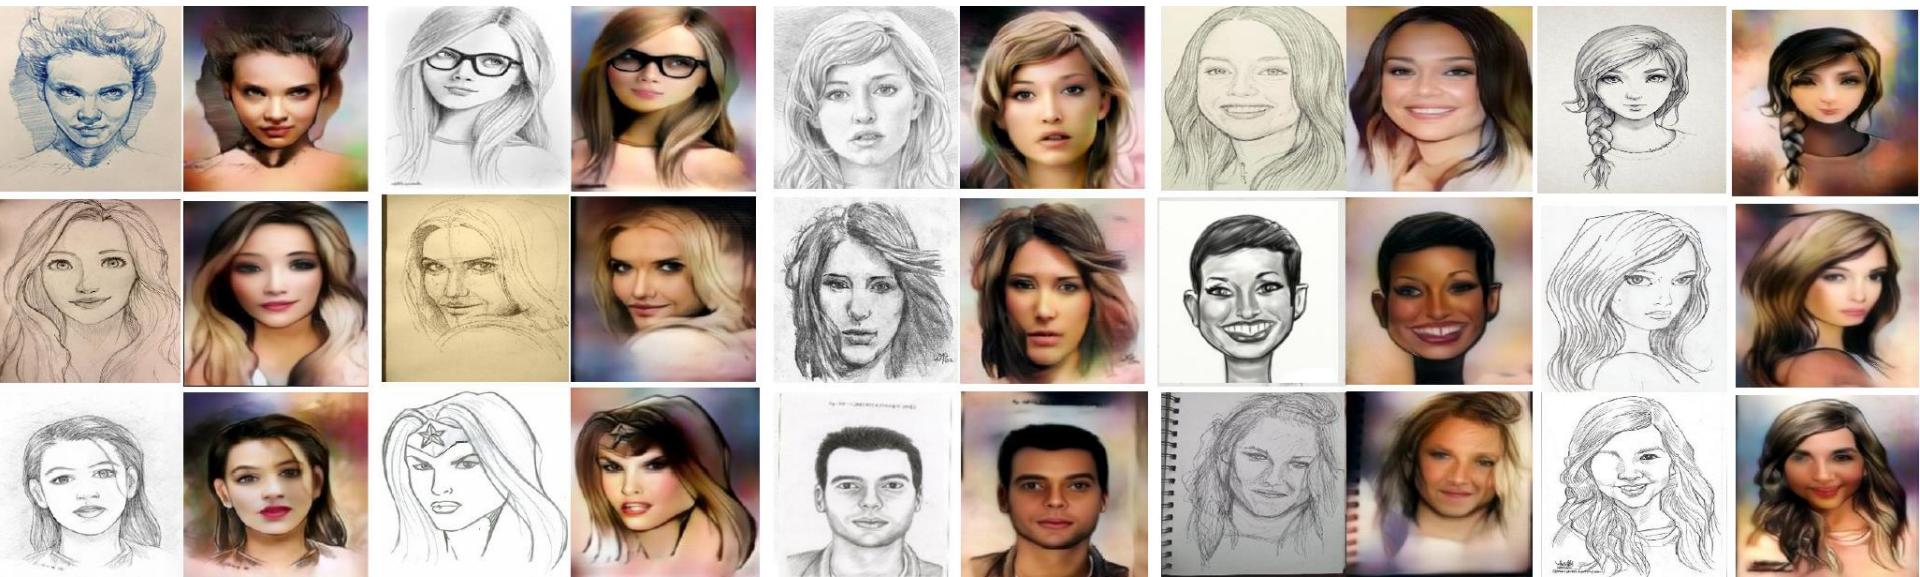
[Wang and Tang 2009]



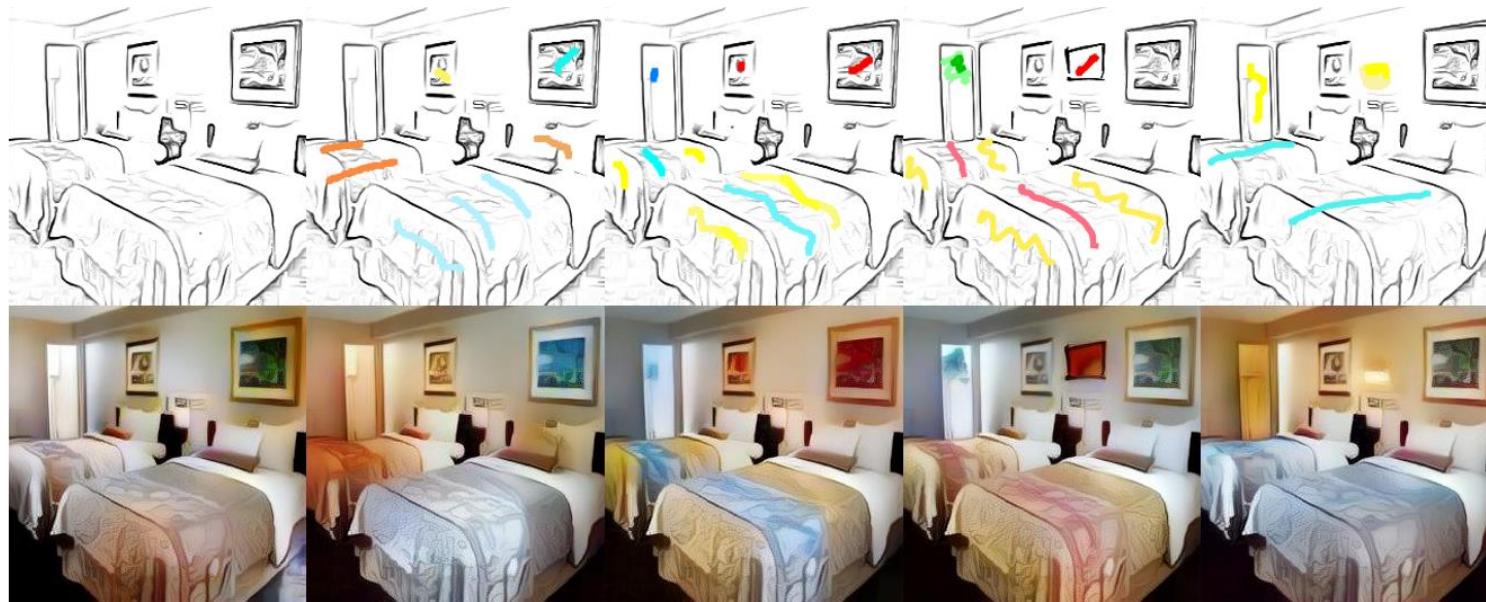
Network Architecture – Adversarial Learning



Results on held out sketches



Results on held out sketchesc



Groundtruth Image

Shoe

Tools



Current Mode: Color Stroke

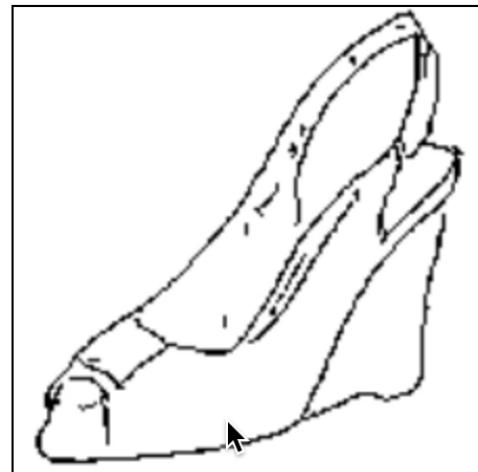
Color Picker

Brush Size



Enter sketch URL

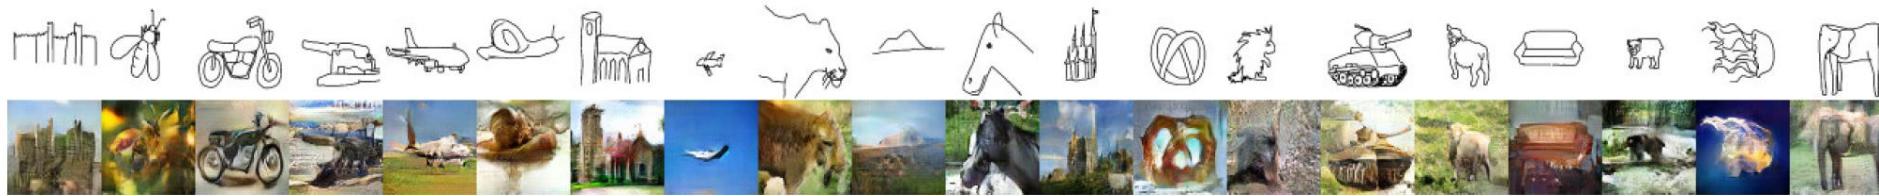
Input Sketch



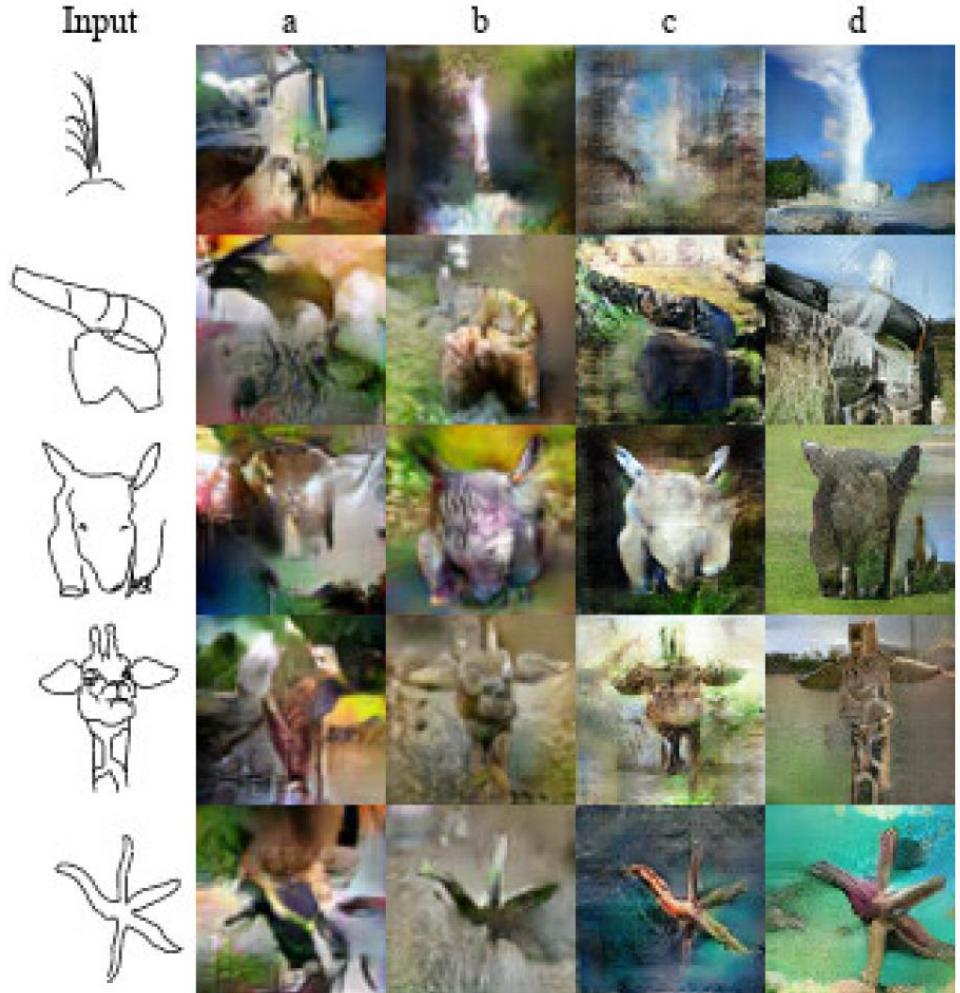
Final Result

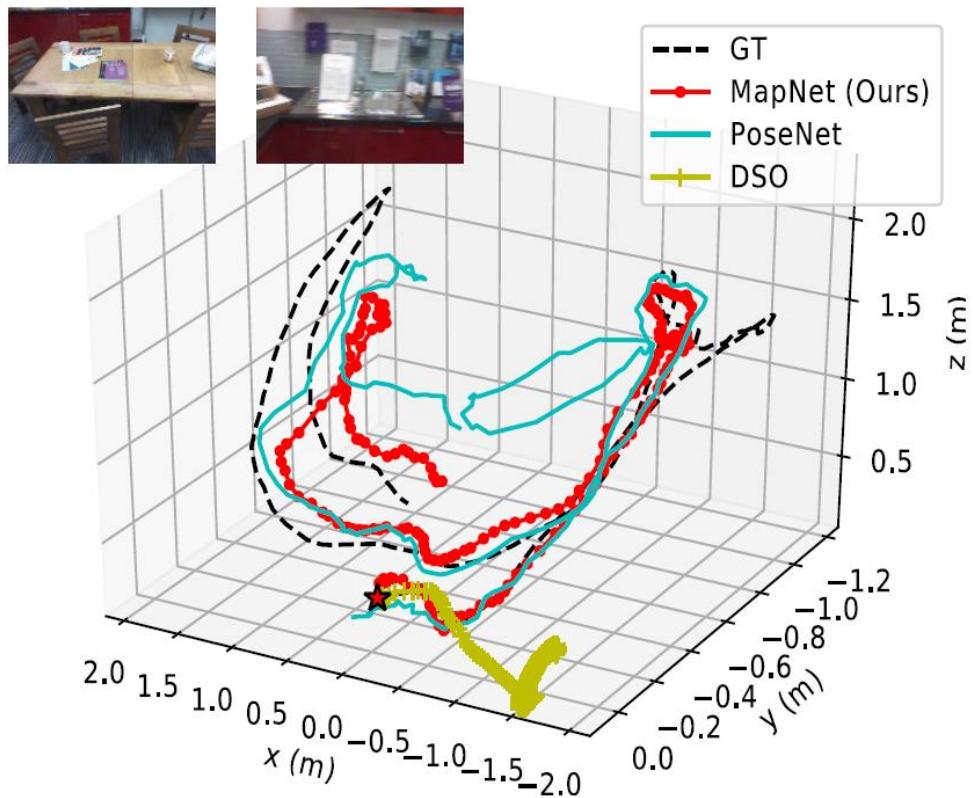
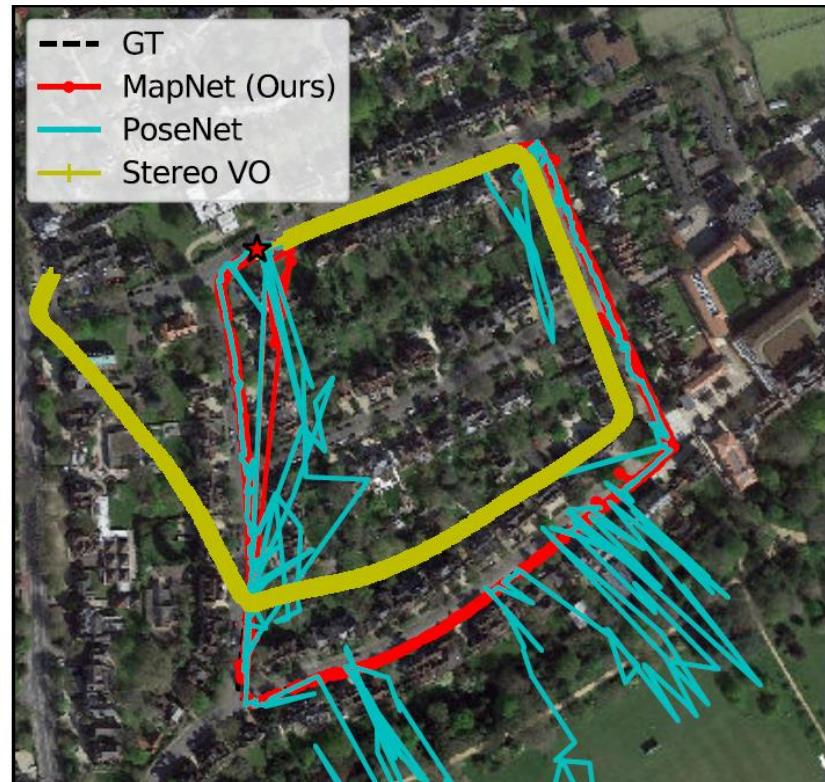


Given a simple sketch drawn by user, we hallucinate colors and shadings



SketchyGAN.
Wengling Chen and
James Hays.
CVPR 2018.





MapNet.
Samarth Brahmbhatt, Jinwei Gu, Kihwan Kim,
James Hays, and Jan Kautz.
CVPR 2018.

What is Computer Vision?

Computer Vision and Nearby Fields

- Computer Graphics: Models to Images
- Comp. Photography: Images to Images
- Computer Vision: Images to Models

Derogatory summary of computer vision:
Machine learning applied to visual data

Computer Vision

Make computers understand images and video **or any visual data.**



What kind of scene?

Where are the cars?

How far is the building?

...

Vision is really hard

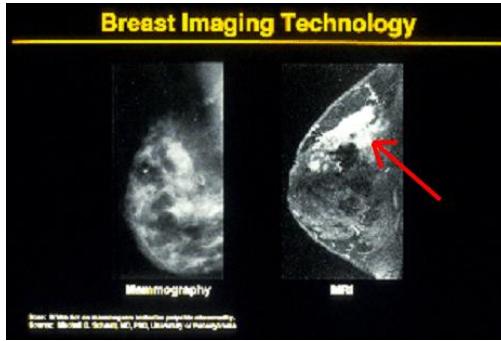
- Vision is an amazing feat of natural intelligence
 - Visual cortex occupies about 50% of Macaque brain
 - One third of human brain devoted to vision (more than anything else)



Why computer vision matters



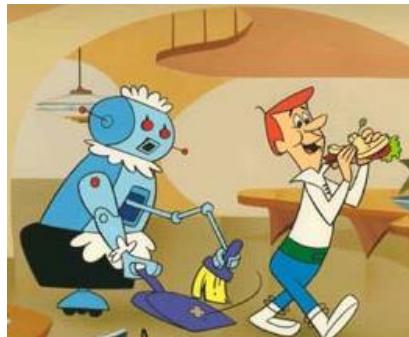
Safety



Health



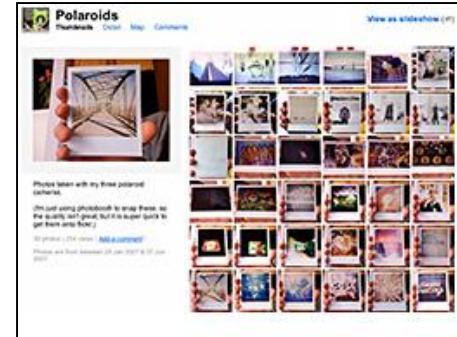
Security



Comfort



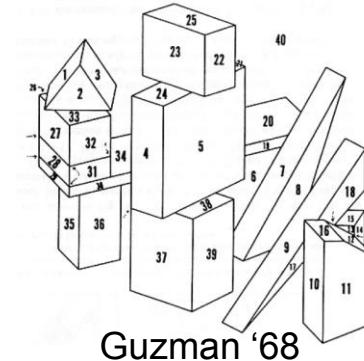
Fun



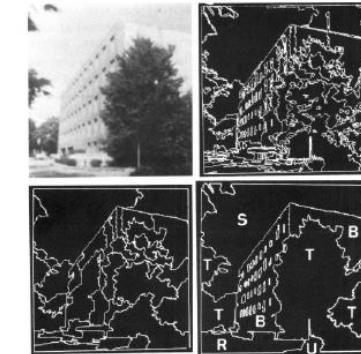
Access

Ridiculously brief history of computer vision

- 1966: Minsky assigns computer vision as an undergrad summer project
- 1960's: interpretation of synthetic worlds
- 1970's: some progress on interpreting selected images
- 1980's: ANNs come and go; shift toward geometry and increased mathematical rigor
- 1990's: face recognition; statistical analysis in vogue
- 2000's: broader recognition; large annotated datasets available; video processing starts
- 2010's: Deep learning with ConvNets
- 2020's: Widespread autonomous vehicles?
- 2030's: robot uprising?



Guzman '68



Ohta Kanade '78



Turk and Pentland '91

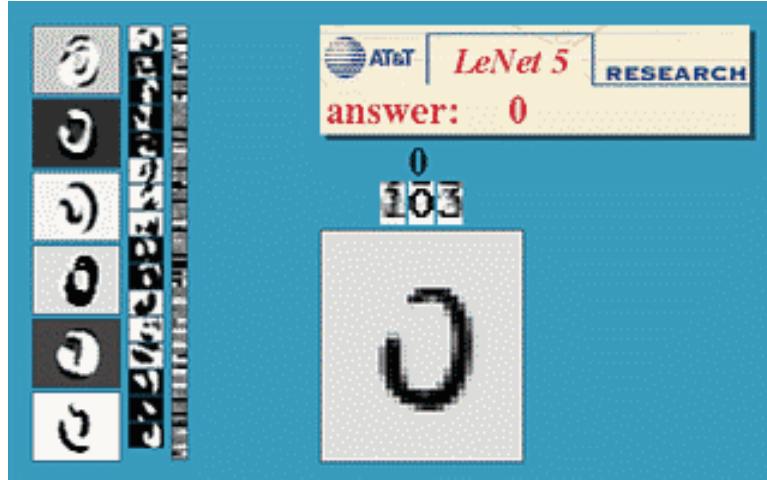
How vision is used now

- Examples of real world applications

Optical character recognition (OCR)

Technology to convert scanned docs to text

- If you have a scanner, it probably came with OCR software



Digit recognition, AT&T labs
<http://www.research.att.com/~yann/>



License plate readers
http://en.wikipedia.org/wiki/Automatic_number_plate_recognition

Face detection

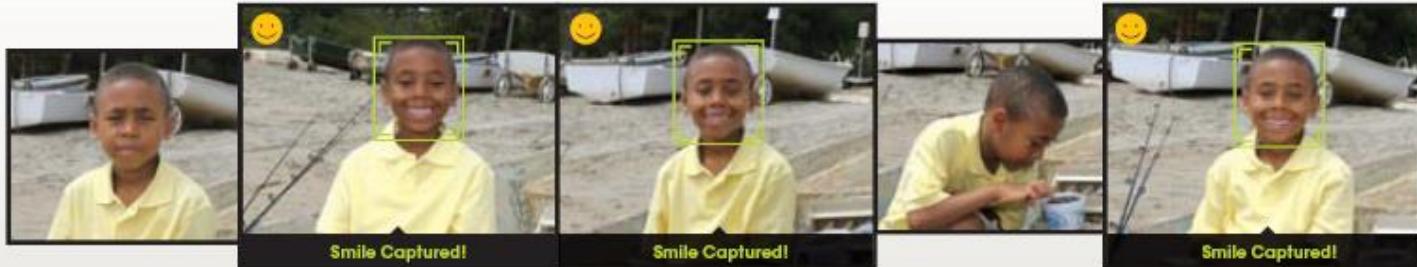


- Digital cameras detect faces

Smile detection

The Smile Shutter flow

Imagine a camera smart enough to catch every smile! In Smile Shutter Mode, your Cyber-shot® camera can automatically trip the shutter at just the right instant to catch the perfect expression.

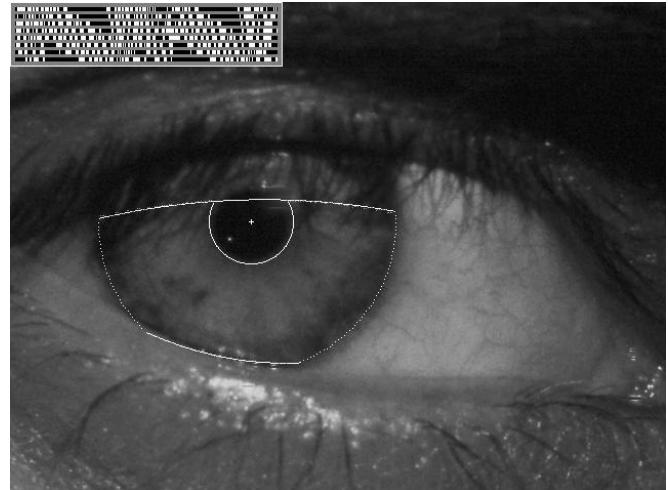
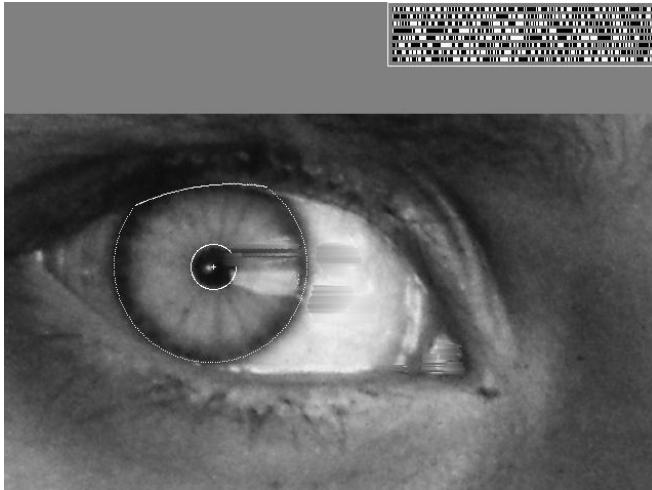


[Sony Cyber-shot® T70 Digital Still Camera](#)

Vision-based biometrics



“How the Afghan Girl was Identified by Her Iris Patterns” Read the [story](#)
[wikipedia](#)



Login without a password...



Fingerprint scanners on
many new laptops,
other devices



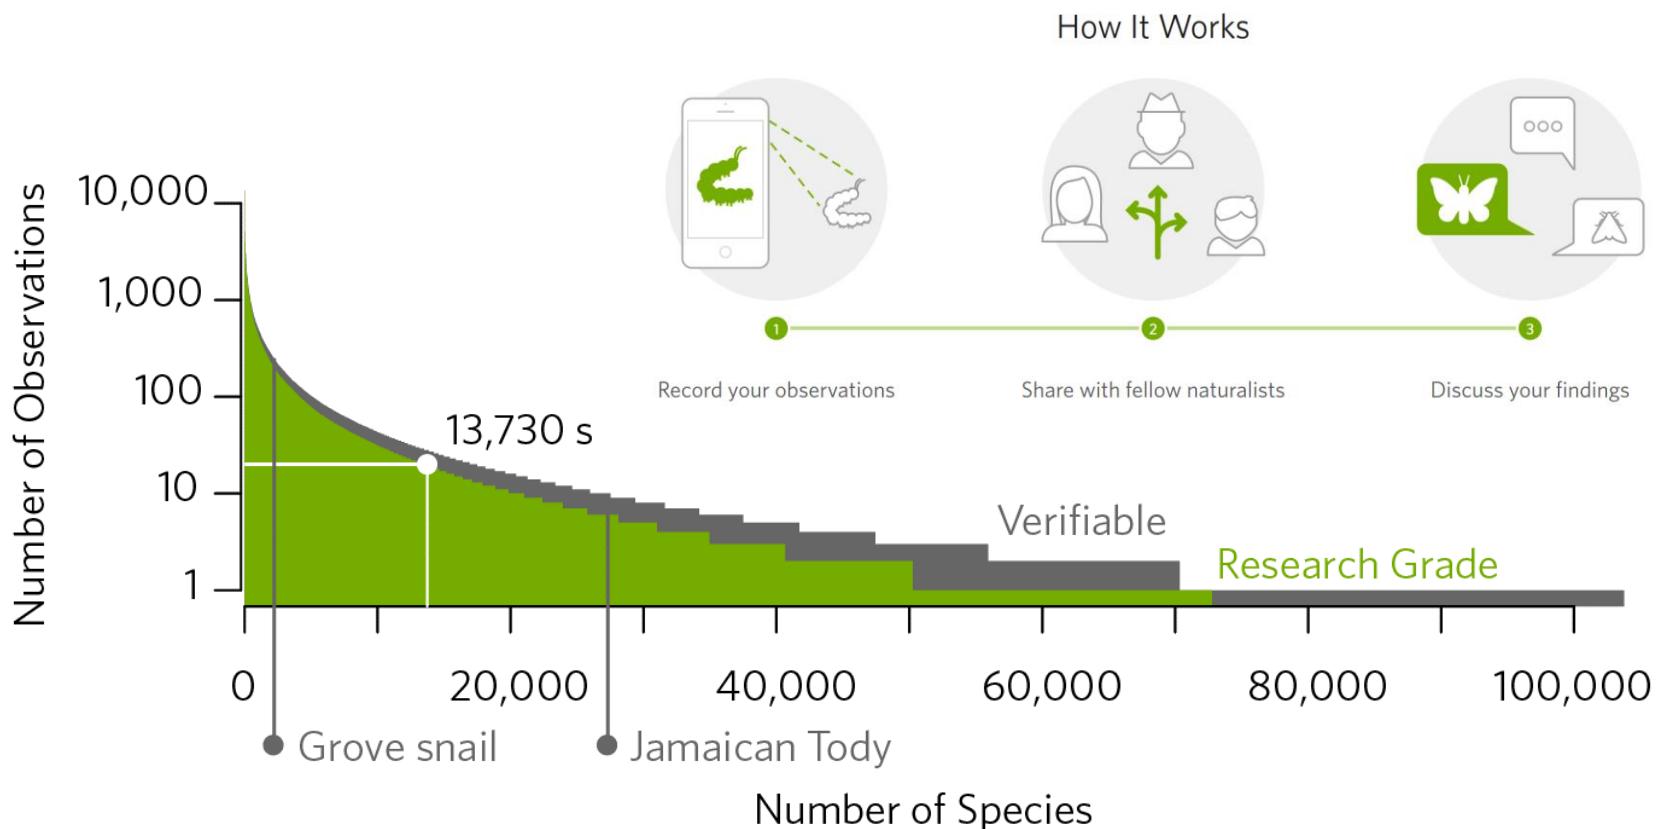
Face recognition systems now
beginning to appear more widely
<http://www.sensiblevision.com/>

Object recognition (in mobile phones)



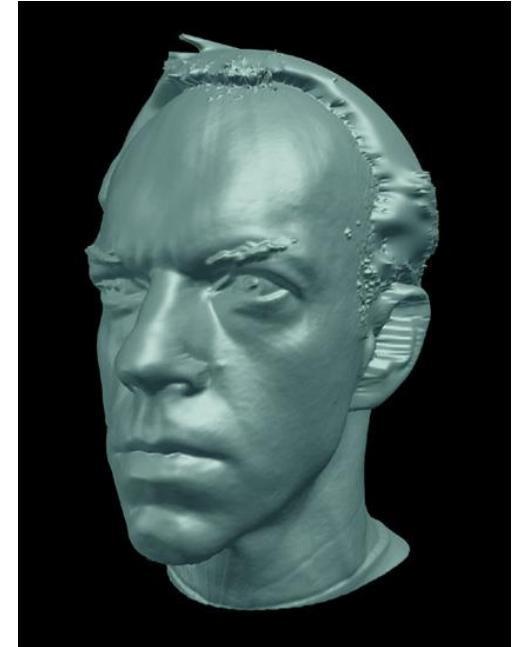
Point & Find, Nokia
Google Goggles

iNaturalist



https://www.inaturalist.org/pages/computer_vision_demo

Special effects: shape capture



The Matrix movies, ESC Entertainment, XYZRGB, NRC

Special effects: motion capture



Pirates of the Caribbean, Industrial Light and Magic

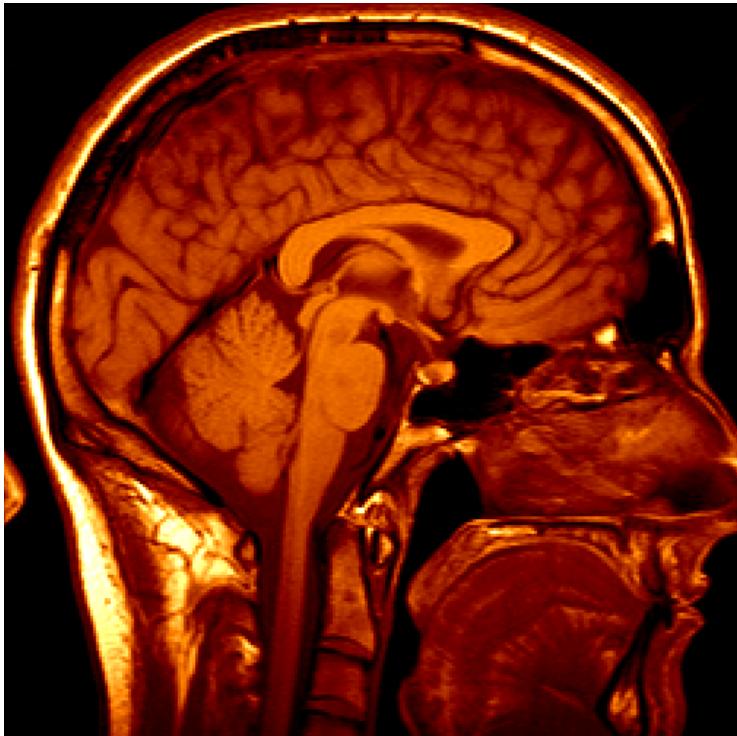
Sports



Sportvision first down line
Nice [explanation](#) on www.howstuffworks.com

<http://www.sportvision.com/video.html>

Medical imaging



3D imaging
MRI, CT



Image guided surgery
[Grimson et al., MIT](#)

Smart cars

Slide content courtesy of Amnon Shashua

The screenshot shows the Mobileye website's homepage. At the top, there are navigation tabs: 'manufacturer products' (with a right arrow), 'consumer products' (with left and right arrows), and 'News'. Below this is a main heading 'Our Vision. Your Safety.' with an image of a car from above showing three cameras: 'rear looking camera' (top left), 'forward looking camera' (top right), and 'side looking camera' (bottom). To the right is a 'News' sidebar with a list of articles and a 'Events' sidebar with links to auto shows.

News

- > [Mobileye Advanced Technologies Power Volvo Cars World First Collision Warning With Auto Brake System](#)
- > [Volvo: New Collision Warning with Auto Brake Helps Prevent Rear-end](#)

[> all news](#)

Events

- > [Mobileye at Equip Auto, Paris, France](#)
- > [Mobileye at SEMA, Las Vegas, NV](#)

[> read more](#)

EyeQ Vision on a Chip



[> read more](#)

Vision Applications



Road, Vehicle, Pedestrian Protection and more

[> read more](#)

AWS Advance Warning System



[> read more](#)

- Mobileye
 - Market Capitalization: 11 Billion dollars
 - Bought by Intel for 15 Billion dollars

Google cars



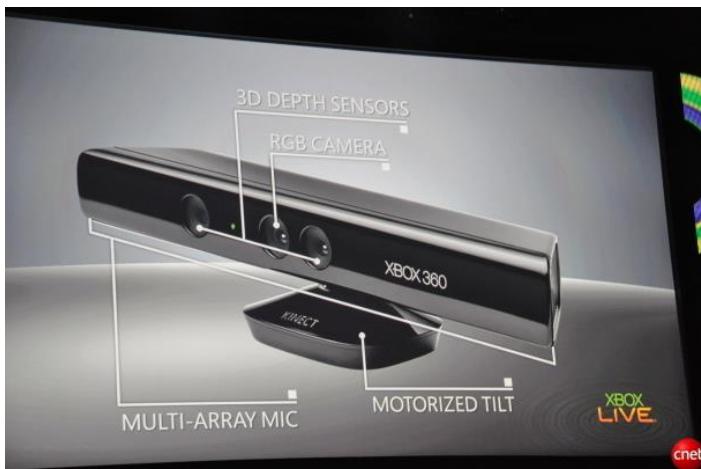
Oct 9, 2010. "[Google Cars Drive Themselves, in Traffic](#)". *The New York Times*. John Markoff

June 24, 2011. "[Nevada state law paves the way for driverless cars](#)". *Financial Post*. Christine Dobby

Aug 9, 2011, "[Human error blamed after Google's driverless car sparks five-vehicle crash](#)". *The Star (Toronto)*

Interactive Games: Kinect

- Object Recognition:
<http://www.youtube.com/watch?feature=iv&v=fQ59dXOo63o>
- Mario: <http://www.youtube.com/watch?v=8CTJL5IUjHg>
- 3D: <http://www.youtube.com/watch?v=7QrnwoO1-8A>
- Robot: <http://www.youtube.com/watch?v=w8BmgtMKFbY>



Augmented Reality and Virtual Reality



Magic Leap, Oculus, Hololens, etc.

Industrial robots



Vision-guided robots position nut runners on wheels

Vision in space



[NASA'S Mars Exploration Rover Spirit](#) captured this westward view from atop a low plateau where Spirit spent the closing months of 2007.

Vision systems (JPL) used for several tasks

- Panorama stitching
- 3D terrain modeling
- Obstacle detection, position tracking
- For more, read "[Computer Vision on Mars](#)" by Matthies et al.

Amazon Prime Air

<https://www.amazon.com/b?node=8037720011>

Skydio

<https://www.skydio.com/>

State of the art today?

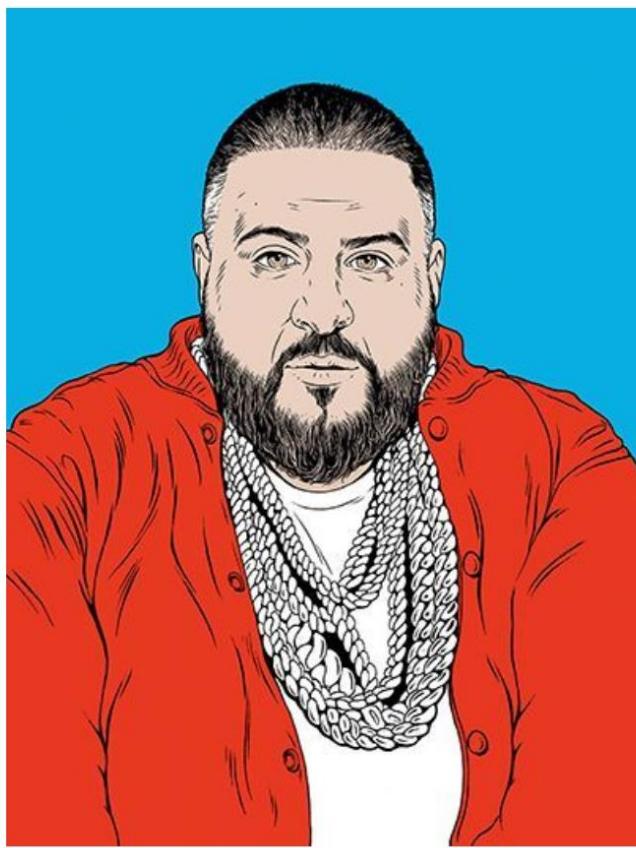
With enough training data, computer vision nearly matches human vision at most recognition tasks

Deep learning has been an enormous disruption to the field. More and more techniques are being “deepified”.

WIRED

100

WHO'S SHAPING THE DIGITAL WORLD?



DJ Khaled

Credit [Louise Zergaeng Pomeroy](#)

73. DJ Khaled

Snapchat icon; DJ and producer

Louisiana-born Khaled Mohamed Khaled, aka DJ Khaled, cut his musical chops in the early 00s as a host for Miami urban music radio WEDR. He proceeded to build a solid if not dazzling career as a mixtape DJ and music producer (he founded his label We The Best Music Group in 2008, and was appointed president of Def Jam South in 2009).

69. Geoffrey Hinton

Psychologist, computer scientist; researcher, Google Toronto

British-born Hinton has been dubbed the "godfather of deep learning". The Cambridge-educated cognitive psychologist and computer scientist started being an ardent believer in the potential of neural networks and deep learning in the 80s, when those technologies enjoyed little support in the wider AI community.

But he soldiered on: in 2004, with support from the Canadian Institute for Advanced Research, he launched a University of Toronto programme in neural computation and adaptive perception, where, with a group of researchers, he carried on investigating how to create computers that could behave like brains.

Hinton's work – in particular his algorithms that train multilayered neural networks – caught the attention of tech giants in Silicon Valley, which realised how deep learning could be applied to voice recognition, predictive search and machine vision.

The spike in interest prompted him to launch a free course on neural networks on e-learning platform Coursera in 2012. Today, 68-year-old Hinton is chair of machine learning at the University of Toronto and moonlights at Google, where he has been using deep learning to help build internet tools since 2013.

63. Yann Lecun

Director of AI research, Facebook, Menlo Park

LeCun is a leading expert in deep learning and heads up what, for Facebook, could be a hugely significant source of revenue: understanding its user's intentions.

62. Richard Branson

Founder, Virgin Group, London

Branson saw his personal fortune grow £550 million when Alaska Air bought Virgin America for \$2.6 billion in April. He is pressing on with civilian space travel with [Virgin Galactic](#).

61. Taylor Swift

Entertainer, Los Angeles





Credit [Google DeepMind](#)

Google-backed startup DeepMind Technologies has built an [artificial intelligence](#) agent that can learn to successfully play 49 classic Atari games by itself, with minimal input.



The story of AlphaGo so far

AlphaGo is the first computer program to defeat a professional human Go player, the first program to defeat a Go world champion, and arguably the strongest Go player in history.

AlphaGo's first formal match was against the reigning 3-times European Champion, Mr Fan Hui, in October 2015. Its 5-0 win was the first ever against a Go professional, and the results were published in full technical detail in the international journal, [Nature](#). AlphaGo then went on to compete against legendary player Mr Lee Sedol, winner of 18 world titles and widely considered to be the greatest player of the past decade.

AlphaGo's 4-1 victory in Seoul, South Korea, in March 2016 was watched by over 200 million people worldwide. It was a landmark achievement that experts agreed was a decade ahead of its time, and earned AlphaGo a 9 dan professional ranking (the highest certification) - the first time a computer Go player had ever received the accolade.

During the games, AlphaGo played a handful of [highly inventive winning moves](#), several of which - including move 37 in game two - were so surprising they overturned hundreds of years of received wisdom, and have since been examined extensively by players of all levels. In the course of winning, AlphaGo somehow taught the world completely new knowledge about perhaps the most studied and contemplated game in history.

Since then, AlphaGo has continued to surprise and amaze. In January 2017, an improved AlphaGo version was revealed as the online player "Master" which achieved [60 straight wins in online fast time-control games](#) against top international Go players.

In May 2017, Alpha Go took part in The Future of Go Summit in the birthplace of Go, China, to delve deeper into the mysteries of Go in a spirit of mutual collaboration with the country's top players. You can read more about the five day summit [here](#).



[> More on The Future of Go Summit in this video](#)



[> Watch the video here](#)

Course Syllabus (tentative)

<http://www.cc.gatech.edu/~hays/compvision>

Grading

- 80% programming projects (6 total)
- 20% quizzes (2 total)

Scope of CS 4476

Image Processing
Geometric Reasoning
Recognition
Deep Learning

Computer Vision

Machine Learning

Robotics

Human Computer Interaction

Graphics

Medical Imaging

Computational Photography

Neuroscience

Optics

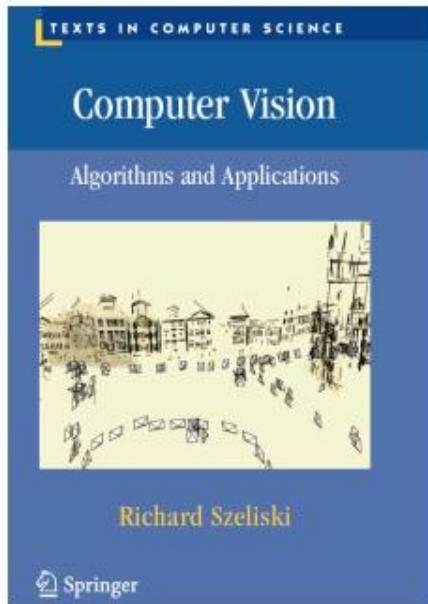
Course Topics

- Interpreting Intensities
 - What determines the brightness and color of a pixel?
 - How can we use image filters to extract meaningful information from the image?
- Correspondence and Alignment
 - How can we find corresponding points in objects or scenes?
 - How can we estimate the transformation between them?
- Grouping and Segmentation
 - How can we group pixels into meaningful regions?
- Categorization and Object Recognition
 - How can we represent images and categorize them?
 - How can we recognize categories of objects?
- Advanced Topics
 - Action recognition, 3D scenes and context, human-in-the-loop vision...

Textbook

Computer Vision: Algorithms and Applications

© 2010 [Richard Szeliski](#), Microsoft Research



<http://szeliski.org/Book/>

Prerequisites

- **Linear algebra**, basic calculus, and probability
- Experience with image processing will help but is not necessary
- Experience with Python or Python-like languages will help

Projects

- Image Filtering and Hybrid Images
- Local Feature Matching
- Camera Calibration and Fundamental Matrix Estimation with RANSAC
- Scene Recognition with Bag of Words
- Object Detection with a Sliding Window
- Recognition with Deep Learning

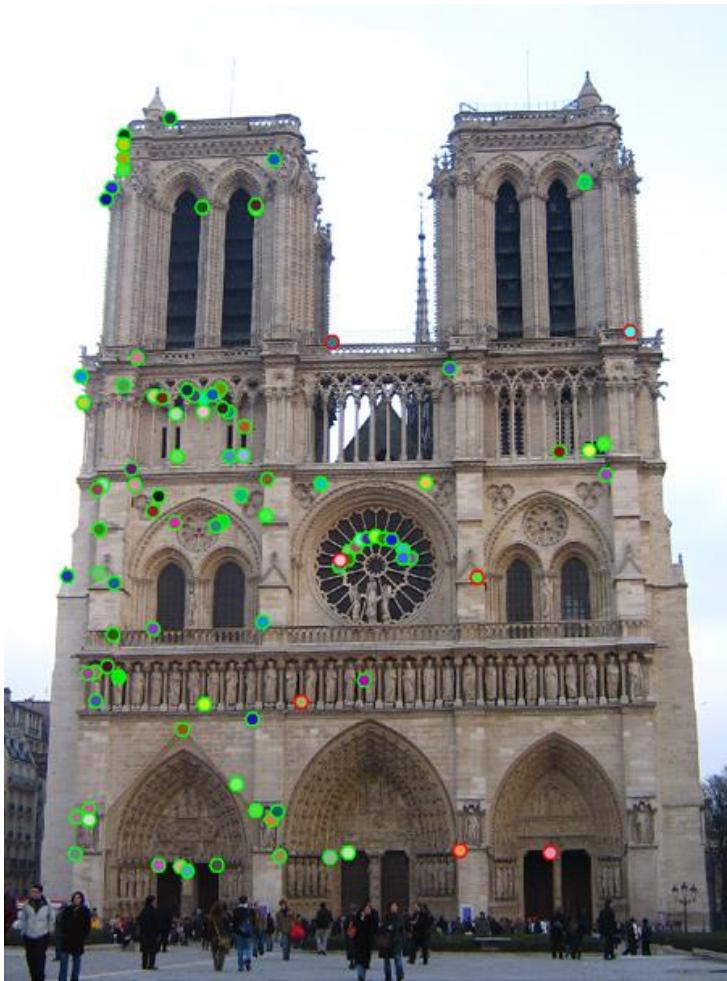
Proj1: Image Filtering and Hybrid Images

- Implement image filtering to separate high and low frequencies
- Combine high frequencies and low frequencies from different images to create an image with scale-dependent interpretation



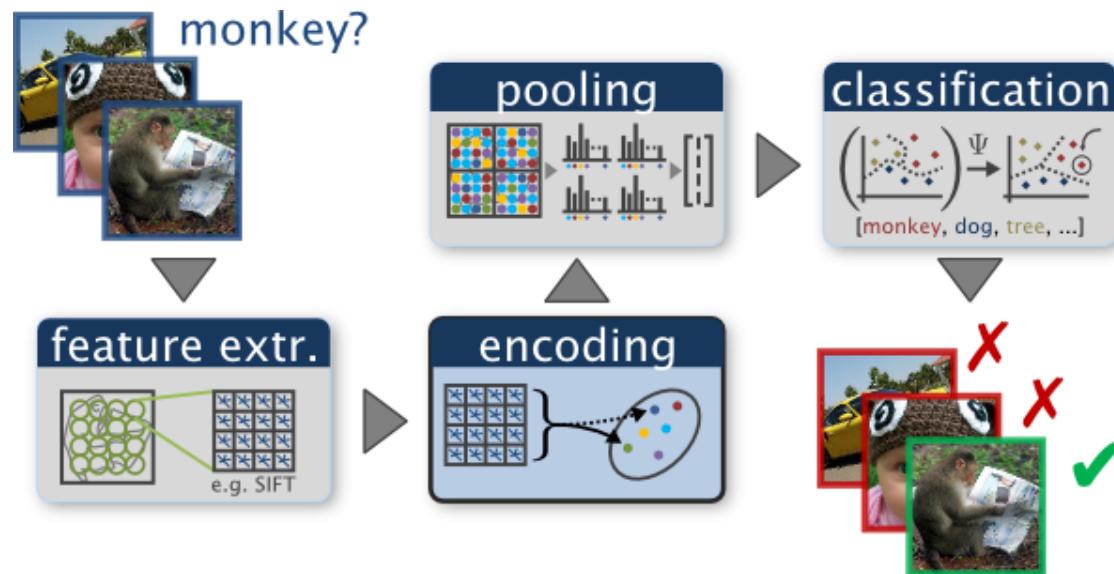
Proj2: Local Feature Matching

- Implement interest point detector, SIFT-like local feature descriptor, and simple matching algorithm.



Proj4: Scene Recognition with Bag of Words

- Quantize local features into a “vocabulary”, describe images as histograms of “visual words”, train classifiers to recognize scenes based on these histograms.



Proj5: Object Detection with a Sliding Window

- Train a face detector based on positive examples and “mined” hard negatives, detect faces at multiple scales and suppress duplicate detections.



Course Syllabus (tentative)

<http://www.cc.gatech.edu/~hays/compvision>