



The blue and green colors are actually the same

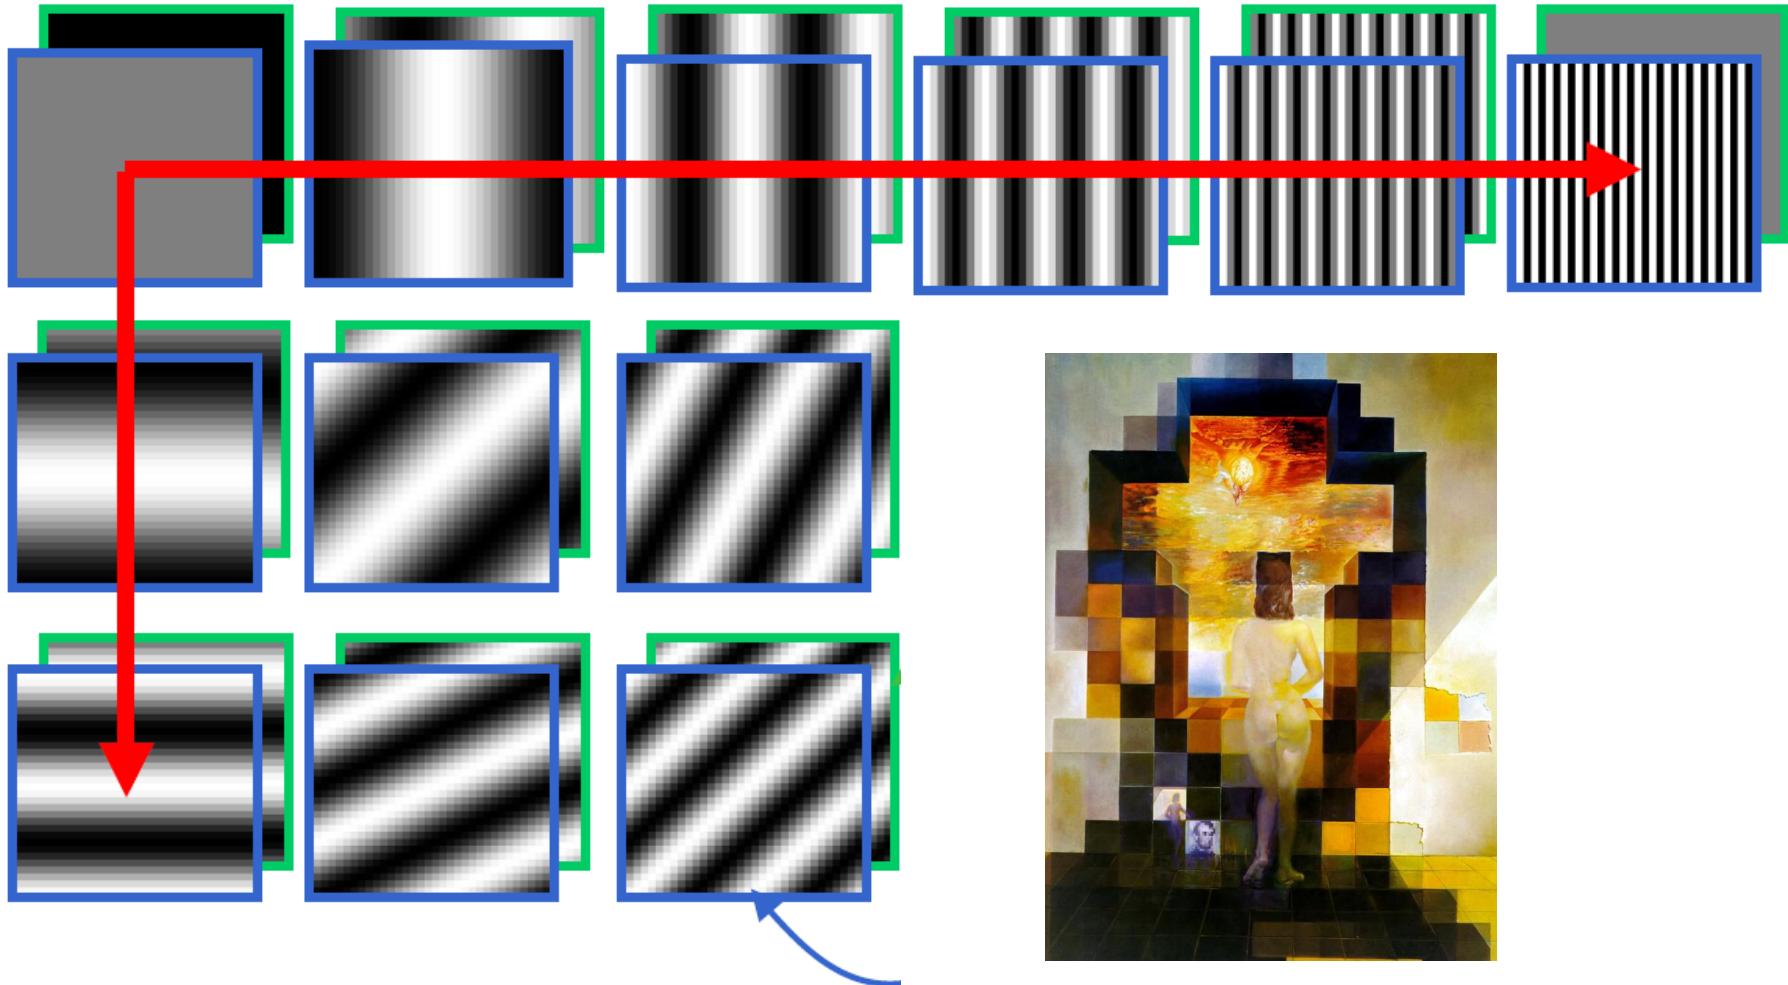


<http://blogs.discovermagazine.com/badastronomy/2009/06/24/the-blue-and-the-green/>

# Recap: Frequency Domain

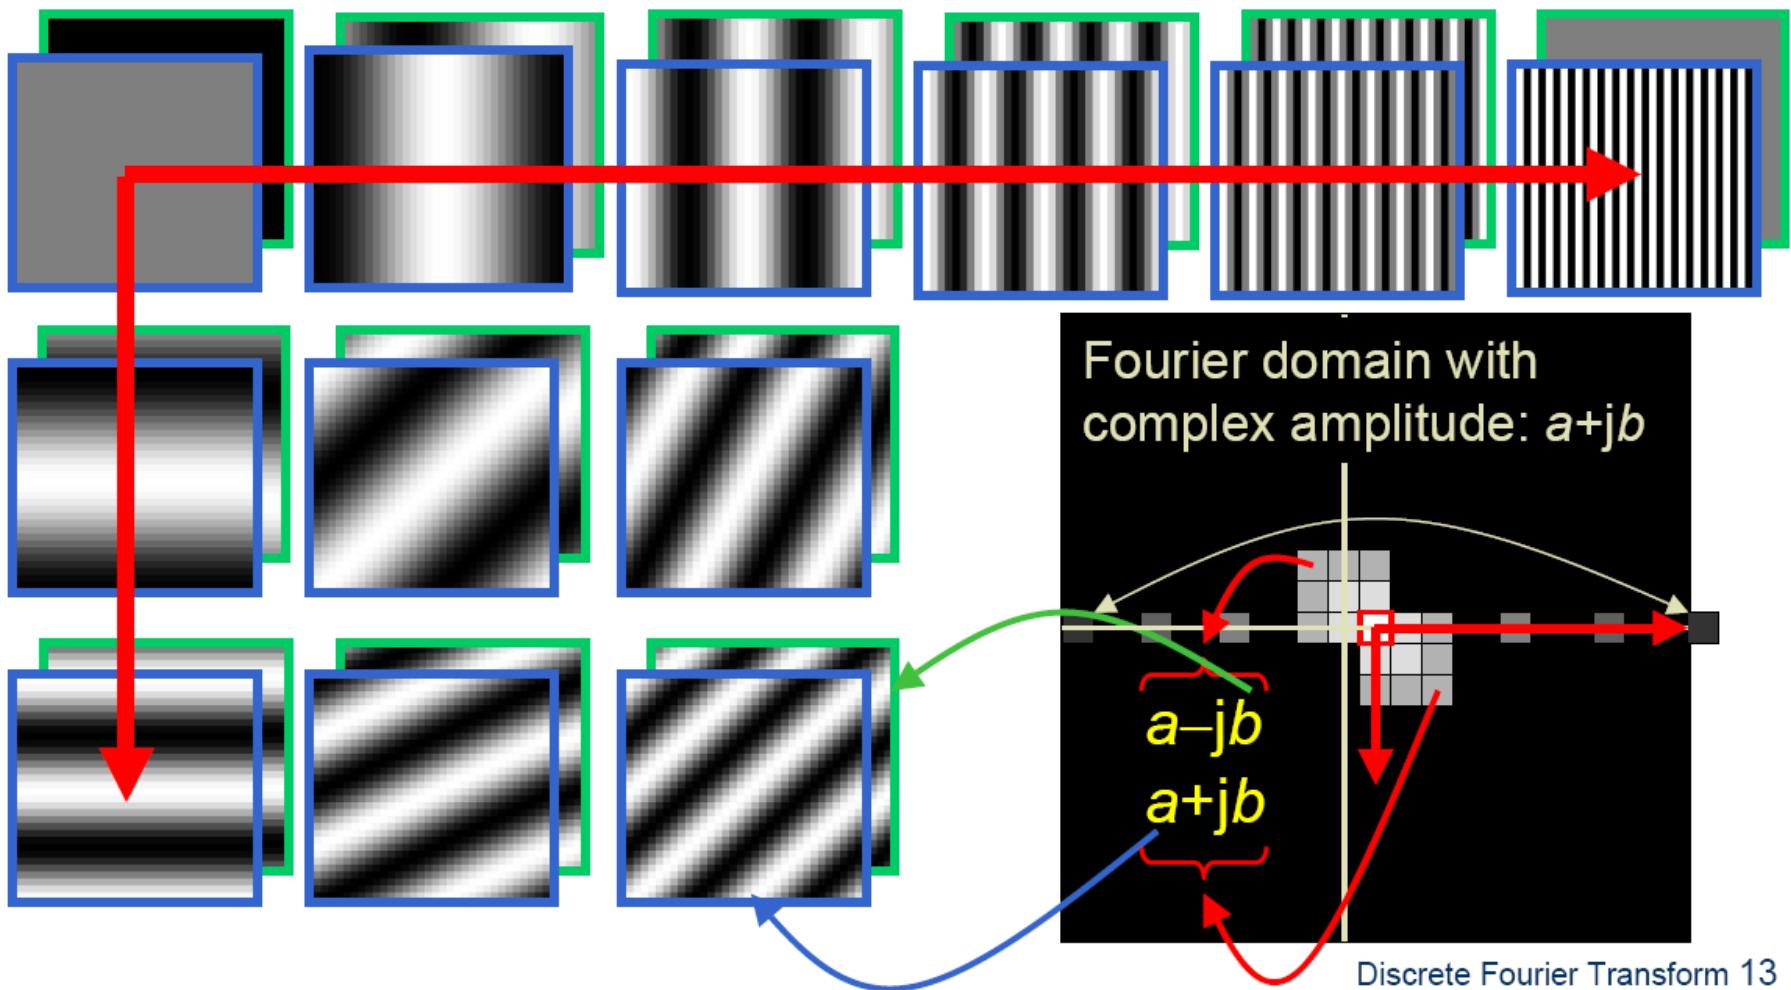
# Fourier Bases

Teases away fast vs. slow changes in the image.



This change of basis is the Fourier Transform

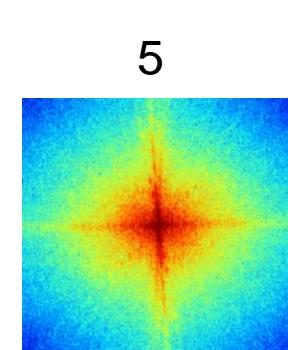
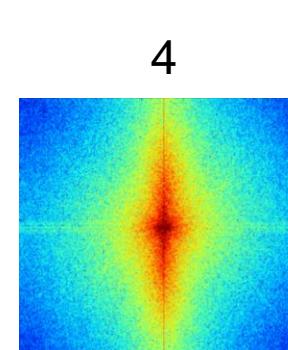
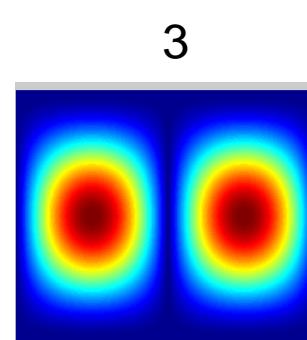
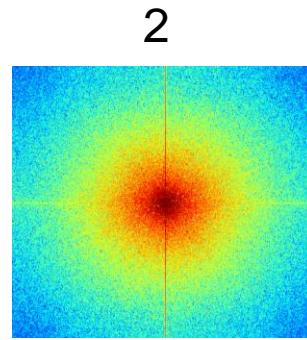
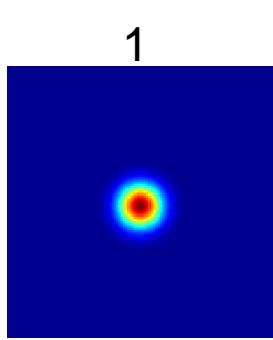
# Fourier Bases



in Matlab, check out: `imagesc(log(abs(fftshift(fft2(im)))));`

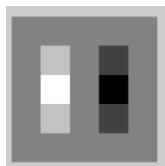
# Review

1. Match the spatial domain image to the Fourier magnitude image



B

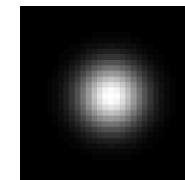
A



C



D



E

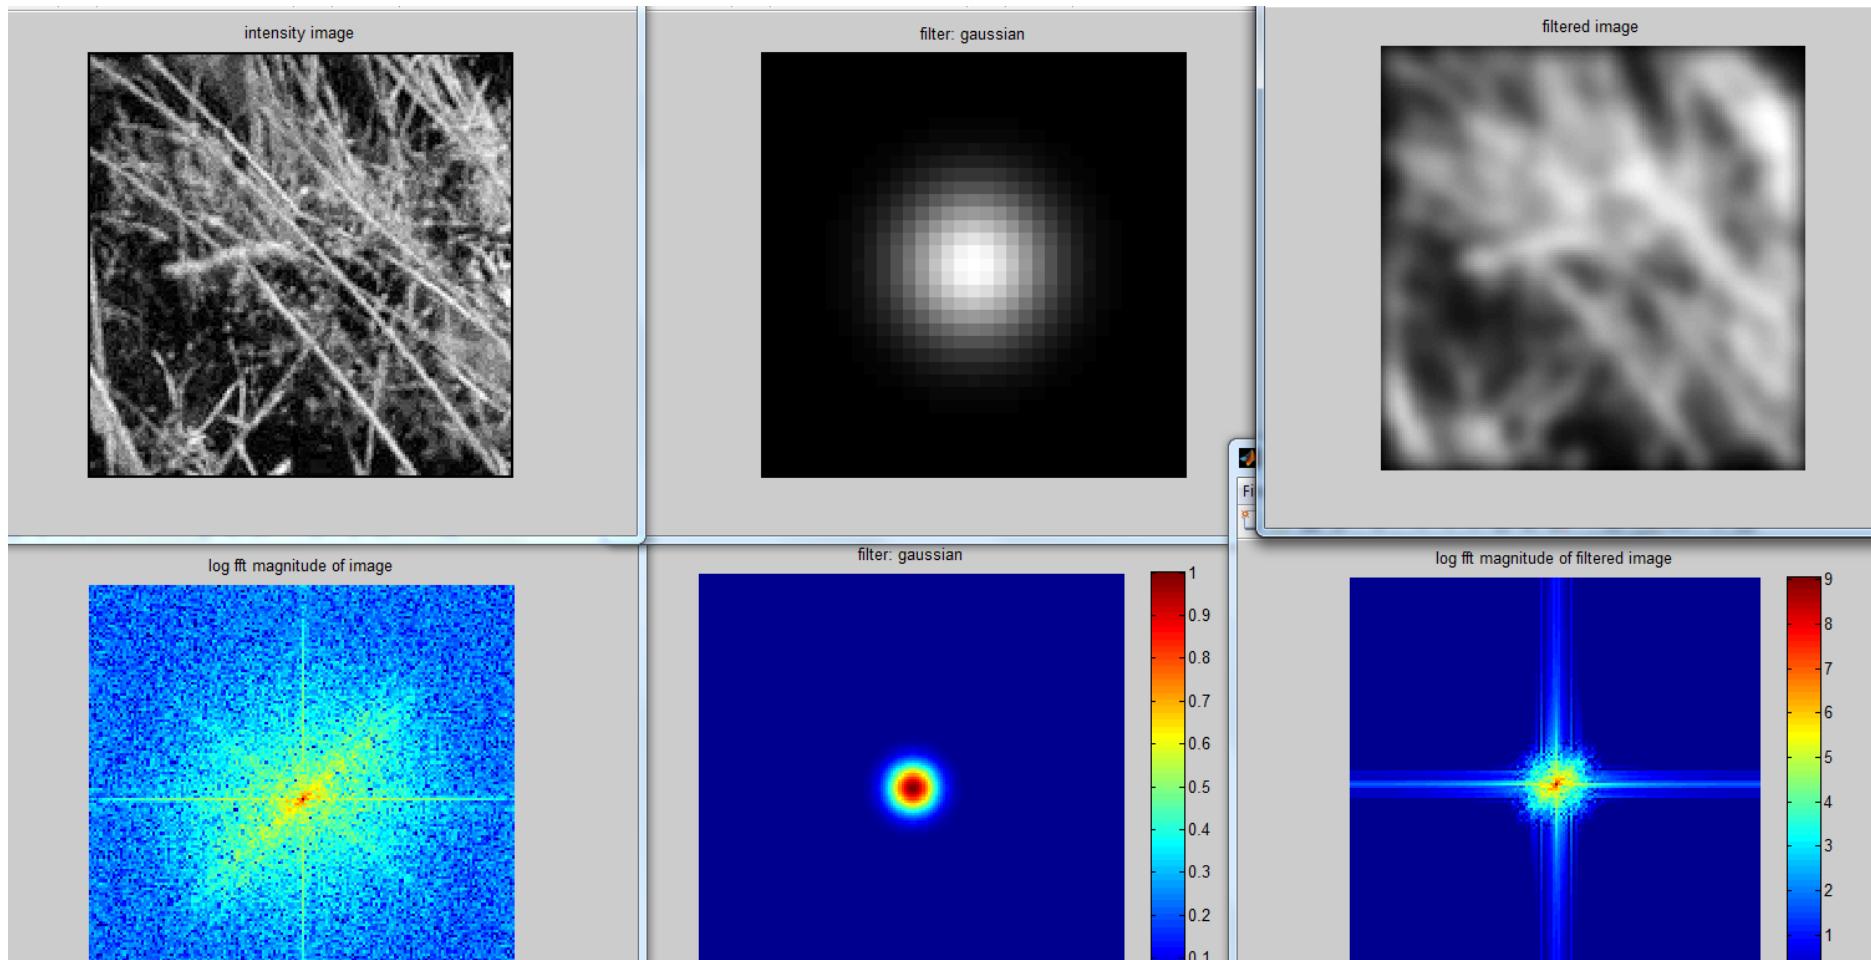


# Is convolution invertible?

- If convolution is just multiplication in the Fourier domain, isn't deconvolution just division?
- Sometimes, it clearly is invertible (e.g. a convolution with an identity filter)
- In one case, it clearly isn't invertible (e.g. convolution with an all zero filter)
- What about for common filters like a Gaussian?

# But you can't invert multiplication by 0

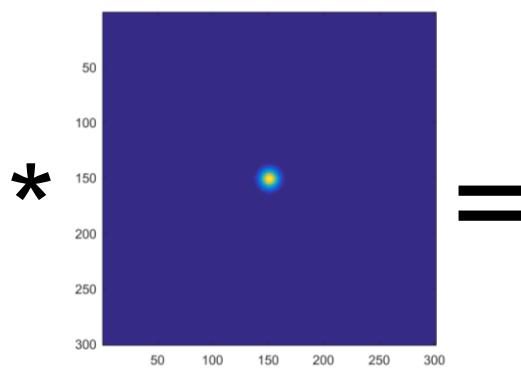
- But it's not quite zero, is it...



# Let's experiment on Novak



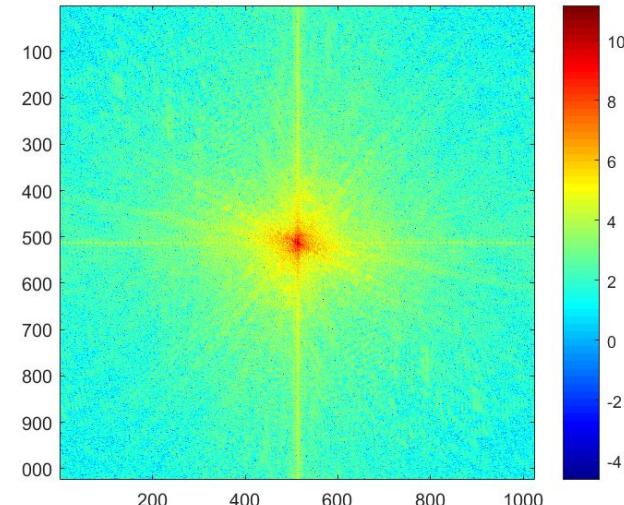
# Convolution



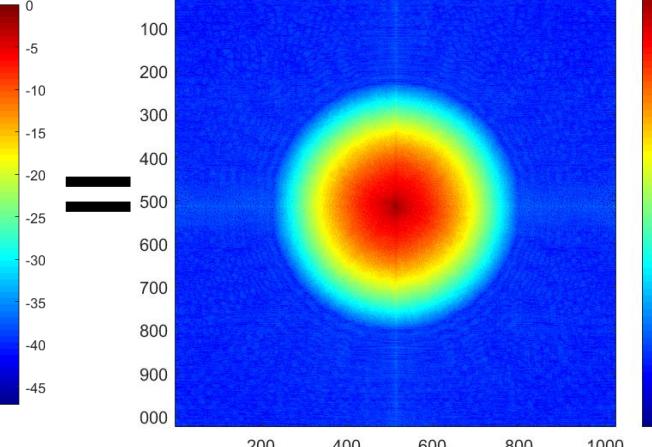
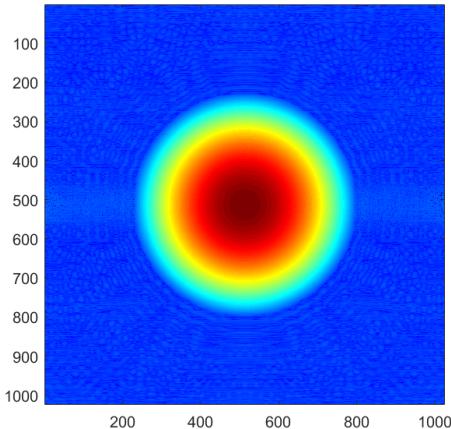
FFT

FFT

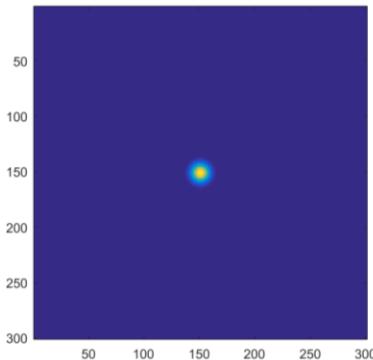
iFFT



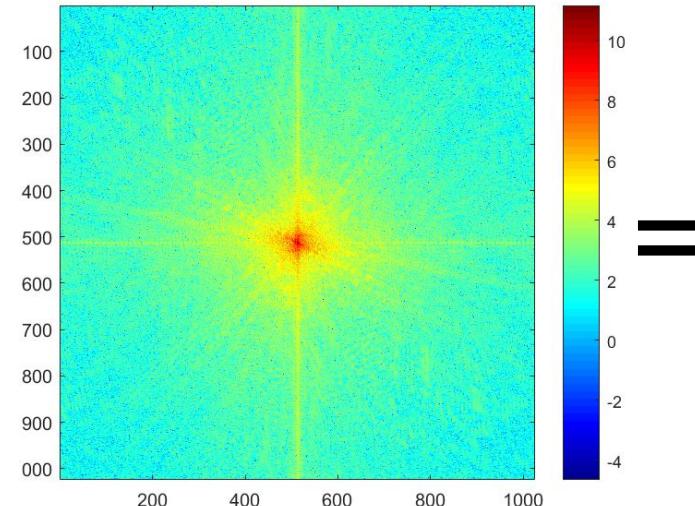
\*



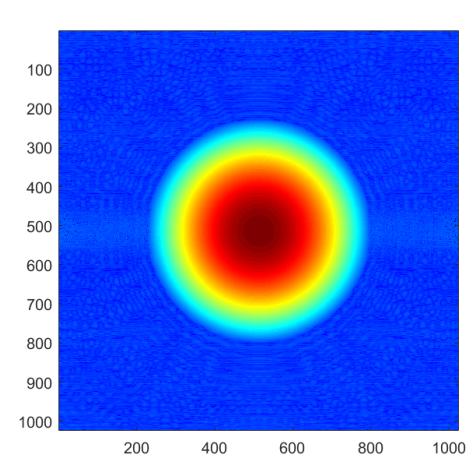
# Deconvolution?



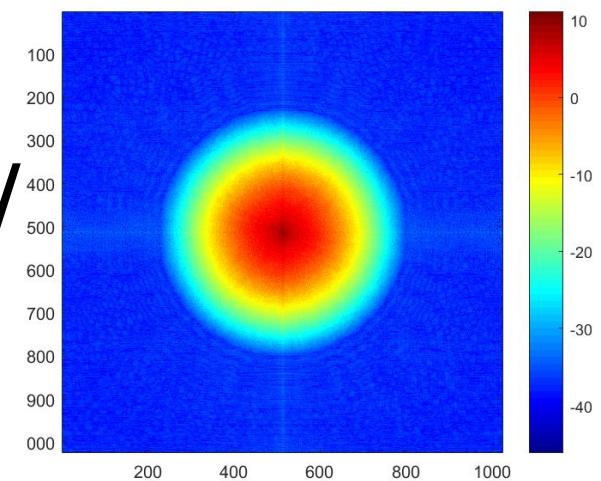
iFFT



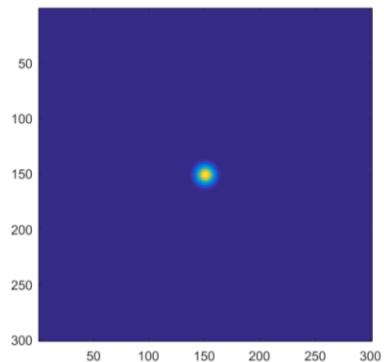
FFT



FFT



# But under more realistic conditions



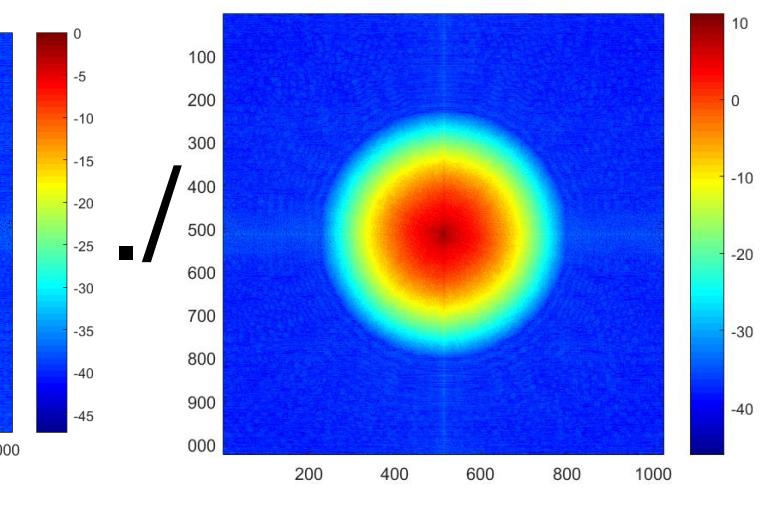
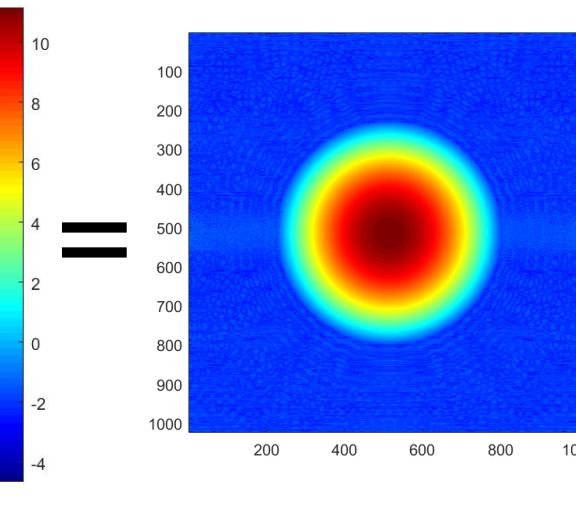
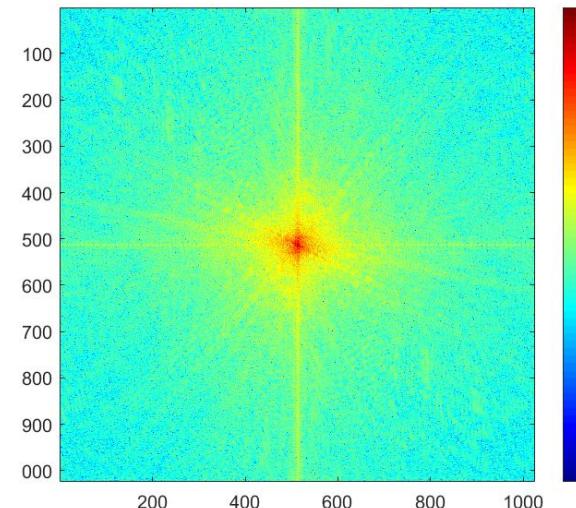
Random noise, .000001 magnitude



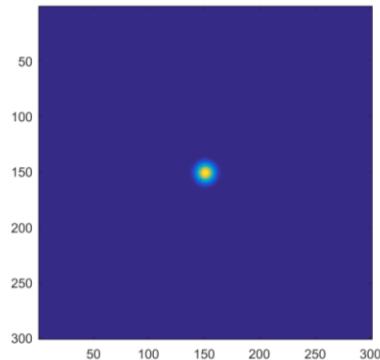
iFFT

FFT

FFT



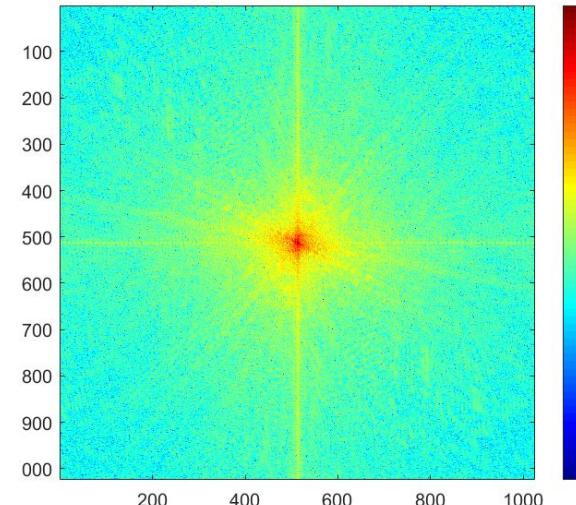
# But under more realistic conditions



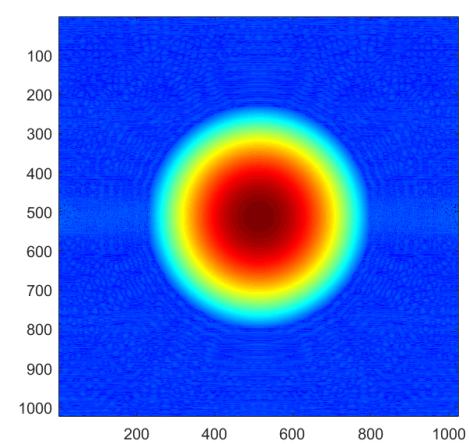
Random noise, .0001 magnitude



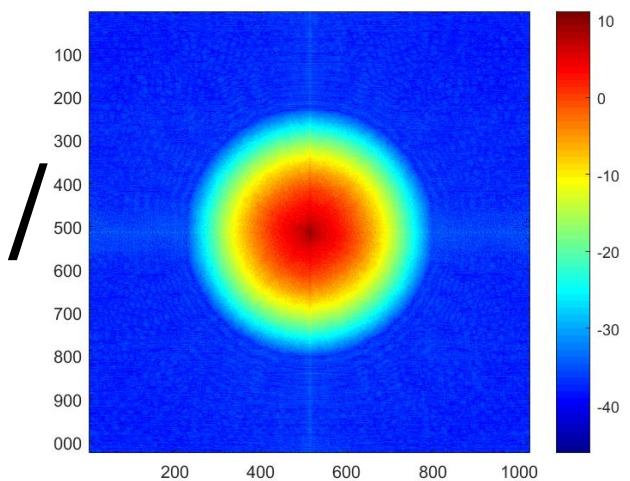
iFFT



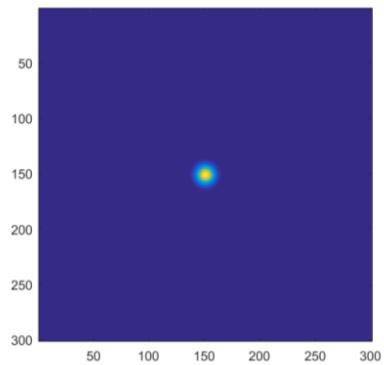
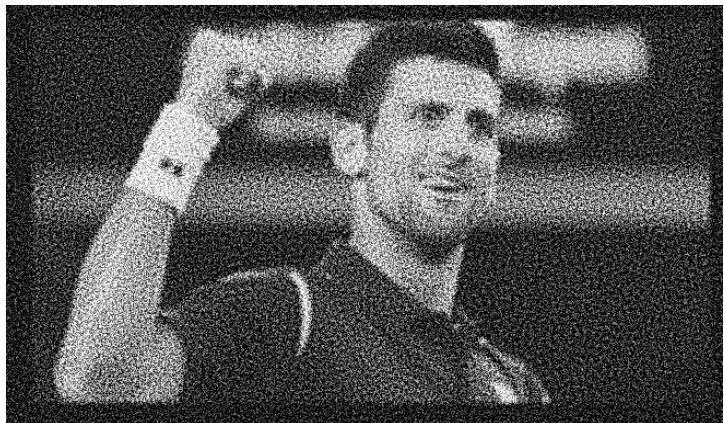
FFT



FFT



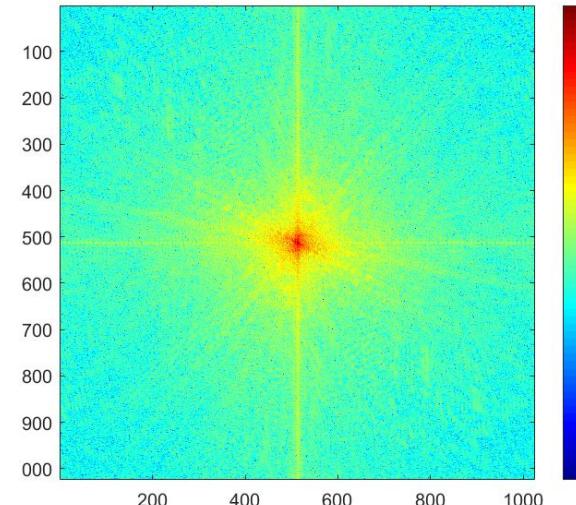
# But under more realistic conditions



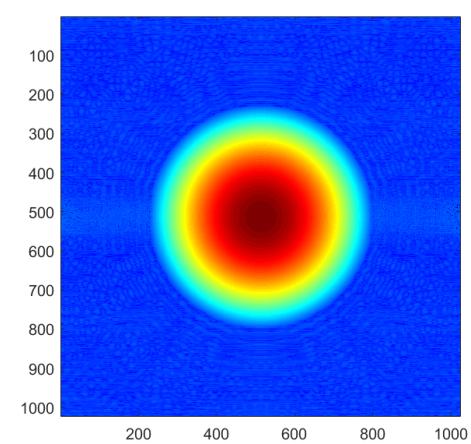
Random noise, .001 magnitude



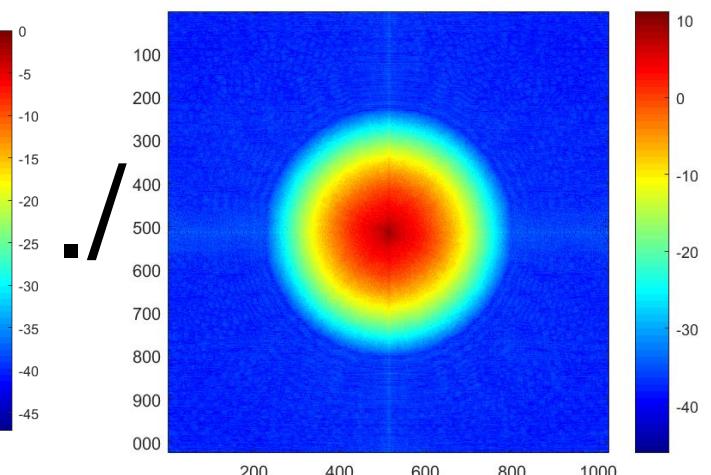
iFFT



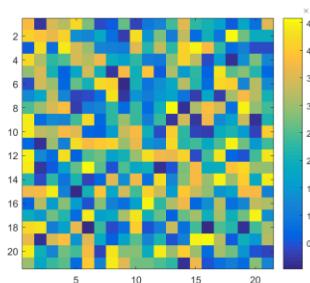
FFT



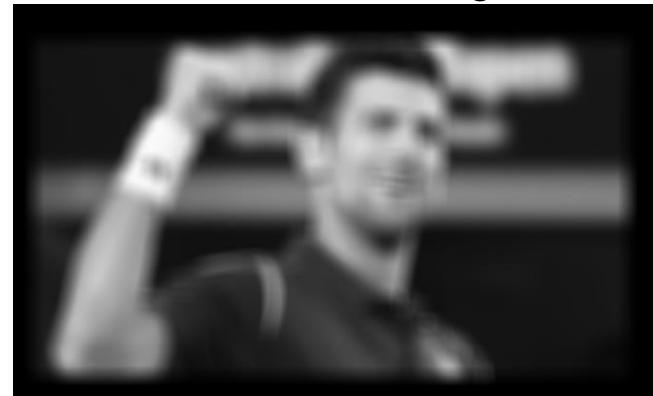
FFT



# With a random filter...



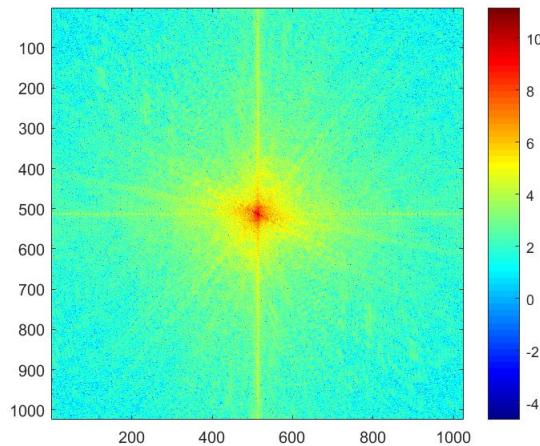
Random noise, .001 magnitude



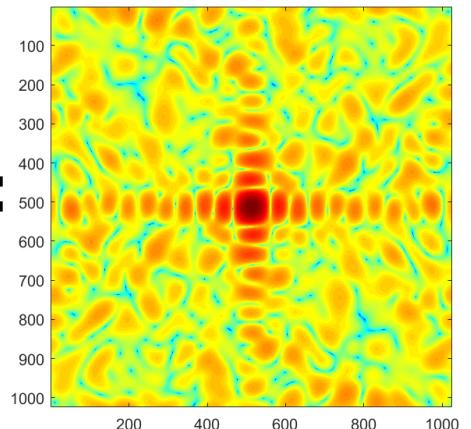
iFFT

FFT

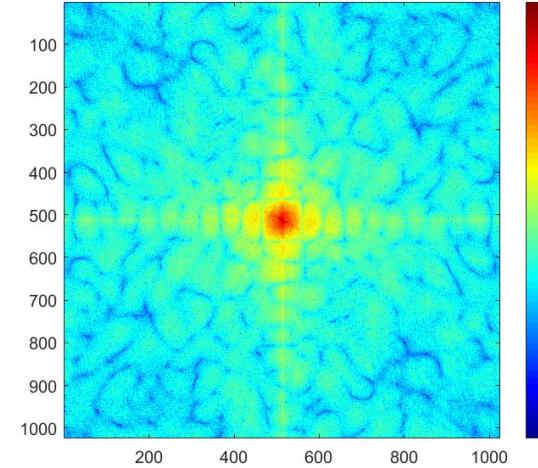
FFT



=



/



# Deconvolution is hard

- Active research area.
- Even if you know the filter (non-blind deconvolution), it is still very hard and requires strong *regularization*.
- If you don't know the filter (blind deconvolution) it is harder still.

# Interest Points and Corners

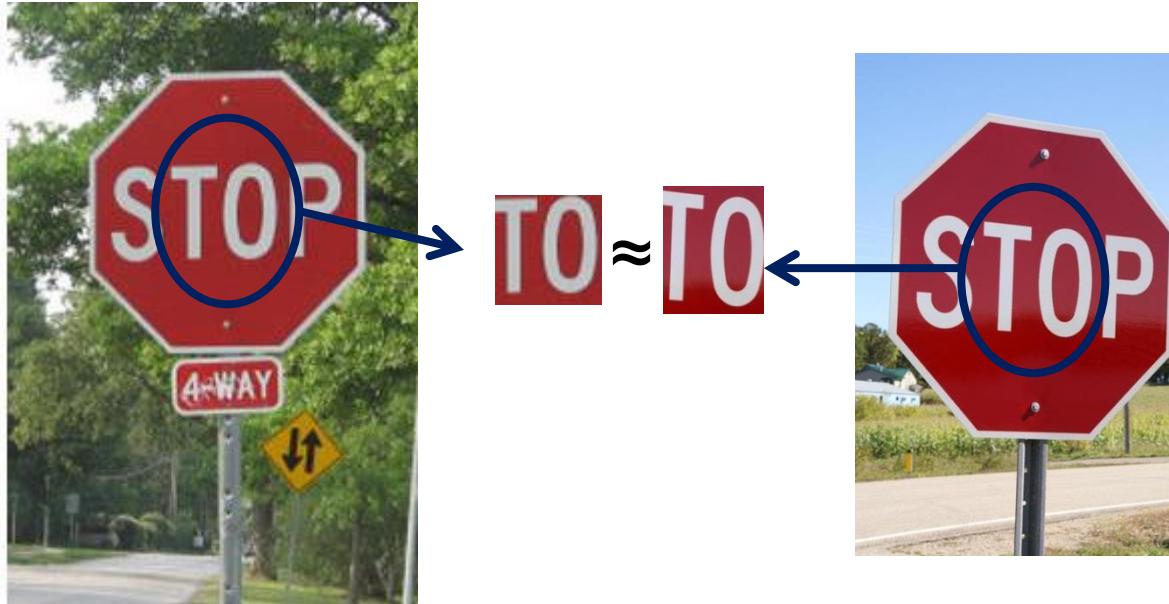
Read Szeliski 4.1

Computer Vision

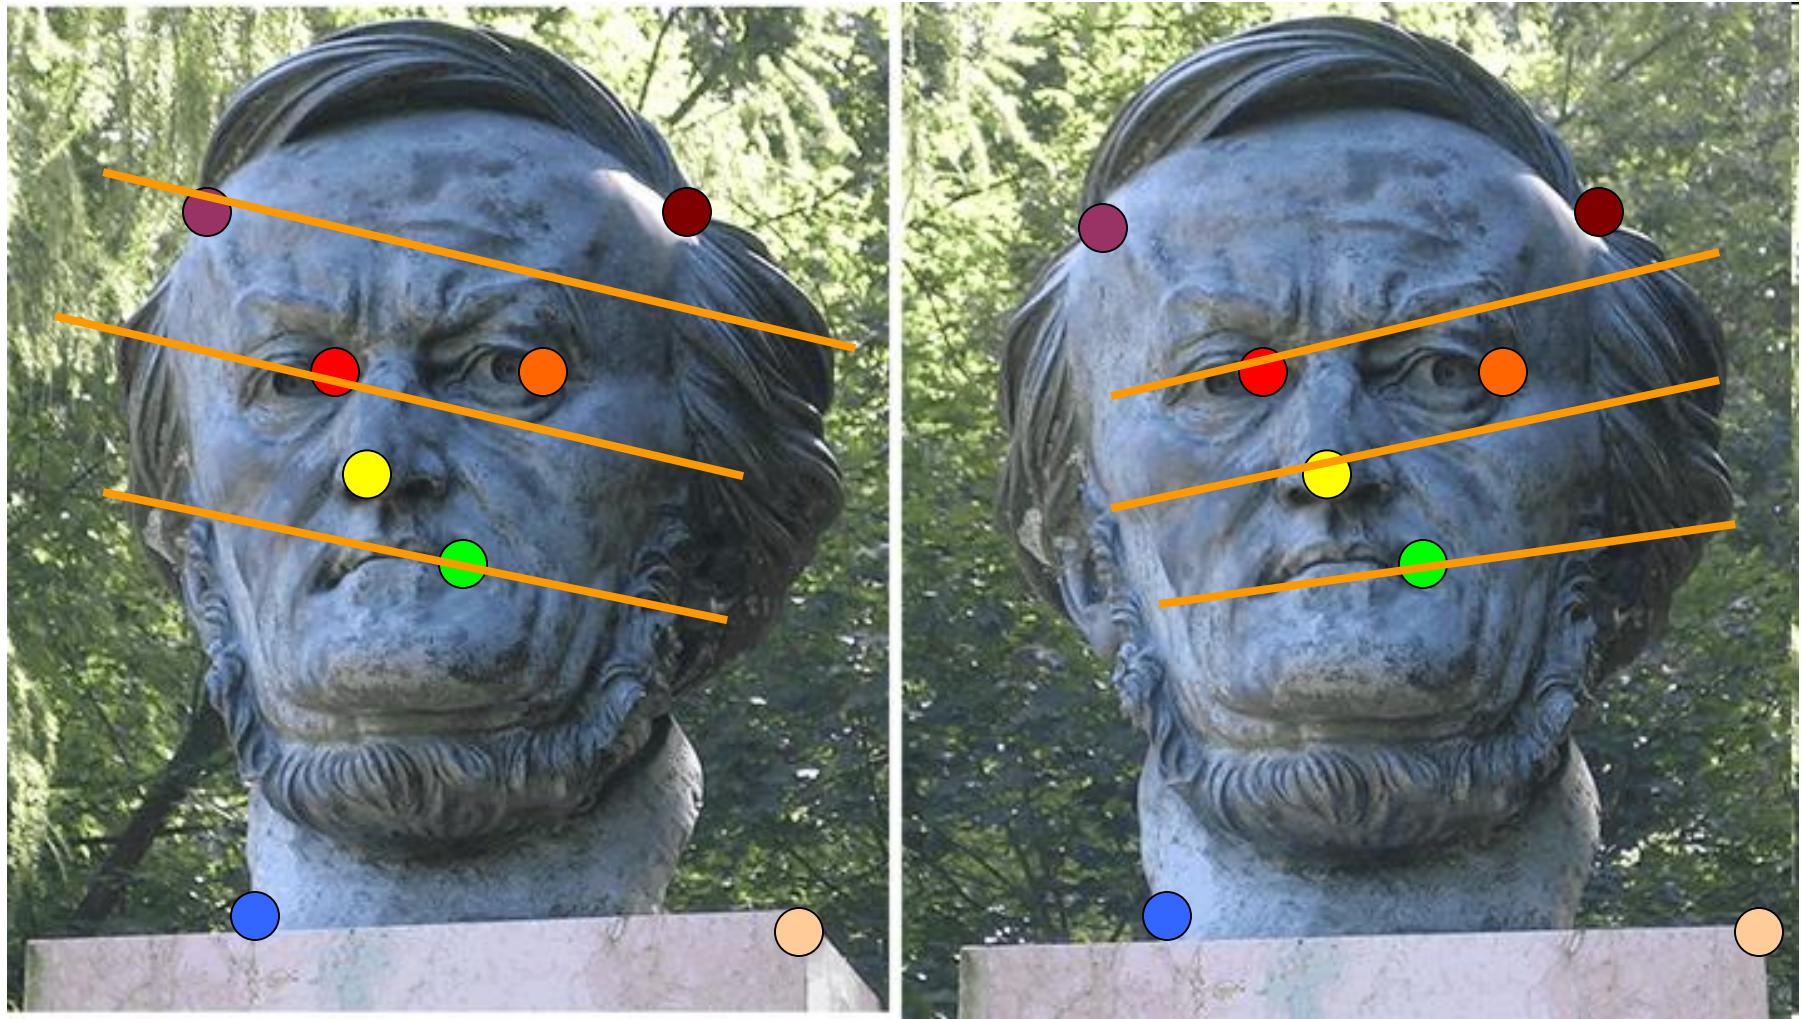
James Hays

# Correspondence across views

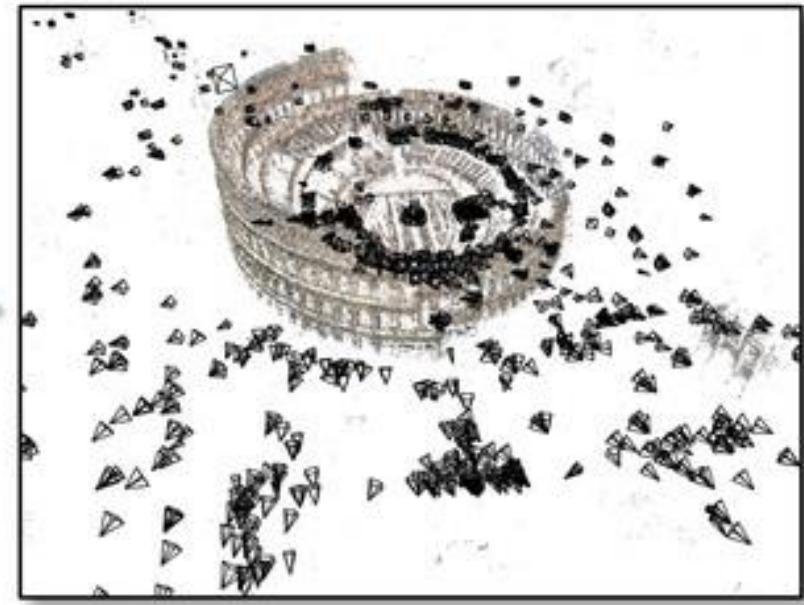
- Correspondence: matching points, patches, edges, or regions across images



# Example: estimating “fundamental matrix” that corresponds two views



# Example: structure from motion



# Applications

- Feature points are used for:
  - Image alignment
  - 3D reconstruction
  - Motion tracking
  - Robot navigation
  - Indexing and database retrieval
  - Object recognition

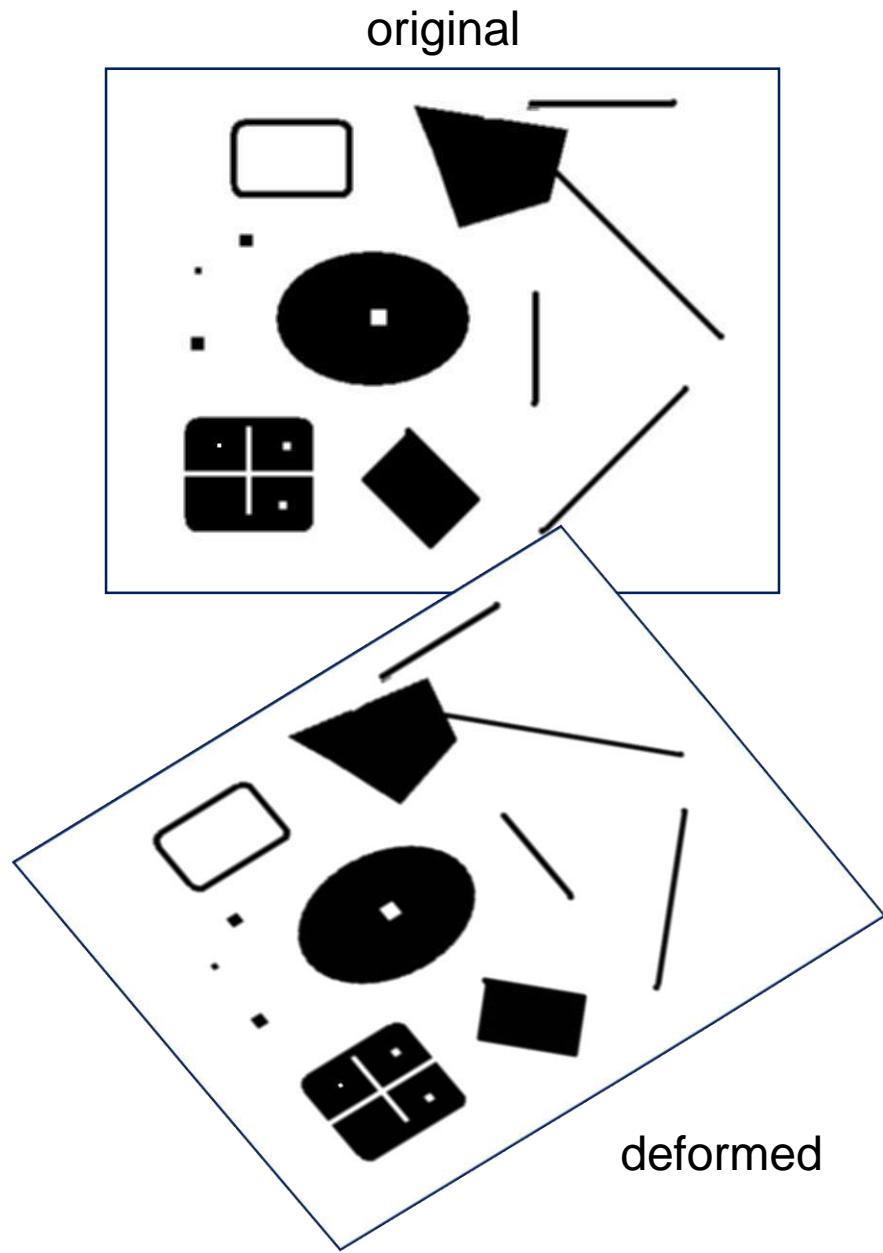


# Project 2: interest points and local features

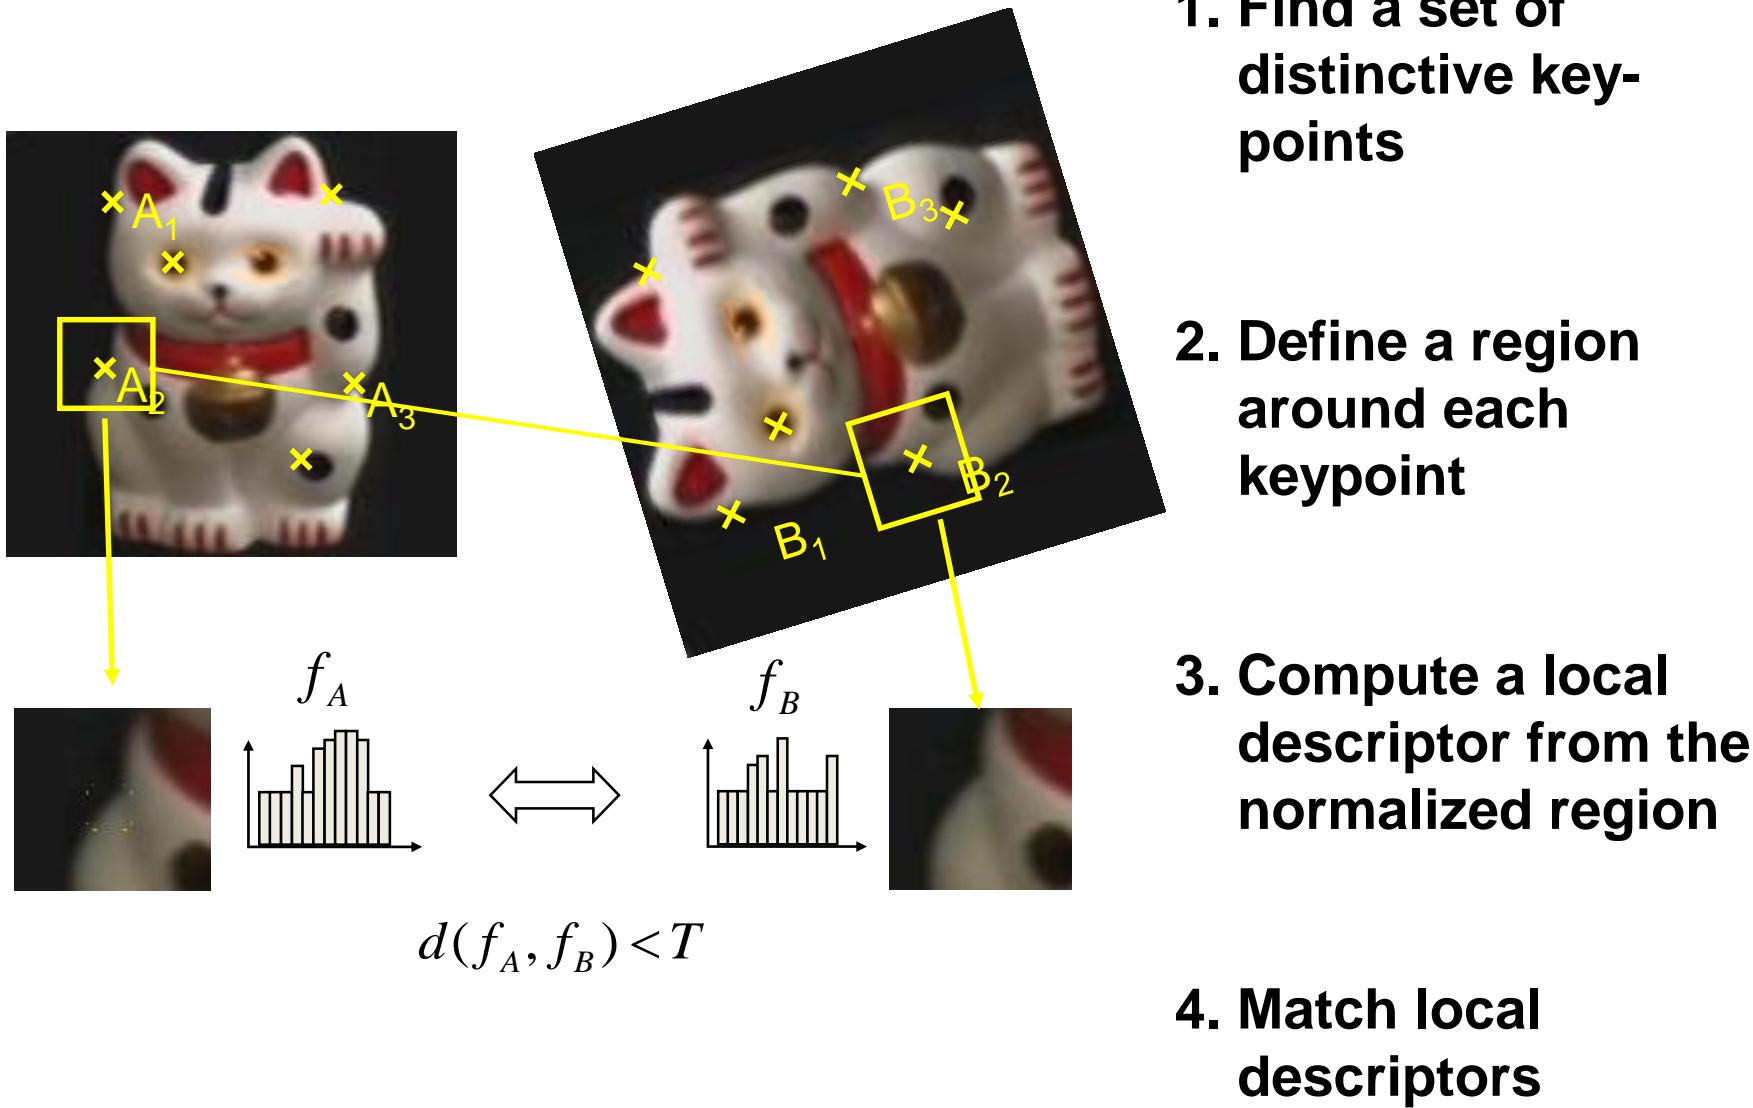
- Note: “interest points” = “keypoints”, also sometimes called “features”

# This class: interest points

- Suppose you have to click on some point, go away and come back after I deform the image, and click on the same points again.
  - Which points would you choose?



# Overview of Keypoint Matching



# Goals for Keypoints

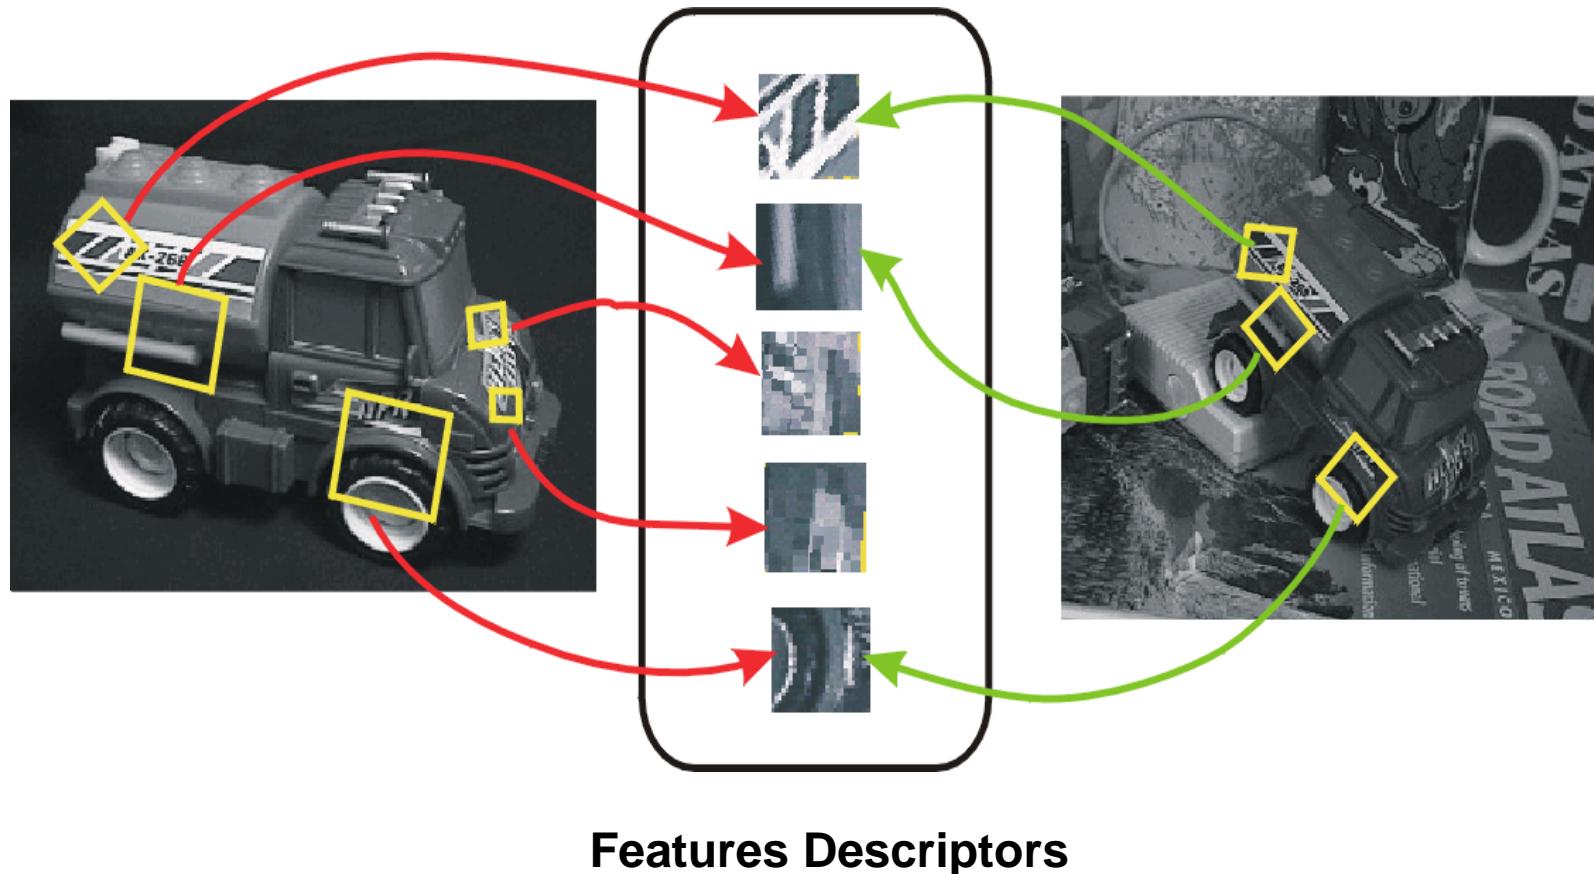


Detect points that are *repeatable* and *distinctive*

# Invariant Local Features

---

Image content is transformed into local feature coordinates that are invariant to translation, rotation, scale, and other imaging parameters



# Why extract features?

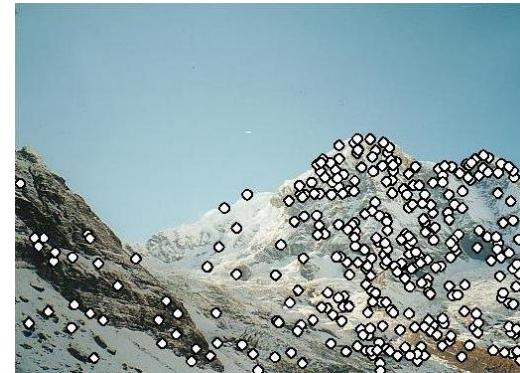
---

- Motivation: panorama stitching
  - We have two images – how do we combine them?



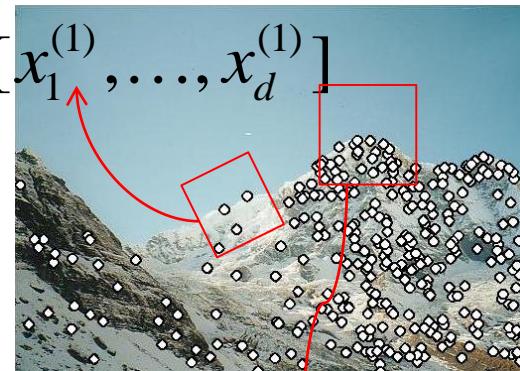
# Local features: main components

1) Detection: Identify the interest points



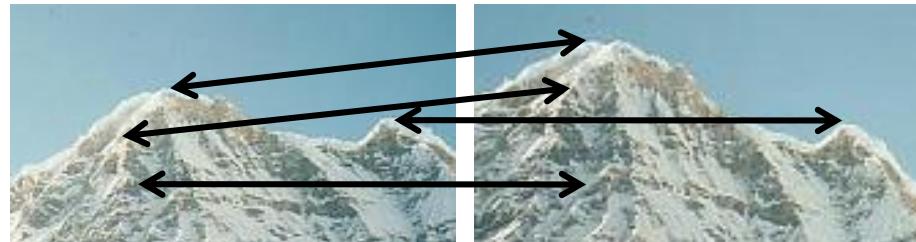
2) Description: Extract vector feature descriptor surrounding each interest point.

$$\mathbf{x}_1 = [x_1^{(1)}, \dots, x_d^{(1)}]$$



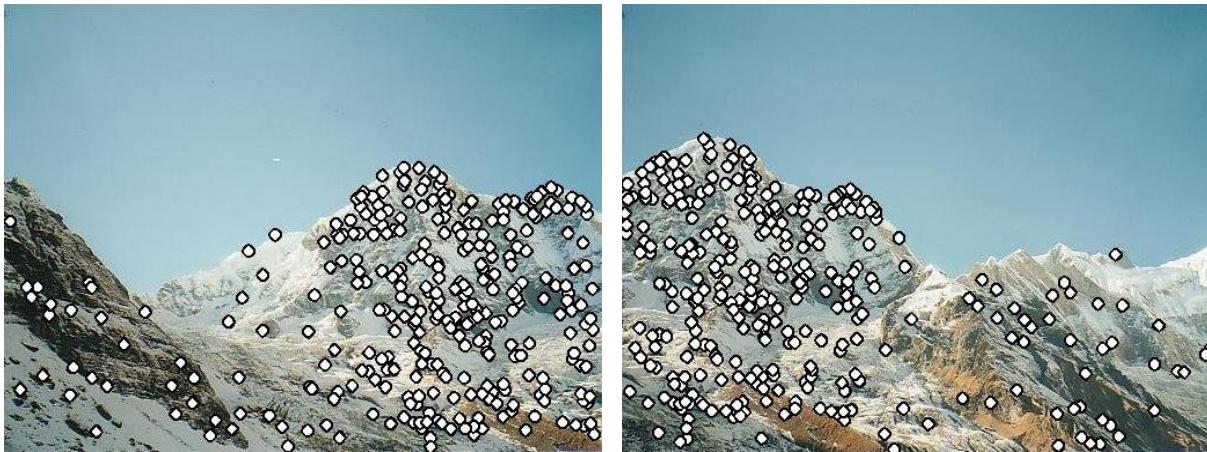
3) Matching: Determine correspondence between descriptors in two views

$$\mathbf{x}_2 = [x_1^{(2)}, \dots, x_d^{(2)}]$$



# Characteristics of good features

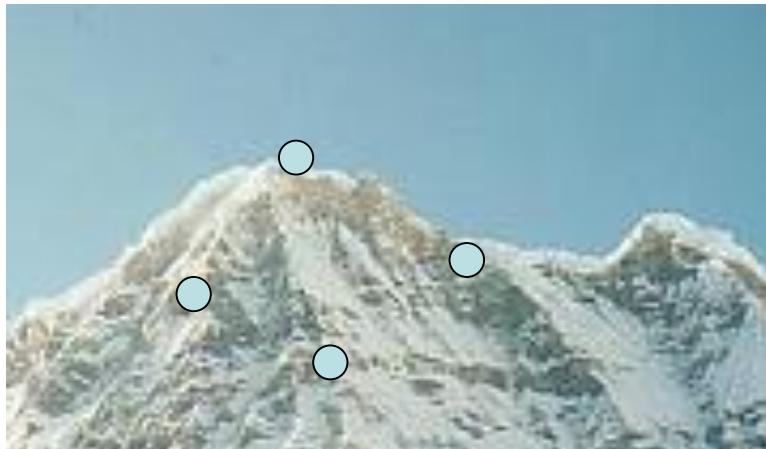
---



- **Repeatability**
  - The same feature can be found in several images despite geometric and photometric transformations
- **Saliency**
  - Each feature is distinctive
- **Compactness and efficiency**
  - Many fewer features than image pixels
- **Locality**
  - A feature occupies a relatively small area of the image; robust to clutter and occlusion

# Goal: interest operator repeatability

- We want to detect (at least some of) the same points in both images.

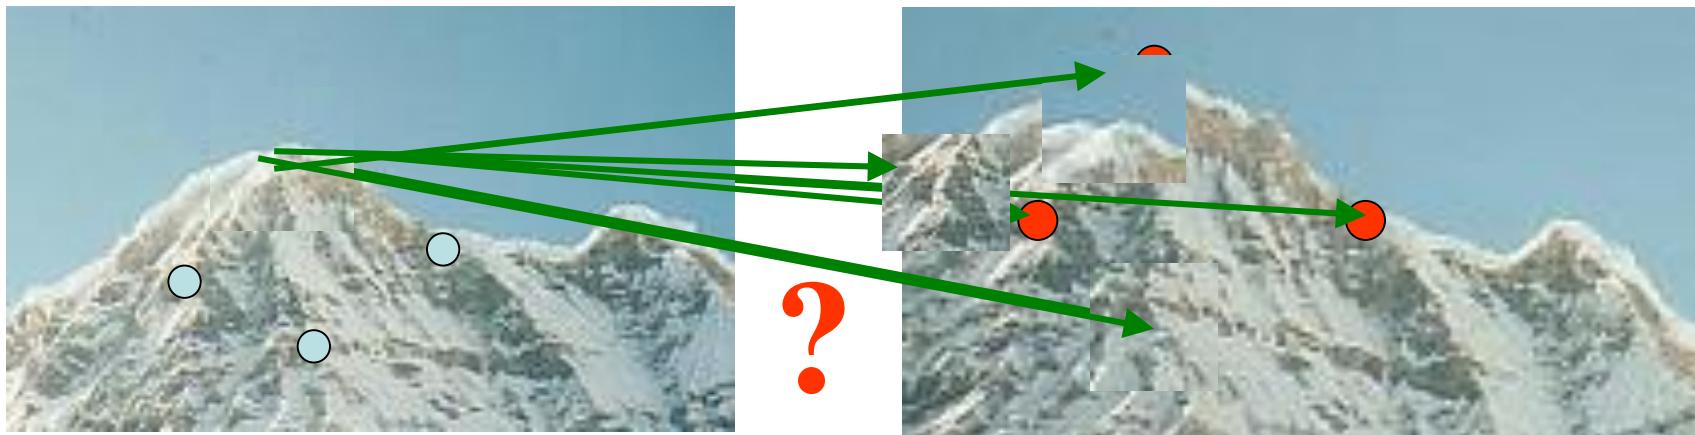


No chance to find true matches!

- Yet we have to be able to run the detection procedure *independently* per image.

# Goal: descriptor distinctiveness

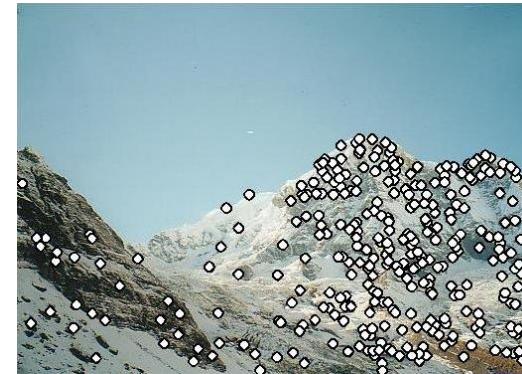
- We want to be able to reliably determine which point goes with which.



- Must provide some invariance to geometric and photometric differences between the two views.

# Local features: main components

1) Detection: Identify the interest points



2) Description: Extract vector feature descriptor surrounding each interest point.

3) Matching: Determine correspondence between descriptors in two views

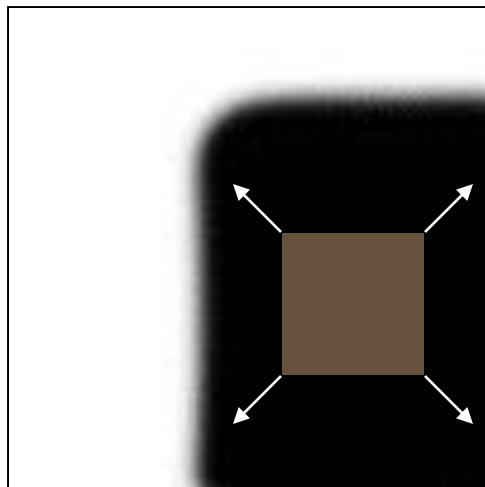
# Many Existing Detectors Available

Hessian & Harris	[Beaudet '78], [Harris '88]
Laplacian, DoG	[Lindeberg '98], [Lowe 1999]
Harris-/Hessian-Laplace	[Mikolajczyk & Schmid '01]
Harris-/Hessian-Affine	[Mikolajczyk & Schmid '04]
EBR and IBR	[Tuytelaars & Van Gool '04]
MSER	[Matas '02]
Salient Regions	[Kadir & Brady '01]
Others...	

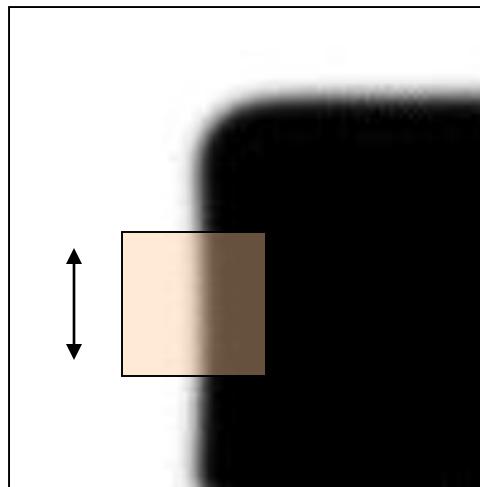
# Corner Detection: Basic Idea

---

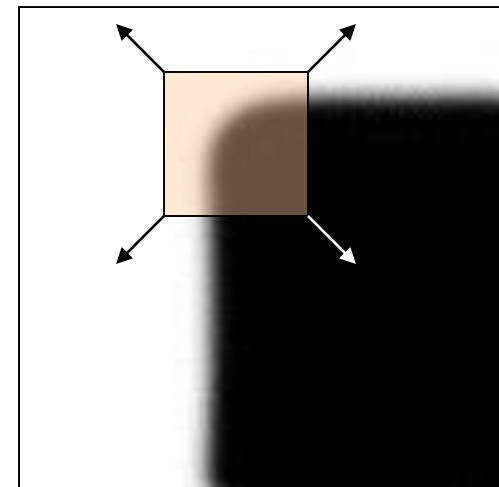
- We should easily recognize the point by looking through a small window
- Shifting a window in *any direction* should give a *large change* in intensity



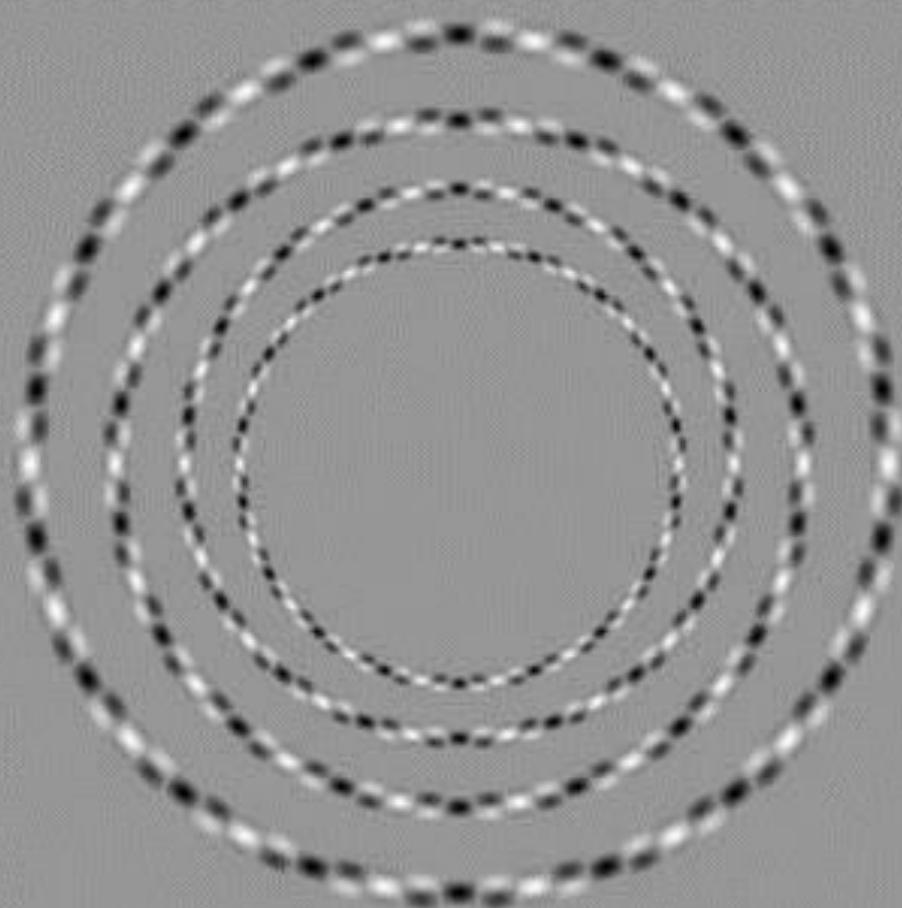
“flat” region:  
no change in  
all directions

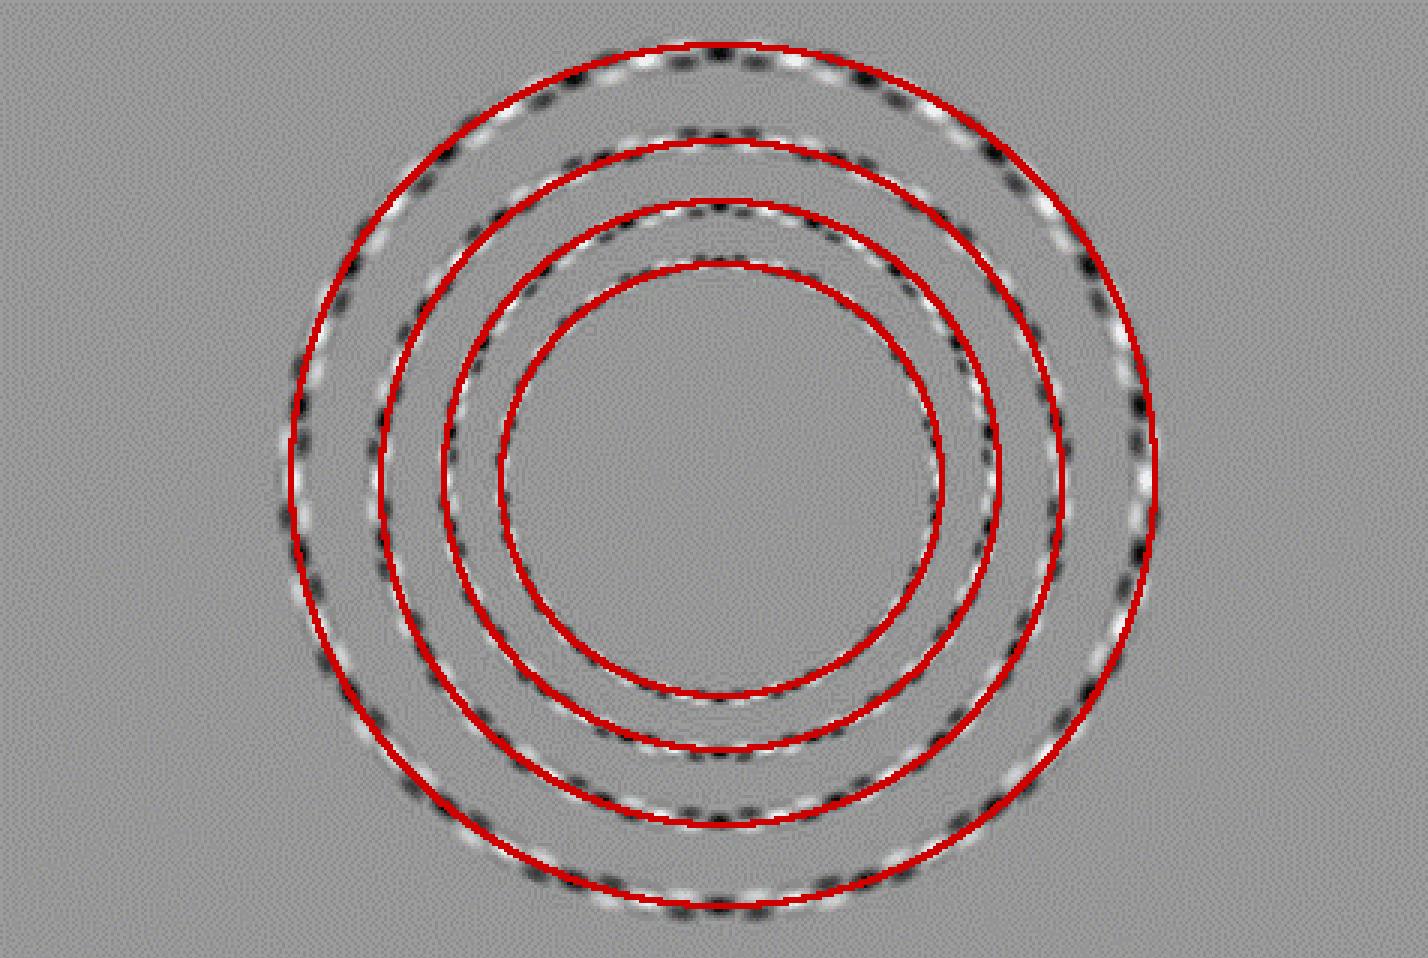


“edge”:  
no change  
along the edge  
direction



“corner”:  
significant  
change in all  
directions



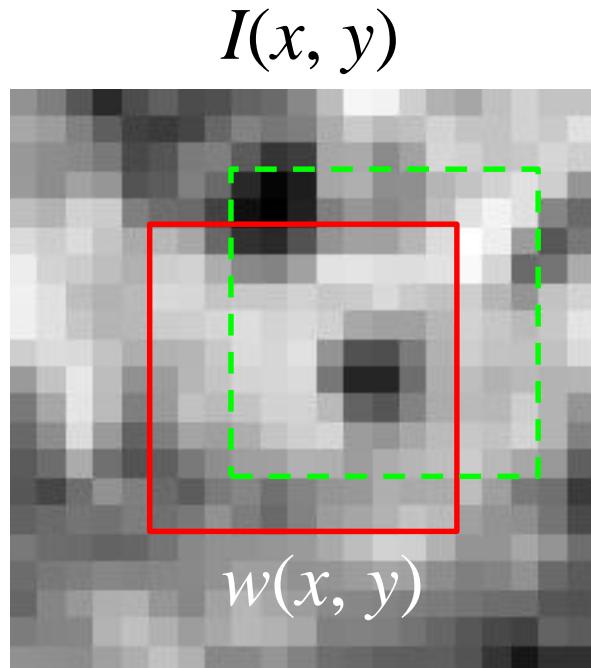


# Corner Detection: Mathematics

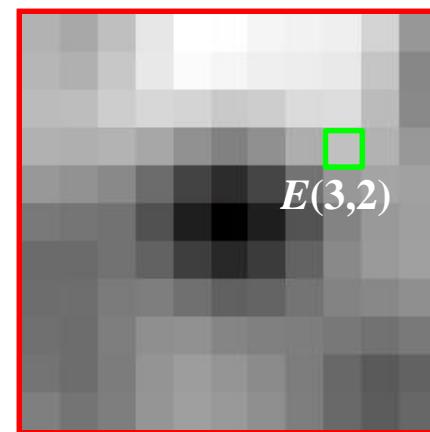
---

Change in appearance of window  $w(x,y)$   
for the shift  $[u,v]$ :

$$E(u, v) = \sum_{x, y} w(x, y) [I(x + u, y + v) - I(x, y)]^2$$



$$E(u, v)$$



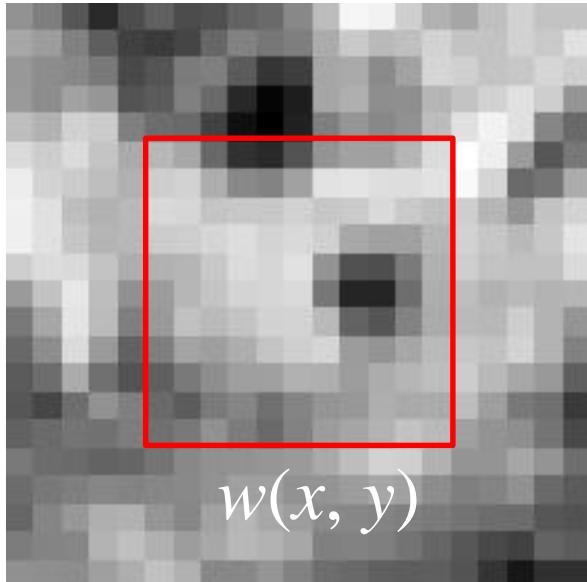
# Corner Detection: Mathematics

---

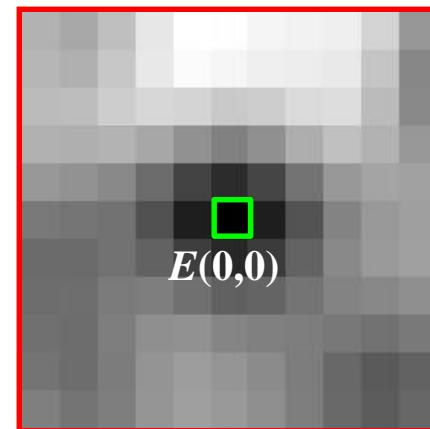
Change in appearance of window  $w(x,y)$   
for the shift  $[u,v]$ :

$$E(u, v) = \sum_{x, y} w(x, y) [I(x + u, y + v) - I(x, y)]^2$$

$I(x, y)$



$E(u, v)$



# Corner Detection: Mathematics

Change in appearance of window  $w(x,y)$   
for the shift  $[u,v]$ :

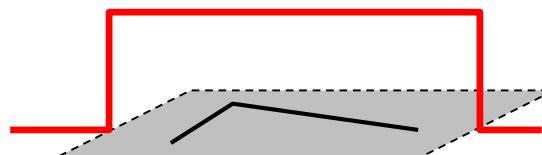
$$E(u,v) = \sum_{x,y} w(x,y) [I(x+u, y+v) - I(x, y)]^2$$

Window  
function

Shifted  
intensity

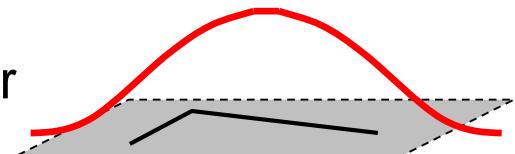
Intensity

Window function  $w(x,y) =$



1 in window, 0 outside

or



Gaussian

# Corner Detection: Mathematics

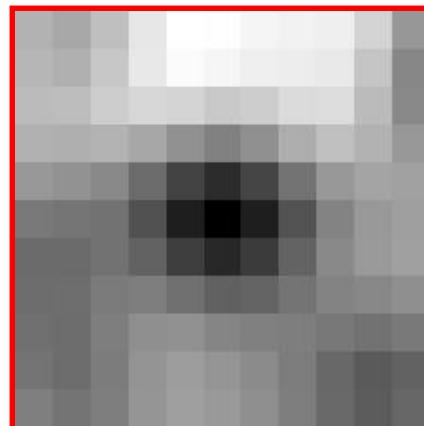
---

Change in appearance of window  $w(x,y)$   
for the shift  $[u,v]$ :

$$E(u, v) = \sum_{x, y} w(x, y) [I(x + u, y + v) - I(x, y)]^2$$

We want to find out how this function behaves for small shifts

$$E(u, v)$$



# Corner Detection: Mathematics

---

Change in appearance of window  $w(x,y)$   
for the shift  $[u,v]$ :

$$E(u,v) = \sum_{x,y} w(x,y) [I(x+u, y+v) - I(x, y)]^2$$

We want to find out how this function behaves for small shifts

But this is very slow to compute naively.

$O(\text{window\_width}^2 * \text{shift\_range}^2 * \text{image\_width}^2)$

$O(11^2 * 11^2 * 600^2) = 5.2$  billion of these  
14.6 thousand per pixel in your image

# Corner Detection: Mathematics

---

Change in appearance of window  $w(x,y)$   
for the shift  $[u,v]$ :

$$E(u,v) = \sum_{x,y} w(x,y) [I(x+u, y+v) - I(x, y)]^2$$

We want to find out how this function behaves for small shifts

Recall Taylor series expansion. A function  $f$  can be approximated around point  $a$  as

$$f(a) + \frac{f'(a)}{1!}(x-a) + \frac{f''(a)}{2!}(x-a)^2 + \frac{f'''(a)}{3!}(x-a)^3 + \dots$$

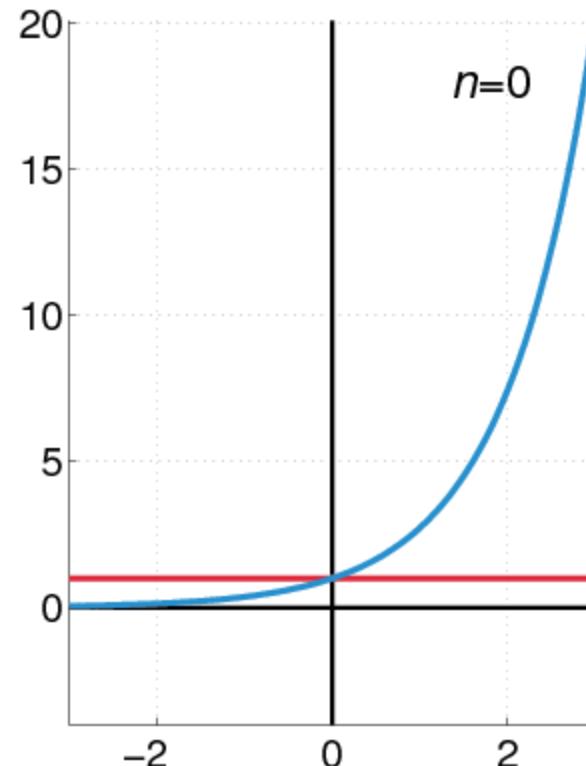
# Recall: Taylor series expansion

---

A function  $f$  can be approximated as

$$f(a) + \frac{f'(a)}{1!}(x - a) + \frac{f''(a)}{2!}(x - a)^2 + \frac{f'''(a)}{3!}(x - a)^3 + \dots$$

Approximation of  
 $f(x) = e^x$   
centered at  $f(0)$



# Corner Detection: Mathematics

---

Change in appearance of window  $w(x,y)$   
for the shift  $[u,v]$ :

$$E(u,v) = \sum_{x,y} w(x,y) [I(x+u, y+v) - I(x, y)]^2$$

We want to find out how this function behaves for small shifts

Local quadratic approximation of  $E(u,v)$  in the neighborhood of  $(0,0)$  is given by the *second-order Taylor expansion*:

$$E(u,v) \approx E(0,0) + [u \ v] \begin{bmatrix} E_u(0,0) \\ E_v(0,0) \end{bmatrix} + \frac{1}{2} [u \ v] \begin{bmatrix} E_{uu}(0,0) & E_{uv}(0,0) \\ E_{uv}(0,0) & E_{vv}(0,0) \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix}$$

# Corner Detection: Mathematics

---

Local quadratic approximation of  $E(u,v)$  in the neighborhood of  $(0,0)$  is given by the *second-order Taylor expansion*:

$$E(u,v) \approx E(0,0) + [u \ v] \begin{bmatrix} E_u(0,0) \\ E_v(0,0) \end{bmatrix} + \frac{1}{2} [u \ v] \begin{bmatrix} E_{uu}(0,0) & E_{uv}(0,0) \\ E_{uv}(0,0) & E_{vv}(0,0) \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix}$$

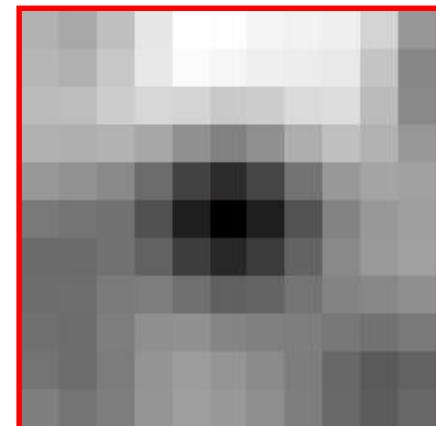


Always 0



First  
derivative  
is 0

$E(u, v)$



# Corner Detection: Mathematics

---

The quadratic approximation simplifies to

$$E(u, v) \approx [u \ v] M \begin{bmatrix} u \\ v \end{bmatrix}$$

where  $M$  is a *second moment matrix* computed from image derivatives:

$$M = \sum_{x,y} w(x, y) \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix}$$

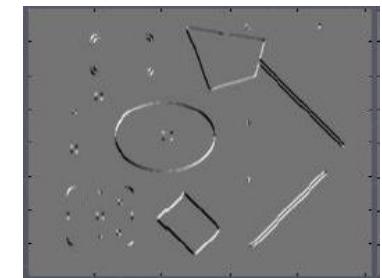
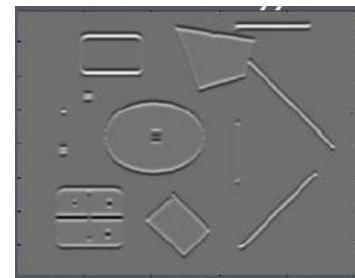
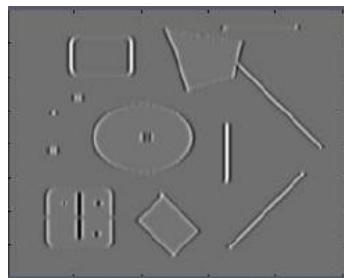
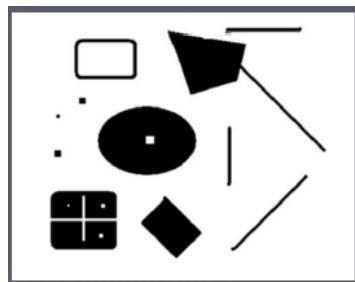
$$M = \begin{bmatrix} \sum I_x I_x & \sum I_x I_y \\ \sum I_x I_y & \sum I_y I_y \end{bmatrix} = \sum \begin{bmatrix} I_x \\ I_y \end{bmatrix} [I_x \ I_y] = \sum \nabla I (\nabla I)^T$$

# Corners as distinctive interest points

---

$$M = \sum w(x, y) \begin{bmatrix} I_x I_x & I_x I_y \\ I_x I_y & I_y I_y \end{bmatrix}$$

2 x 2 matrix of image derivatives (averaged in neighborhood of a point).



Notation:

$$I_x \Leftrightarrow \frac{\partial I}{\partial x}$$

$$I_y \Leftrightarrow \frac{\partial I}{\partial y}$$

$$I_x I_y \Leftrightarrow \frac{\partial I}{\partial x} \frac{\partial I}{\partial y}$$

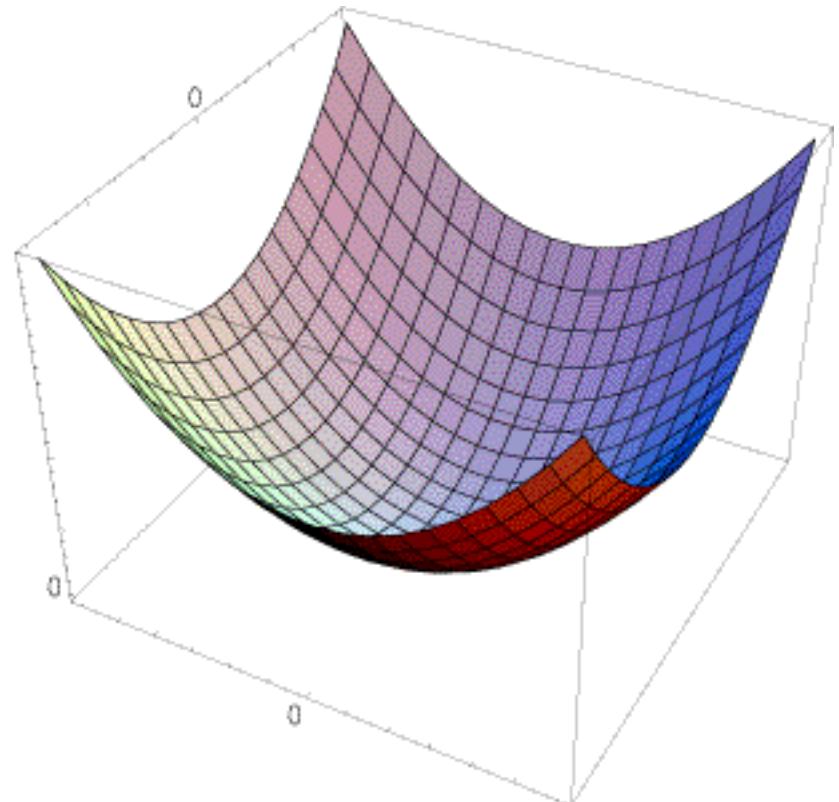
# Interpreting the second moment matrix

---

The surface  $E(u,v)$  is locally approximated by a quadratic form. Let's try to understand its shape.

$$E(u,v) \approx [u \ v] M \begin{bmatrix} u \\ v \end{bmatrix}$$

$$M = \sum_{x,y} w(x,y) \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix}$$

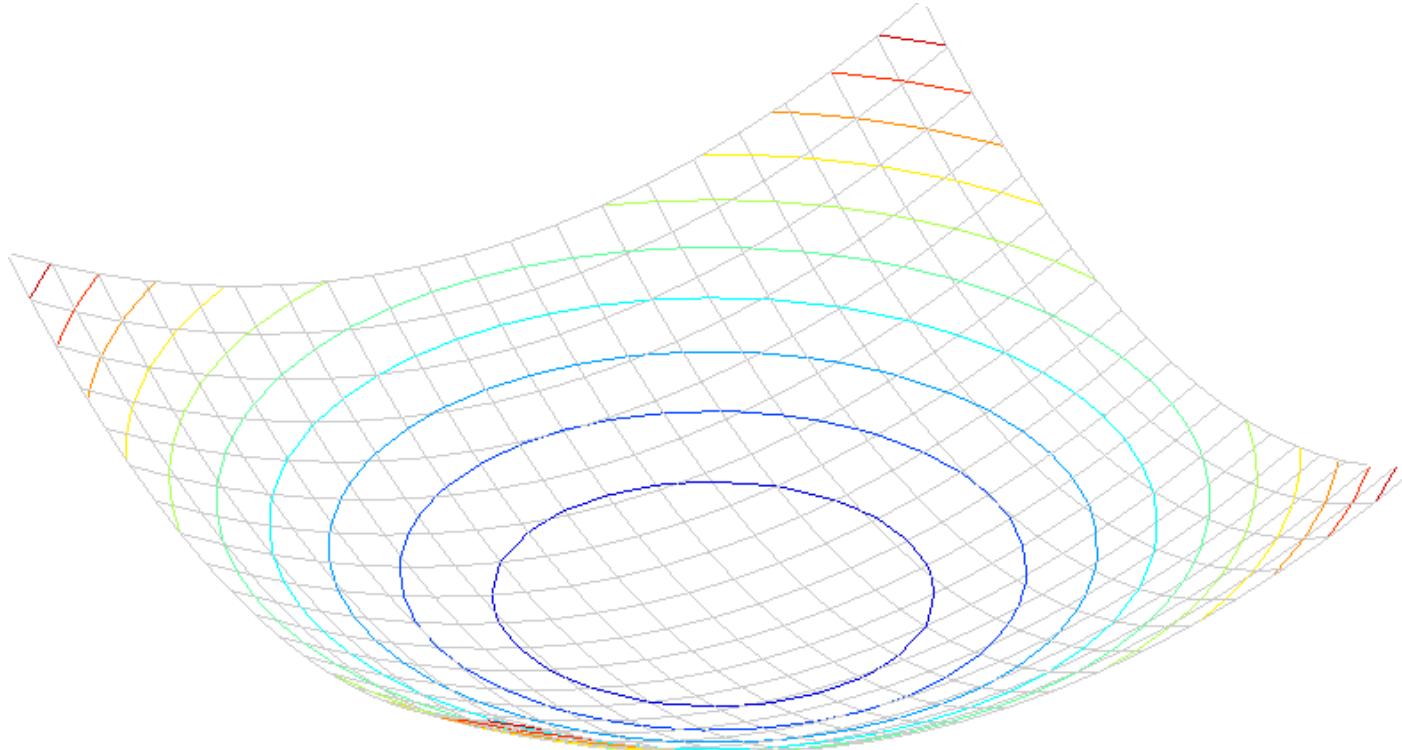


# Interpreting the second moment matrix

---

Consider a horizontal “slice” of  $E(u, v)$ :  $[u \ v] M \begin{bmatrix} u \\ v \end{bmatrix} = \text{const}$

This is the equation of an ellipse.



# Interpreting the second moment matrix

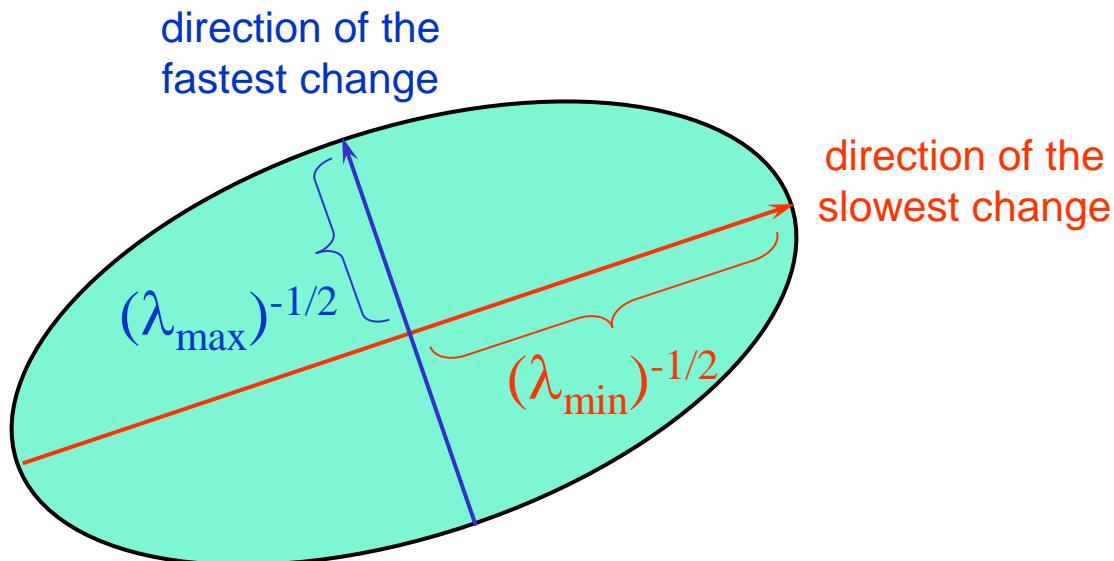
---

Consider a horizontal “slice” of  $E(u, v)$ :  $[u \ v] M \begin{bmatrix} u \\ v \end{bmatrix} = \text{const}$

This is the equation of an ellipse.

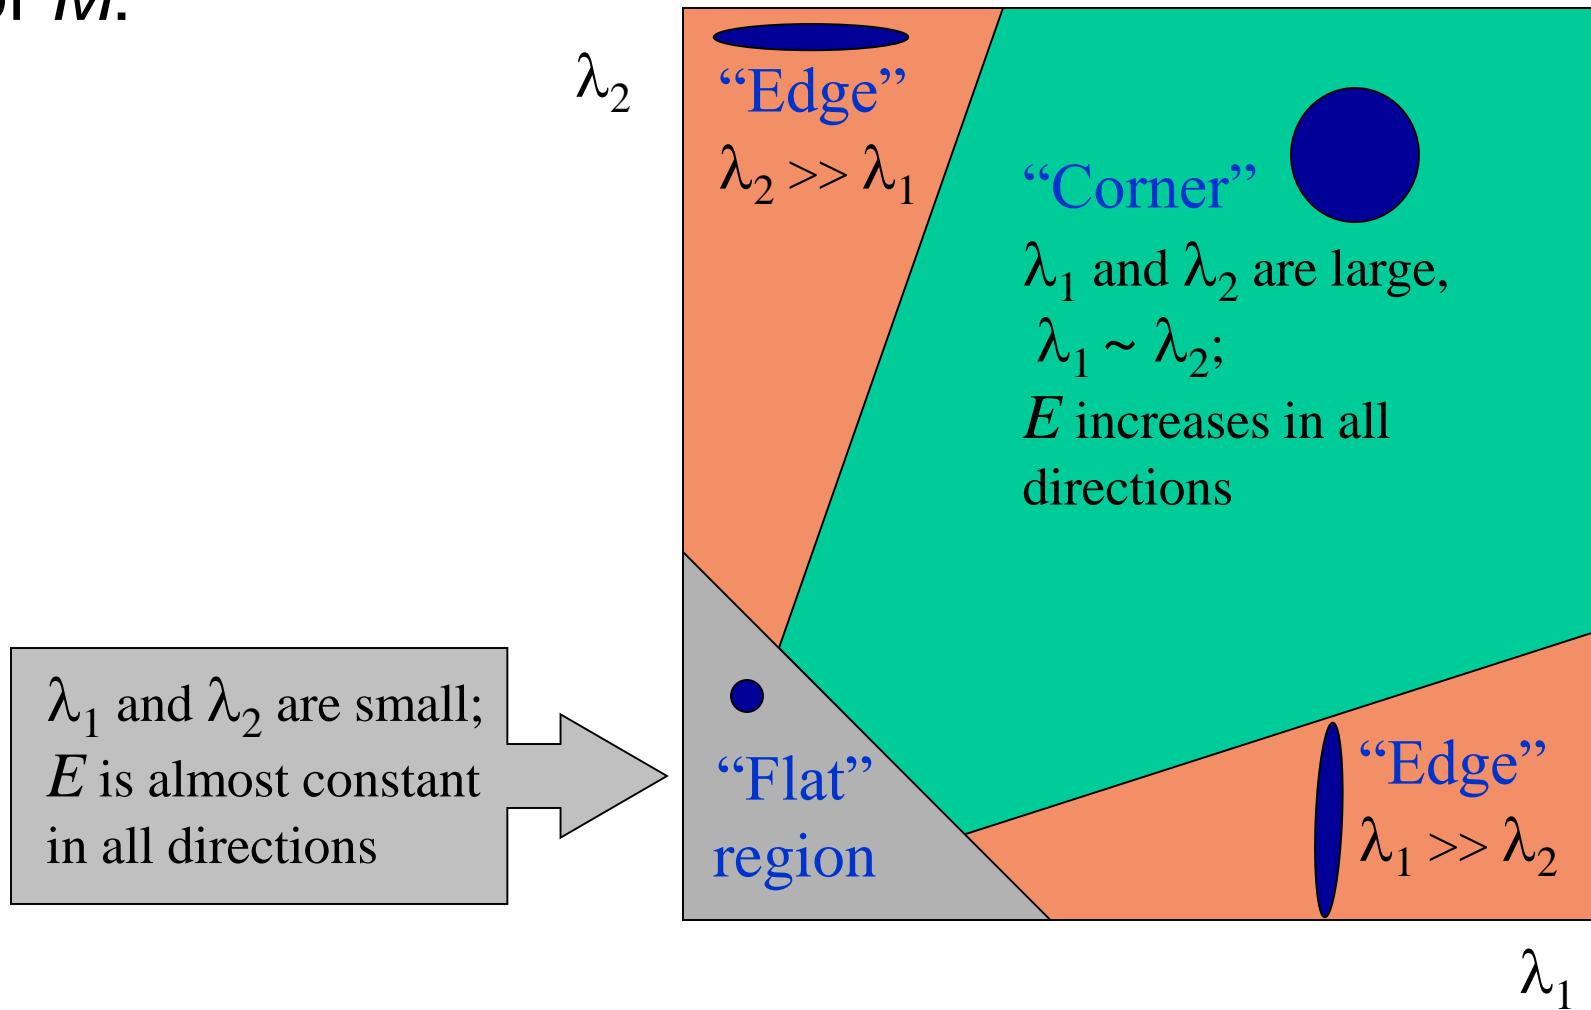
Diagonalization of  $M$ :  $M = R^{-1} \begin{bmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{bmatrix} R$

The axis lengths of the ellipse are determined by the eigenvalues and the orientation is determined by  $R$



# Interpreting the eigenvalues

Classification of image points using eigenvalues of  $M$ :

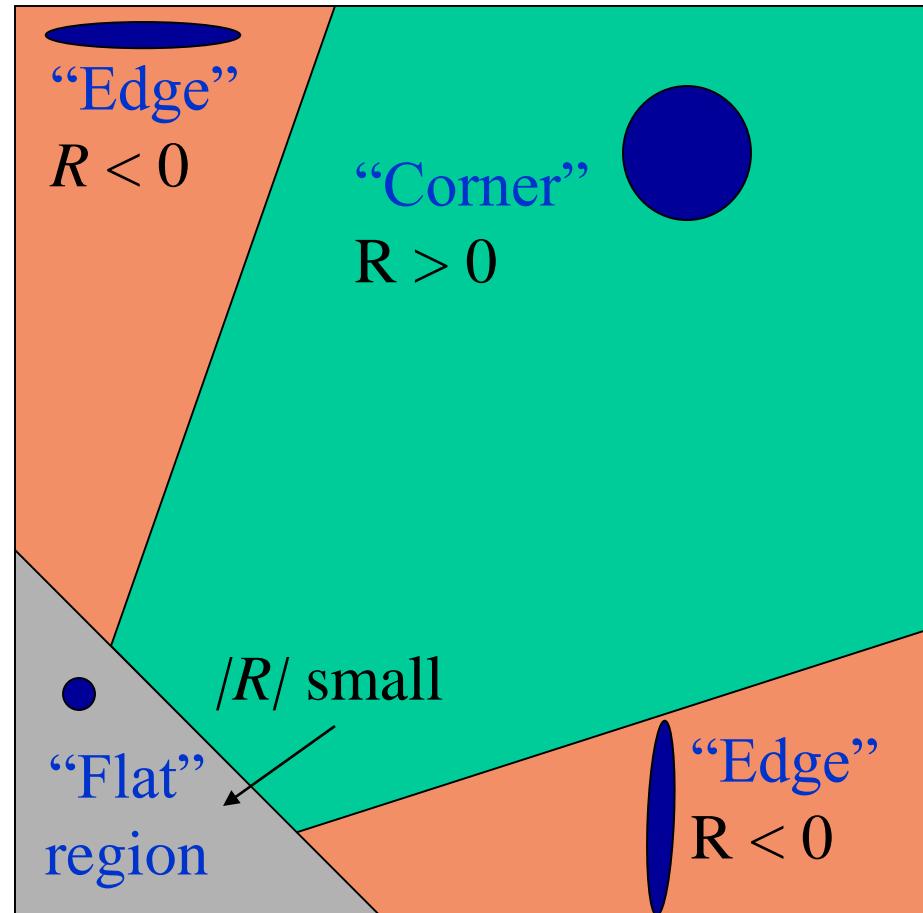


# Corner response function

---

$$R = \det(M) - \alpha \operatorname{trace}(M)^2 = \lambda_1 \lambda_2 - \alpha(\lambda_1 + \lambda_2)^2$$

$\alpha$ : constant (0.04 to 0.06)



# Harris corner detector

---

- 1) Compute  $M$  matrix for each image window to get their *cornerness* scores.
- 2) Find points whose surrounding window gave large corner response ( $f >$  threshold)
- 3) Take the points of local maxima, i.e., perform non-maximum suppression

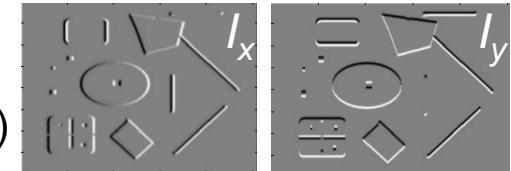
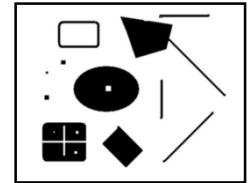
C.Harris and M.Stephens. "[A Combined Corner and Edge Detector.](#)"  
*Proceedings of the 4th Alvey Vision Conference*: pages 147—151, 1988.

# Harris Detector [Harris88]

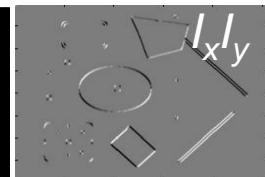
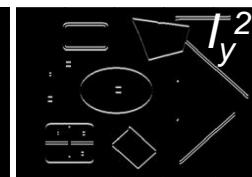
- Second moment matrix

$$\mu(\sigma_I, \sigma_D) = g(\sigma_I) * \begin{bmatrix} I_x^2(\sigma_D) & I_x I_y(\sigma_D) \\ I_x I_y(\sigma_D) & I_y^2(\sigma_D) \end{bmatrix}$$

1. Image derivatives  
(optionally, blur first)



2. Square of derivatives



3. Gaussian filter  $g(\sigma_\nu)$



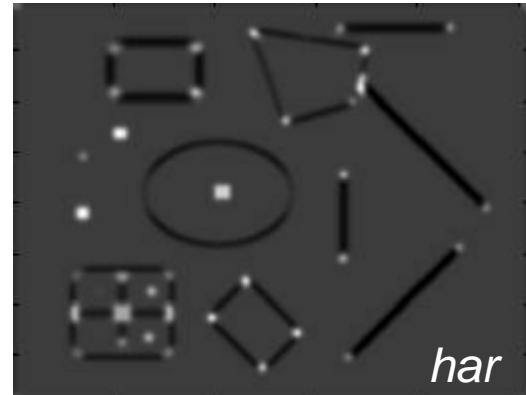
$$\det M = \lambda_1 \lambda_2$$

$$\text{trace } M = \lambda_1 + \lambda_2$$

4. Cornerness function – both eigenvalues are strong

$$\begin{aligned} har &= \det[\mu(\sigma_I, \sigma_D)] - \alpha[\text{trace}(\mu(\sigma_I, \sigma_D))^2] = \\ &= g(I_x^2)g(I_y^2) - [g(I_x I_y)]^2 - \alpha[g(I_x^2) + g(I_y^2)]^2 \end{aligned}$$

5. Non-maxima suppression



# Harris Detector: Steps

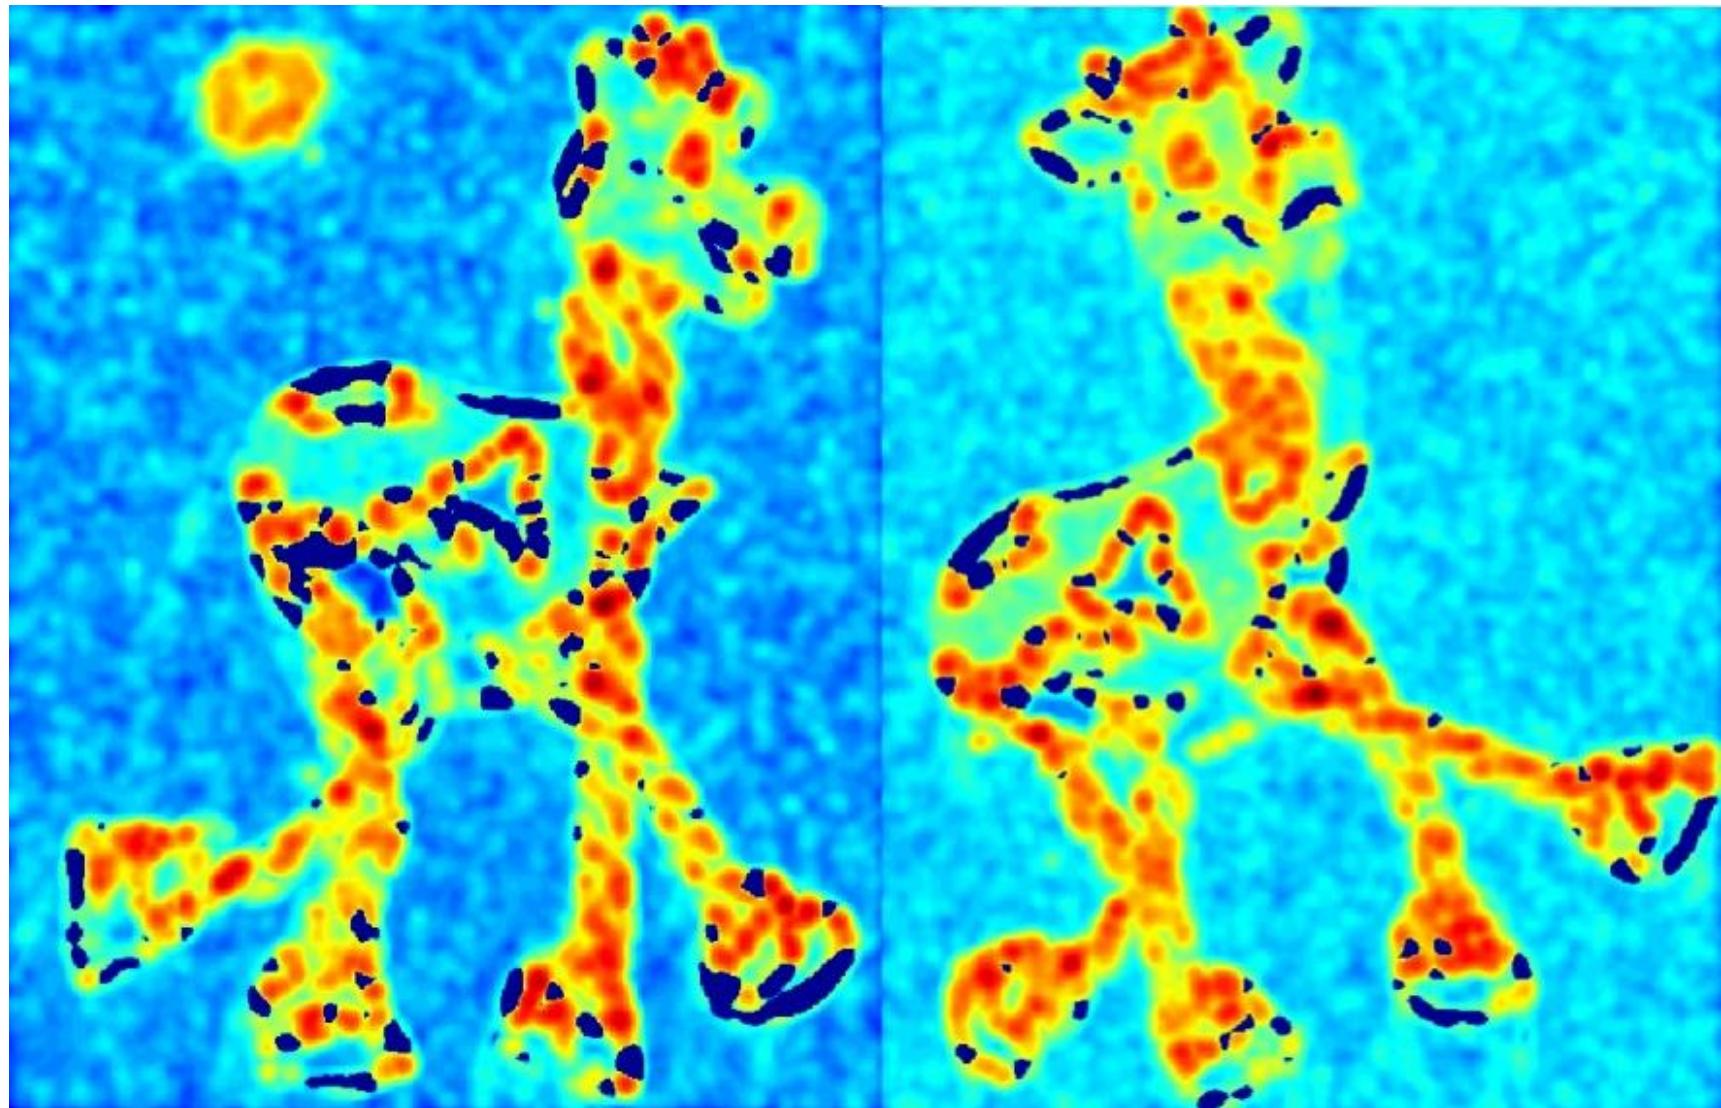
---



# Harris Detector: Steps

---

Compute corner response  $R$



# Harris Detector: Steps

---

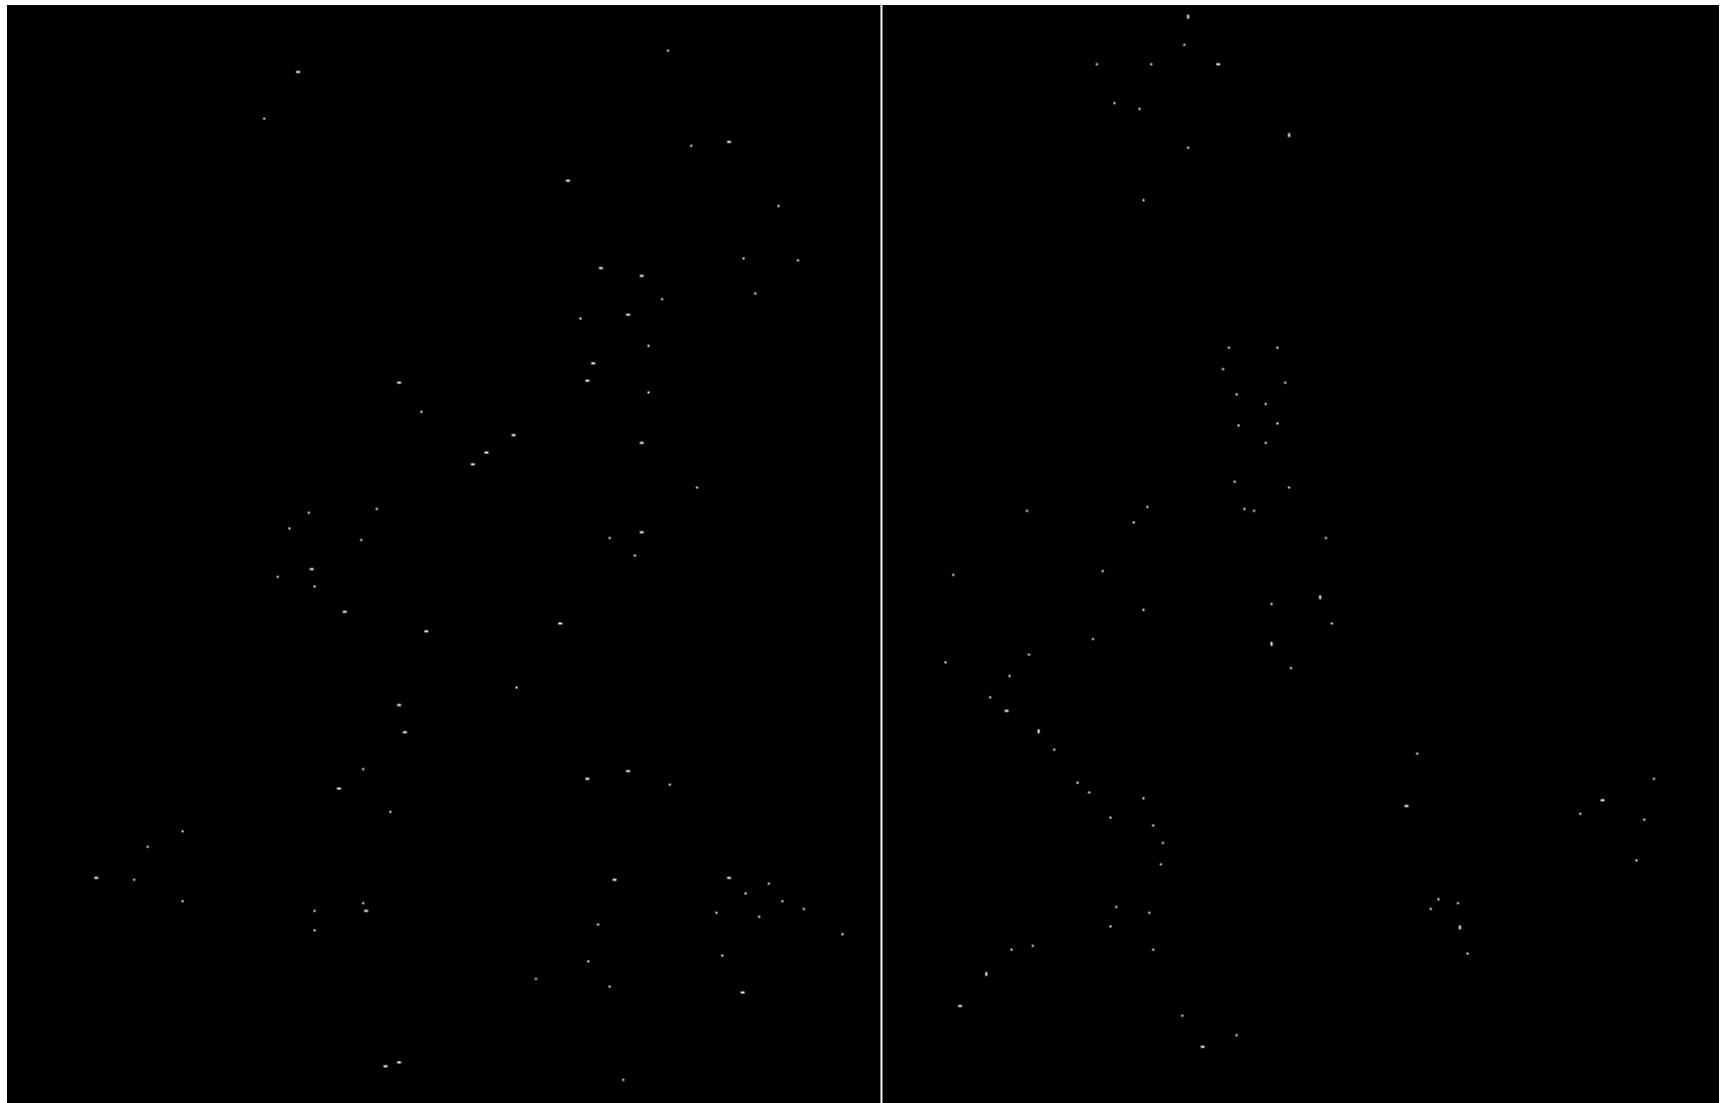
Find points with large corner response:  $R>\text{threshold}$



# Harris Detector: Steps

---

Take only the points of local maxima of  $R$



# Harris Detector: Steps

---



# Invariance and covariance

---

- We want corner locations to be *invariant* to photometric transformations and *covariant* to geometric transformations
  - **Invariance:** image is transformed and corner locations do not change
  - **Covariance:** if we have two transformed versions of the same image, features should be detected in corresponding locations



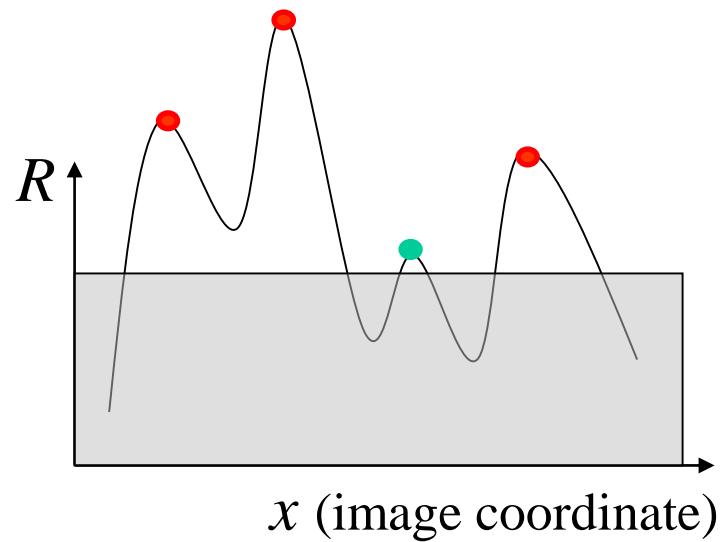
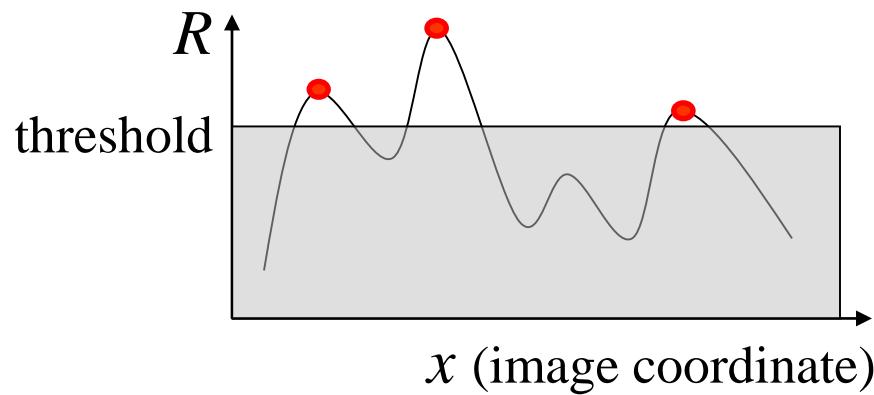
# Affine intensity change

---



$$I \rightarrow a I + b$$

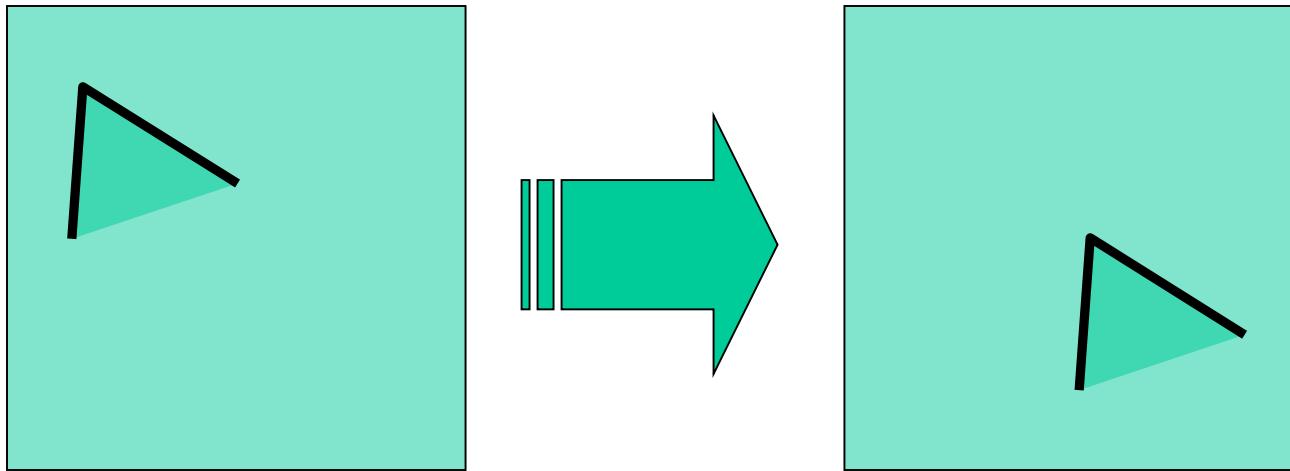
- Only derivatives are used => invariance to intensity shift  $I \rightarrow I + b$
- Intensity scaling:  $I \rightarrow a I$



*Partially invariant to affine intensity change*

# Image translation

---

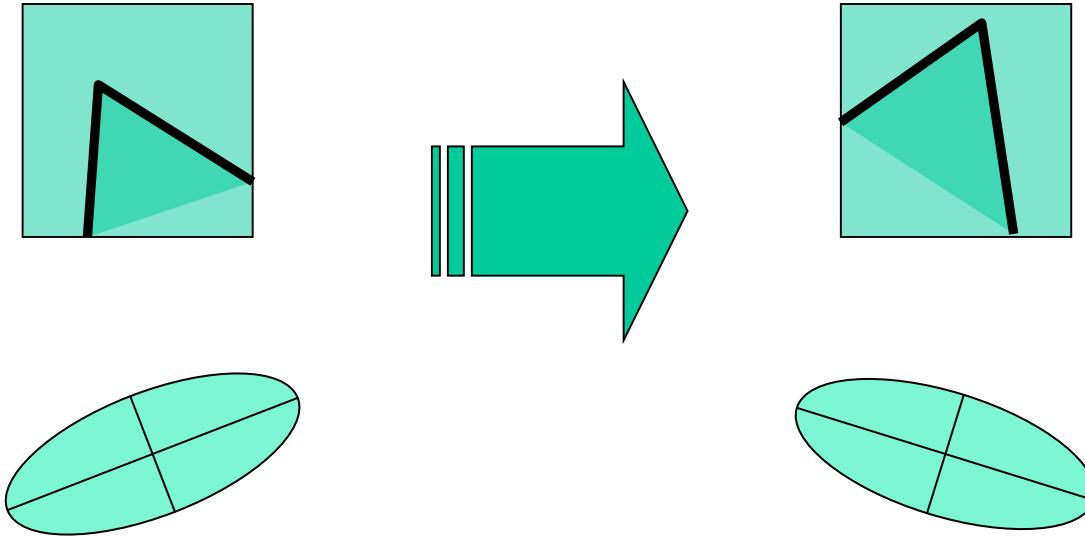


- Derivatives and window function are shift-invariant

Corner location is covariant w.r.t. translation

# Image rotation

---

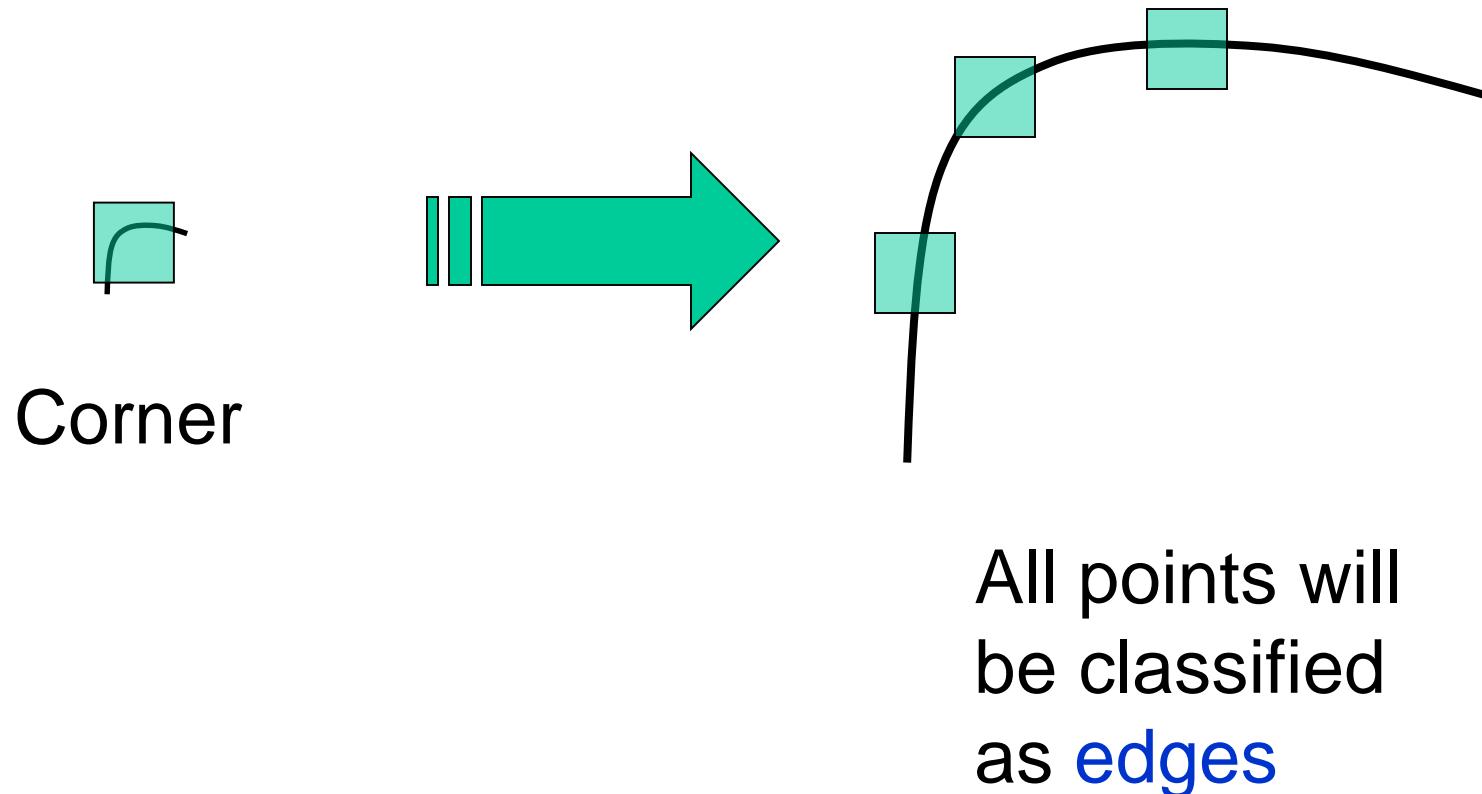


Second moment ellipse rotates but its shape  
(i.e. eigenvalues) remains the same

Corner location is covariant w.r.t. rotation

# Scaling

---



Corner location is not covariant to scaling!