

# Let's look at some lakefront property

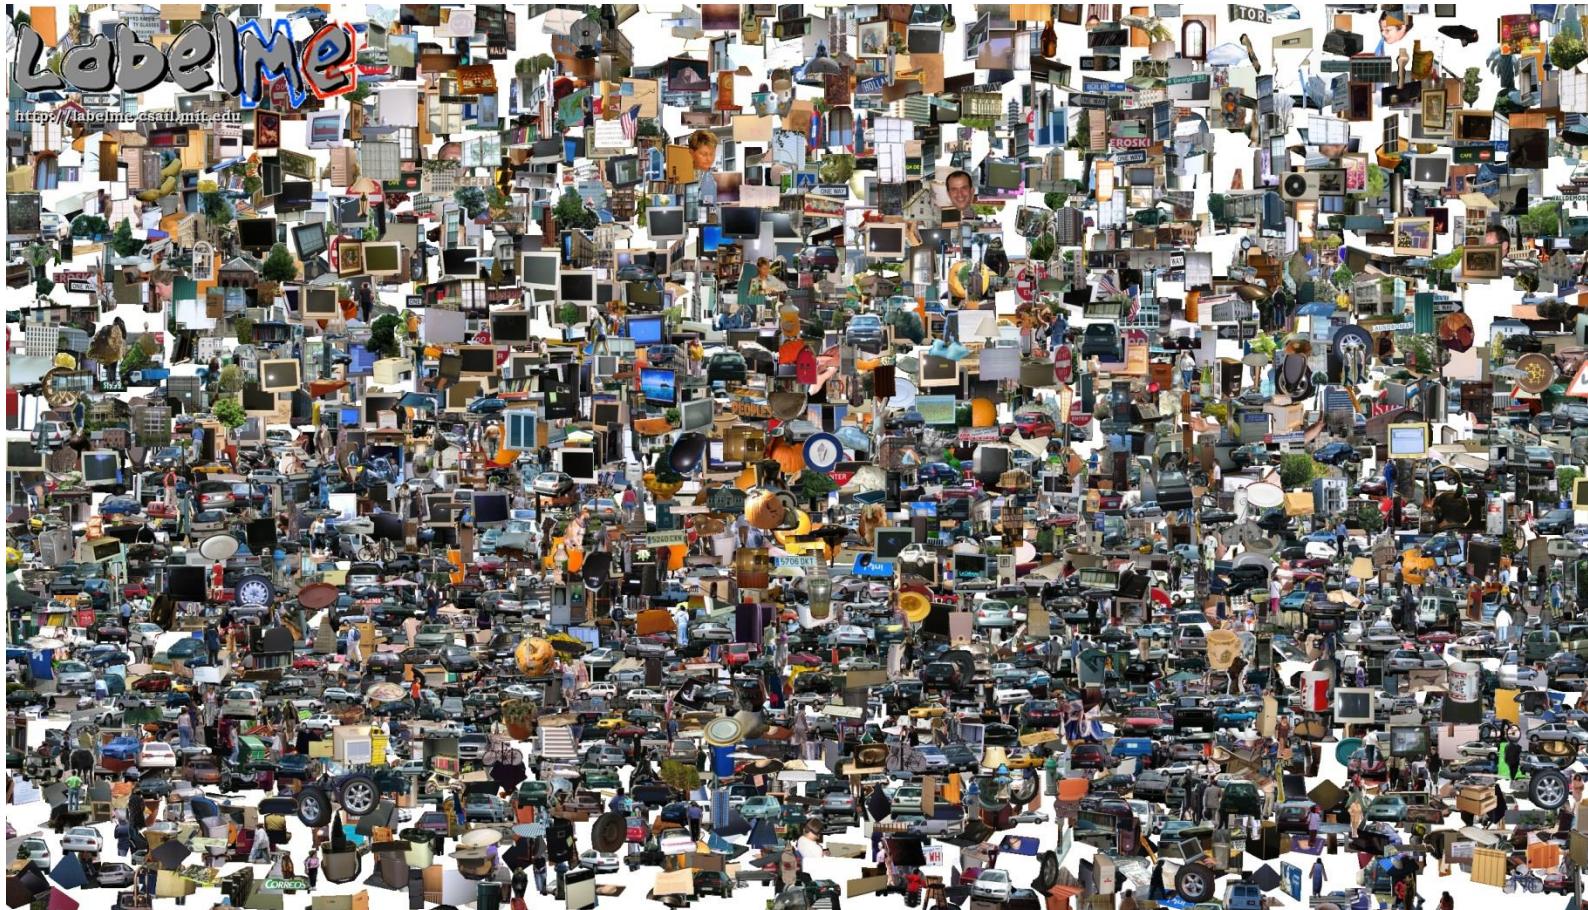


\*actually fences / walls





# Opportunities of Scale



Computer Vision  
James Hays

# Outline

## Opportunities of Scale: Data-driven methods

- The Unreasonable Effectiveness of Data
- Scene Completion
- Im2gps
- Recognition via Tiny Images
- Project 5 Intro

# Computer Vision so far

- The geometry of image formation
  - Ancient / Renaissance
- Signal processing / Convolution
  - 1800, but really the 50's and 60's
- Hand-designed Features for recognition,  
either instance-level or categorical
  - 1999 (SIFT), 2003 (Video Google), 2005 (Dalal-Triggs), 2006 (spatial pyramid)
- Learning from Data
  - 1991 (EigenFaces) but late 90's to now especially

# What has changed in the last decade?

- The Internet
- Crowdsourcing
- Learning representations from the data these sources provide (deep learning)

# Google and massive data-driven algorithms

## A.I. for the postmodern world:

- all questions have already been answered...many times, in many ways
- Google is dumb, the “intelligence” is in the data



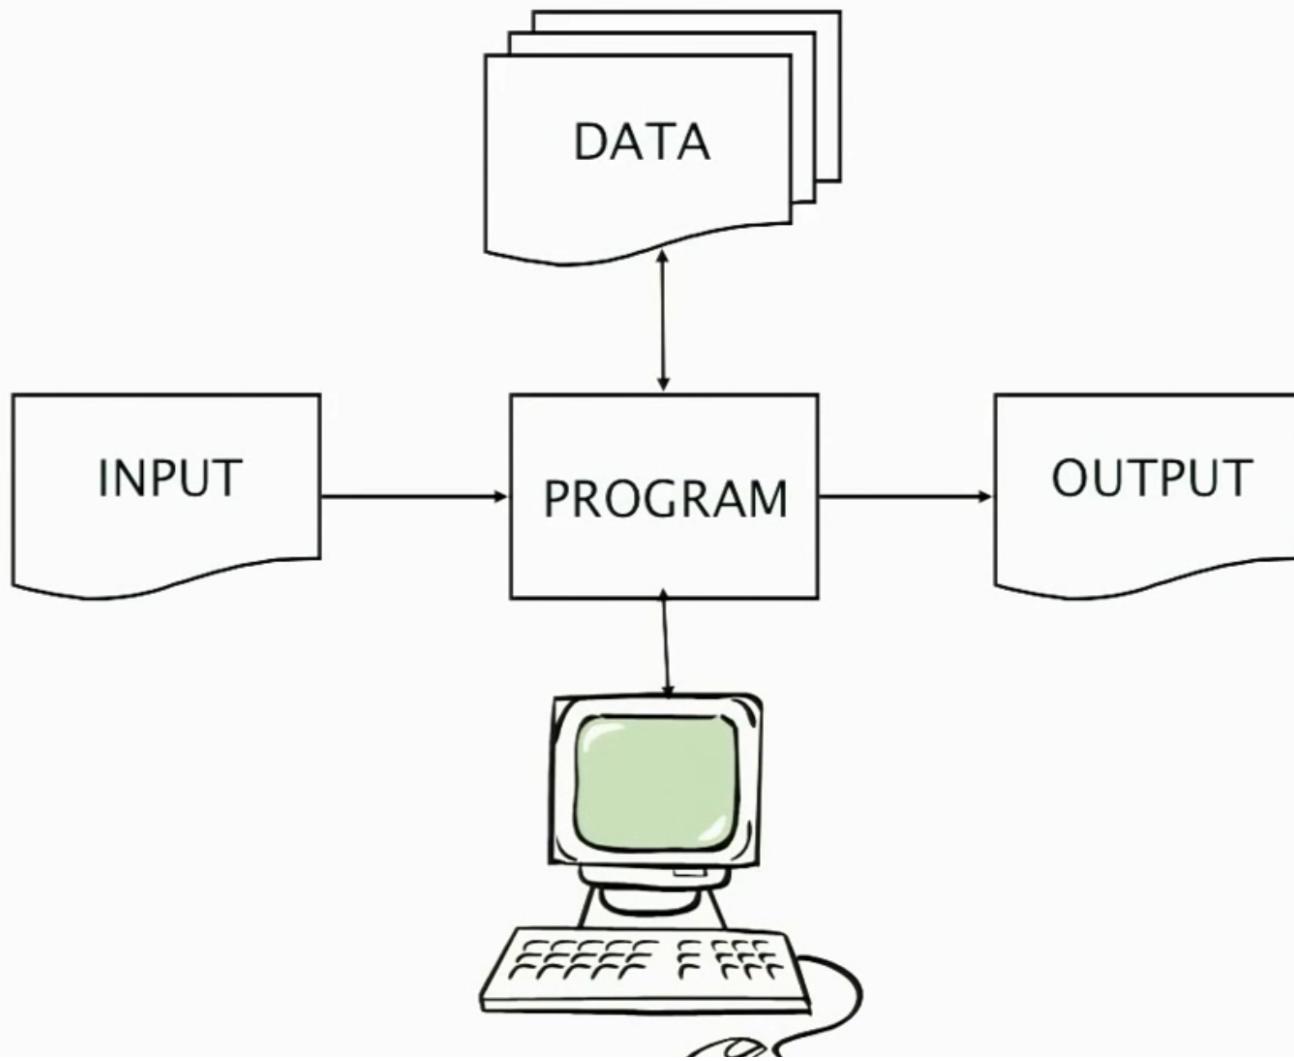
# The Unreasonable Effectiveness of Data

Peter Norvig  
Google



Peter Norvig

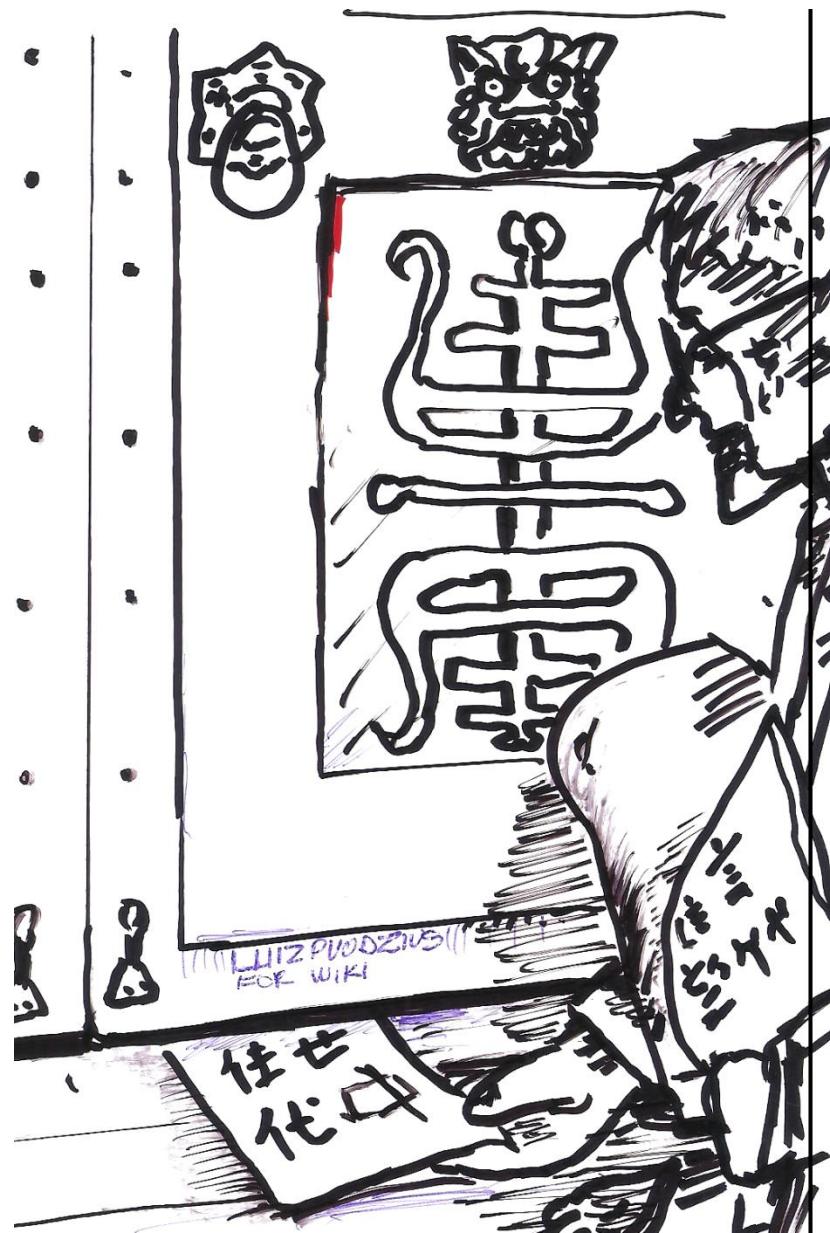
The Unreasonable Effectiveness of Data



# Chinese Room, John Searle (1980)

If a machine can convincingly simulate an intelligent conversation, does it necessarily understand? In the experiment, Searle imagines himself in a room, acting as a computer by manually executing a program that convincingly simulates the behavior of a native Chinese speaker.

Most of the discussion consists of attempts to refute it. "The overwhelming majority," notes *BBS* editor Stevan Harnad, "still think that the Chinese Room Argument is dead wrong." The sheer volume of the literature that has grown up around it inspired Pat Hayes to quip that the field of cognitive science ought to be redefined as "the ongoing research program of showing Searle's Chinese Room Argument to be false."





**Yann LeCun**

October 23 at 9:58pm ·

Questions from the piece:

Q1. Does the Chinese Room argument prove the impossibility of machine consciousness?

A1: Hell no. ... See More



## Can Machines Become Moral?

The question is heard more and more often, both from those who think that machines cannot become moral, and who think that to believe otherwise is a dangerous illusion, and from those who think that machines must become moral,...

BIGQUESTIONSONLINE.COM | BY DON HOWARD



You and 156 others

30 Comments 20 Shares

Like

Comment

Share

# Big Idea

- Do we need computer vision systems to have strong AI-like reasoning about our world?
- What if invariance / generalization isn't actually the core difficulty of computer vision?
- What if we can perform high level reasoning with brute-force, data-driven algorithms?

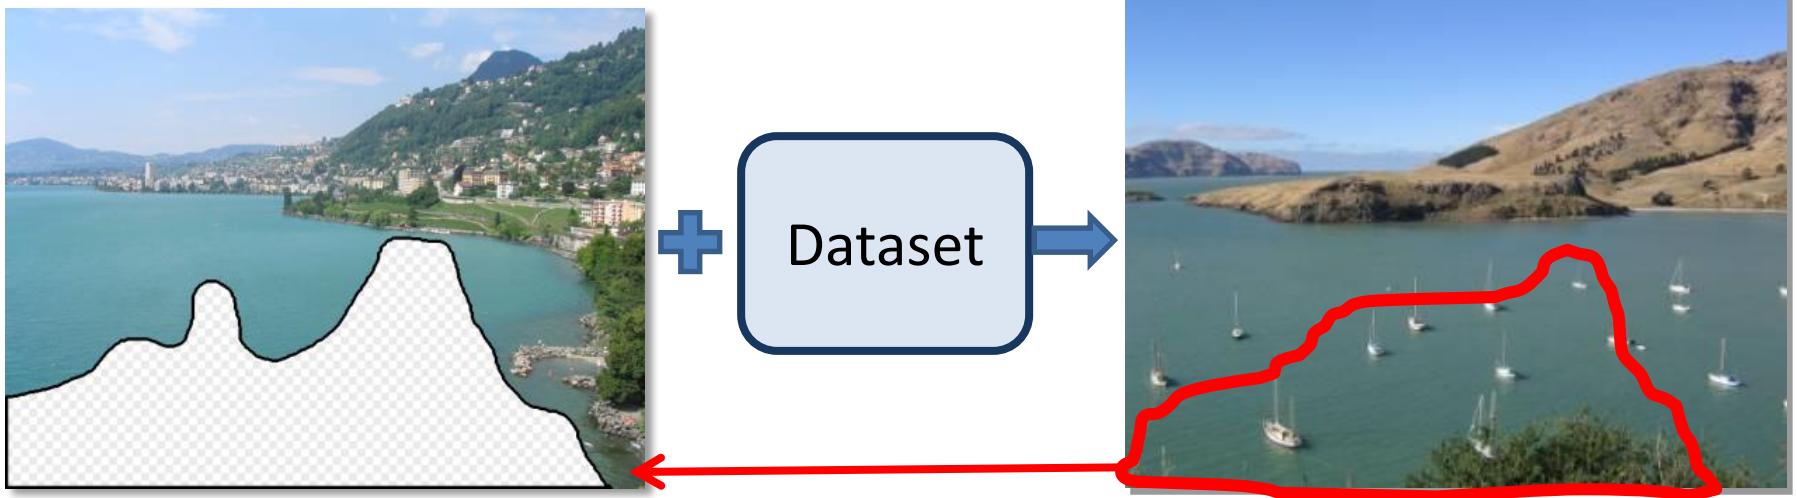
# Scene Completion

[Hays and Efros. Scene Completion Using Millions of Photographs.  
SIGGRAPH 2007 and CACM October 2008.]

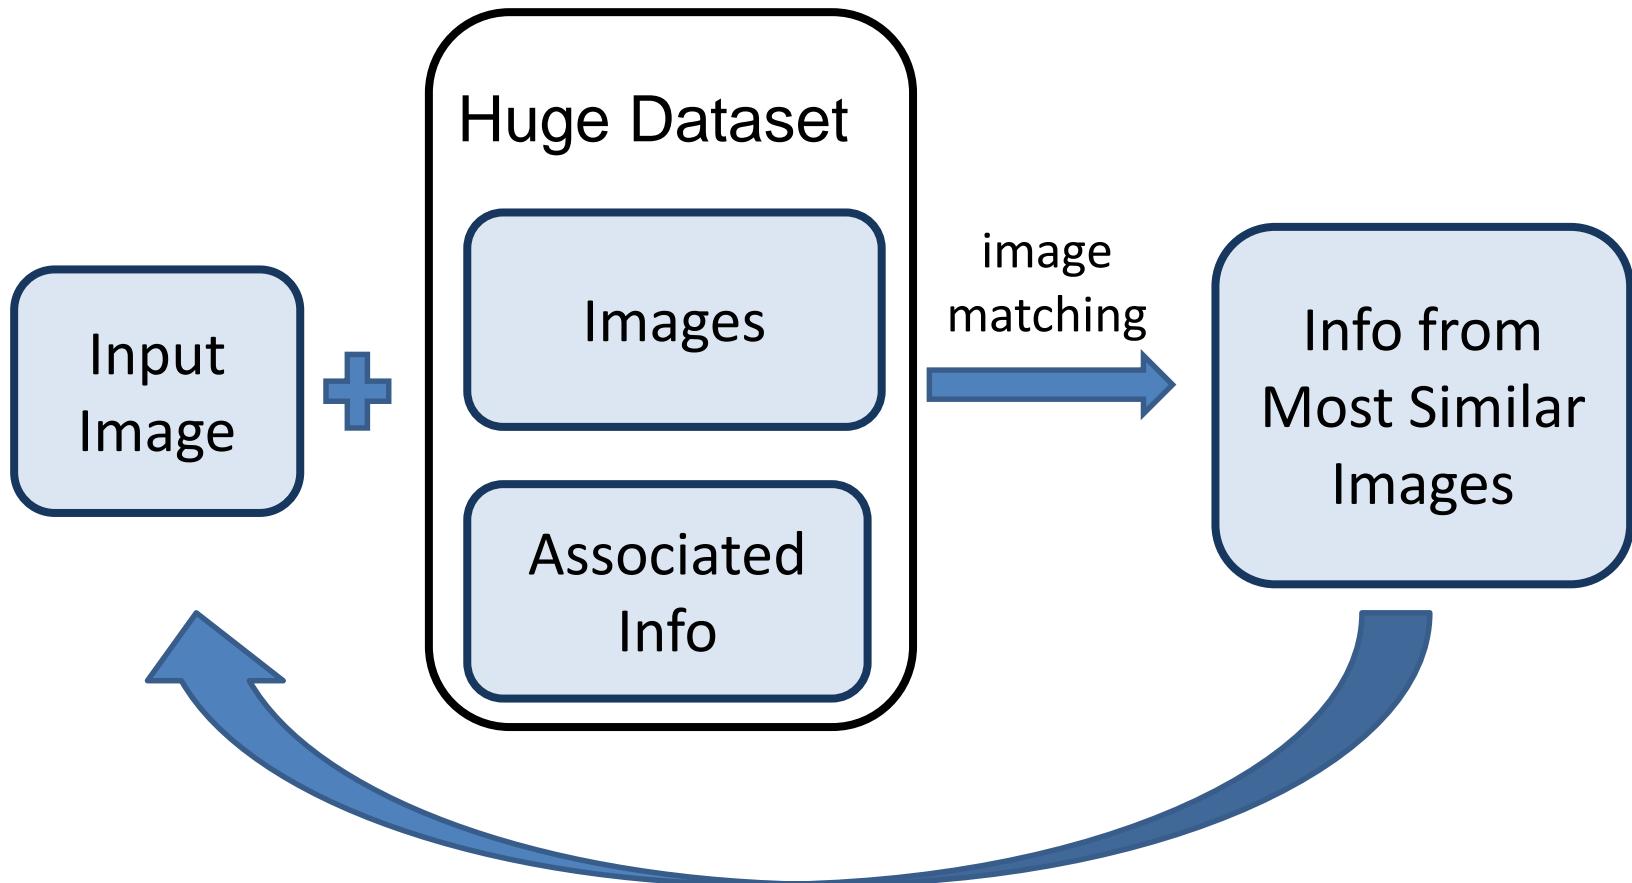
<http://graphics.cs.cmu.edu/projects/scene-completion/>

# How it works

- Find a similar image from a large dataset
- Blend a region from that image into the hole

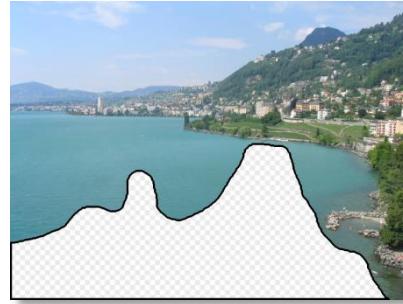


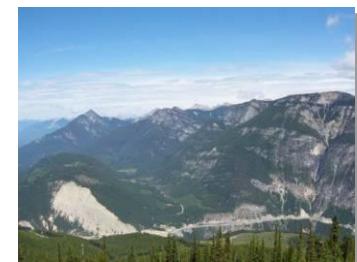
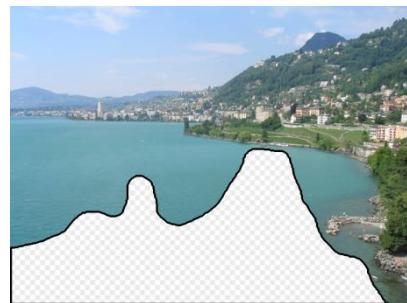
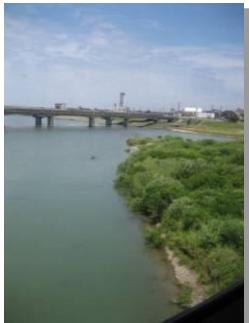
# General Principal



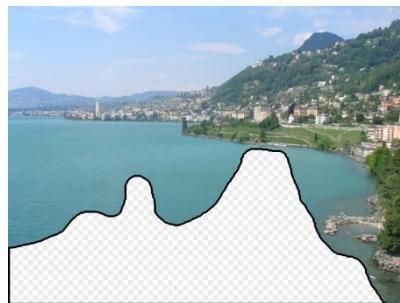
Hopefully, If you have enough images, the dataset will contain very similar images that you can find with simple matching methods.

# How many images is enough?





Nearest neighbors from a  
collection of 20 thousand images



Nearest neighbors from a  
collection of 2 million images

# Image Data on the Internet

- Flickr (as of Sept. 19<sup>th</sup>, 2010)
  - 5 billion photographs
  - 100+ million geotagged images
- Facebook (as of 2009)
  - 15 billion

# Image Data on the Internet

- Flickr (as of Nov 2013)
  - 10 billion photographs
  - 100+ million geotagged images
  - 3.5 million a day
- Facebook (as of Sept 2013)
  - 250 billion+
  - 300 million a day
- Instagram
  - 55 million a day

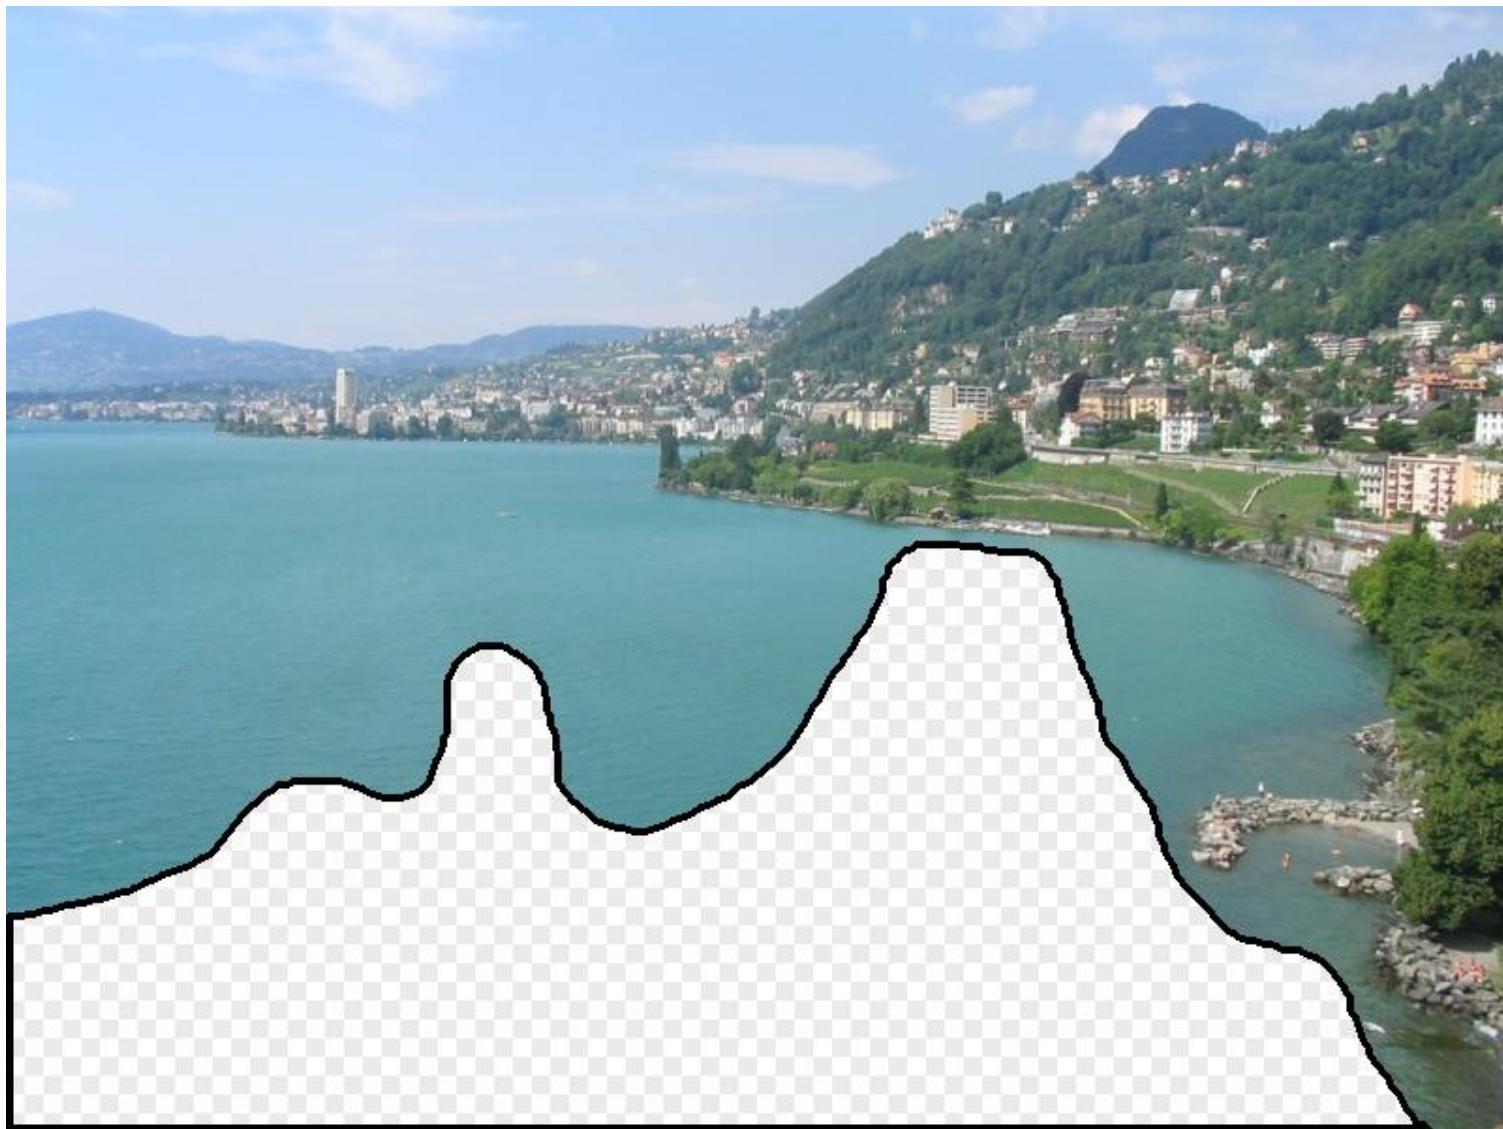
# Scene Completion: how it works

[Hays and Efros. Scene Completion Using Millions of Photographs.  
SIGGRAPH 2007 and CACM October 2008.]

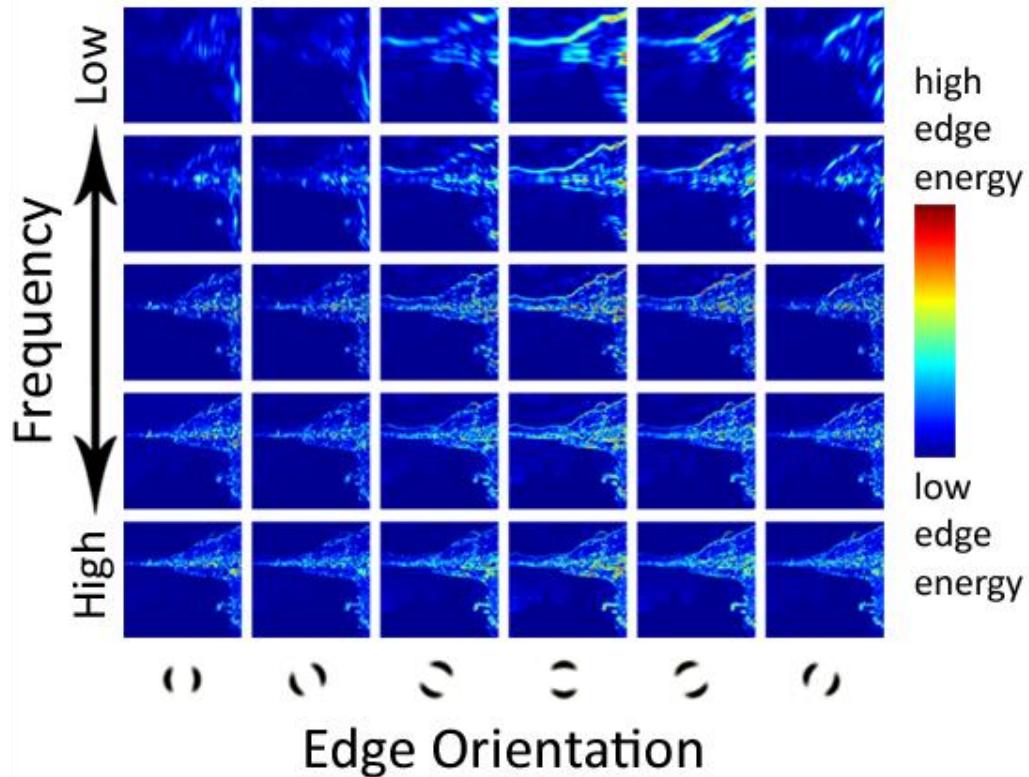
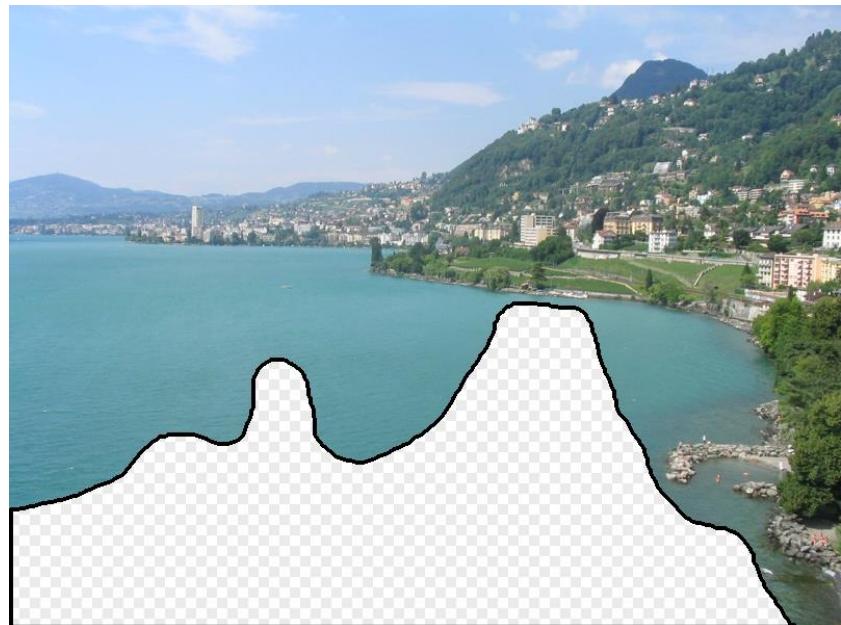
# The Algorithm



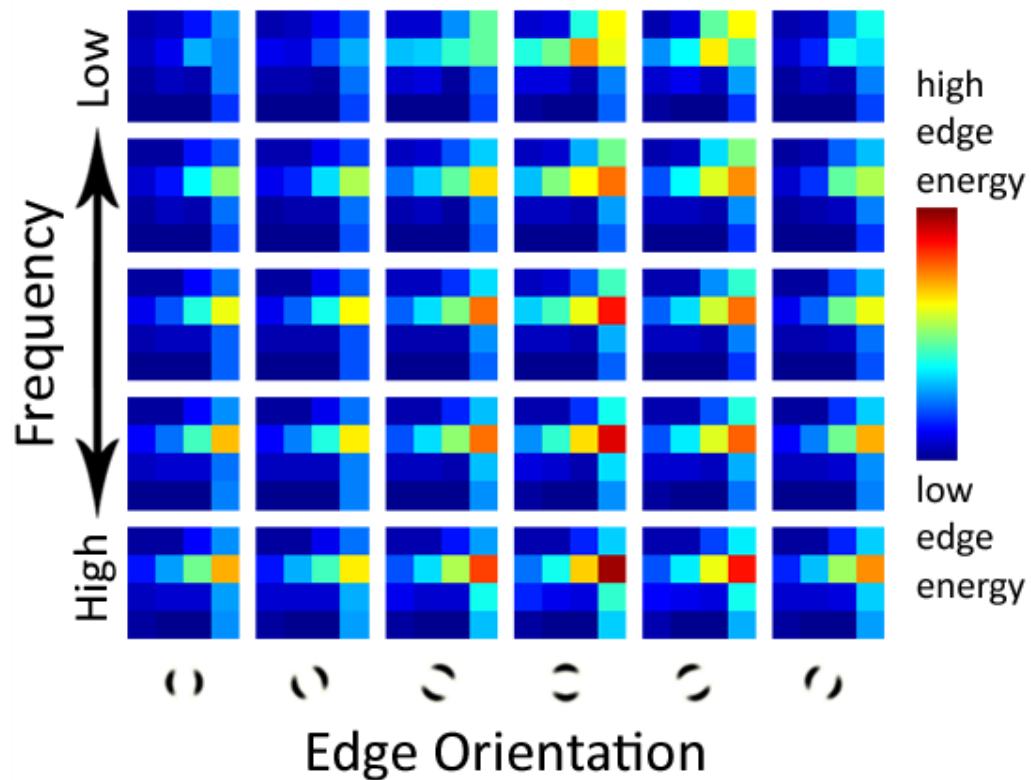
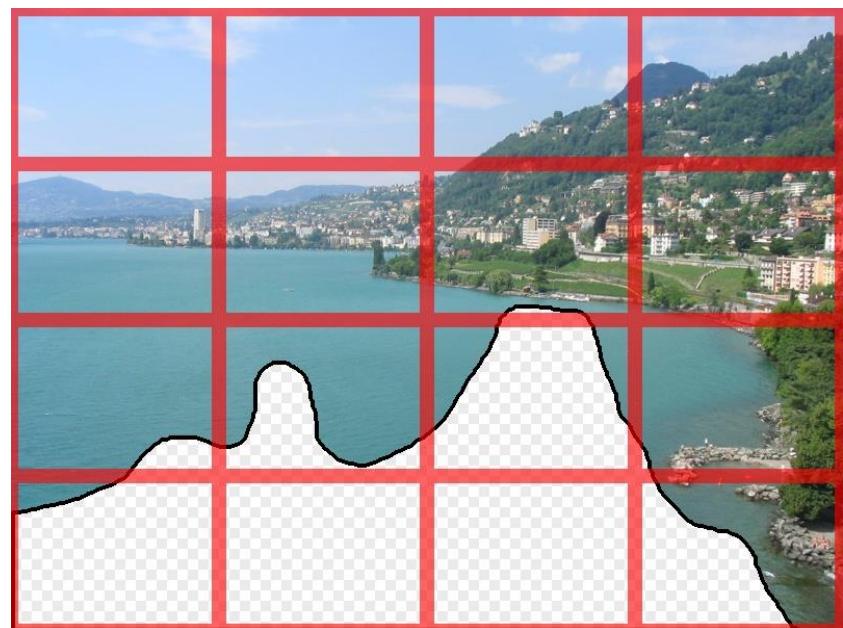
# Scene Matching



# Scene Descriptor

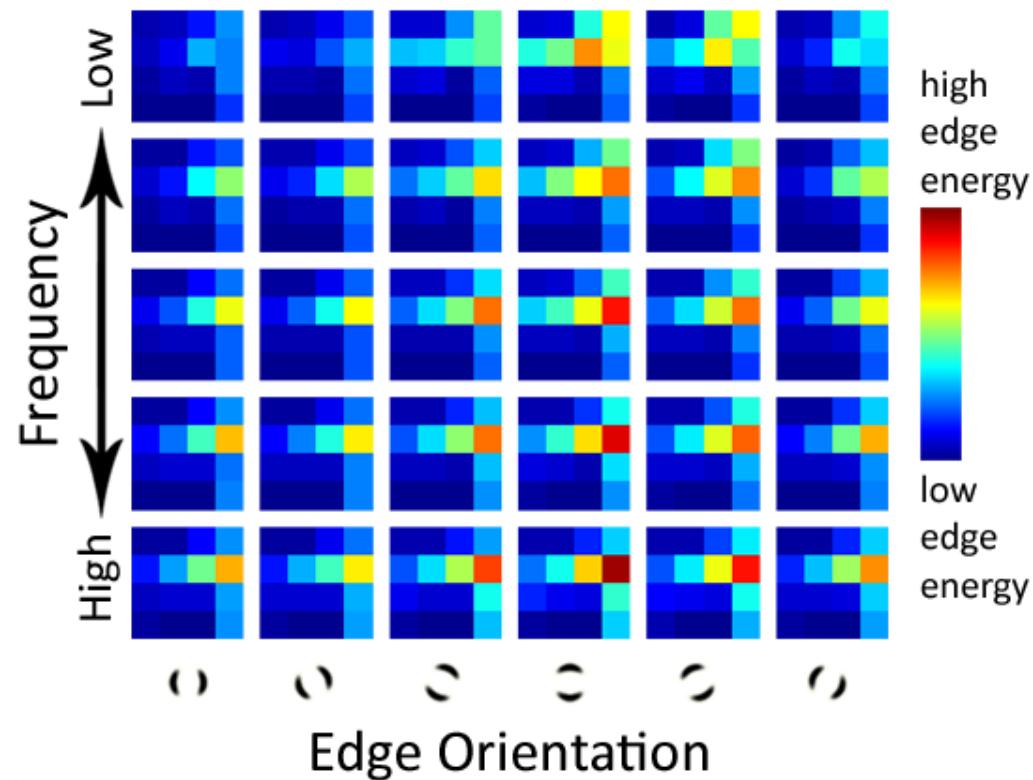
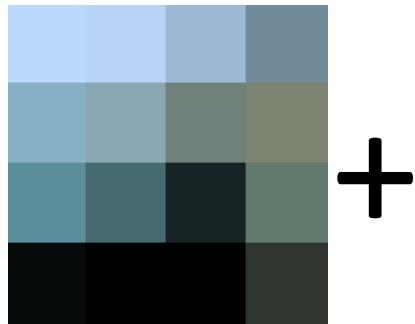


# Scene Descriptor



Scene Gist Descriptor  
(Oliva and Torralba 2001)

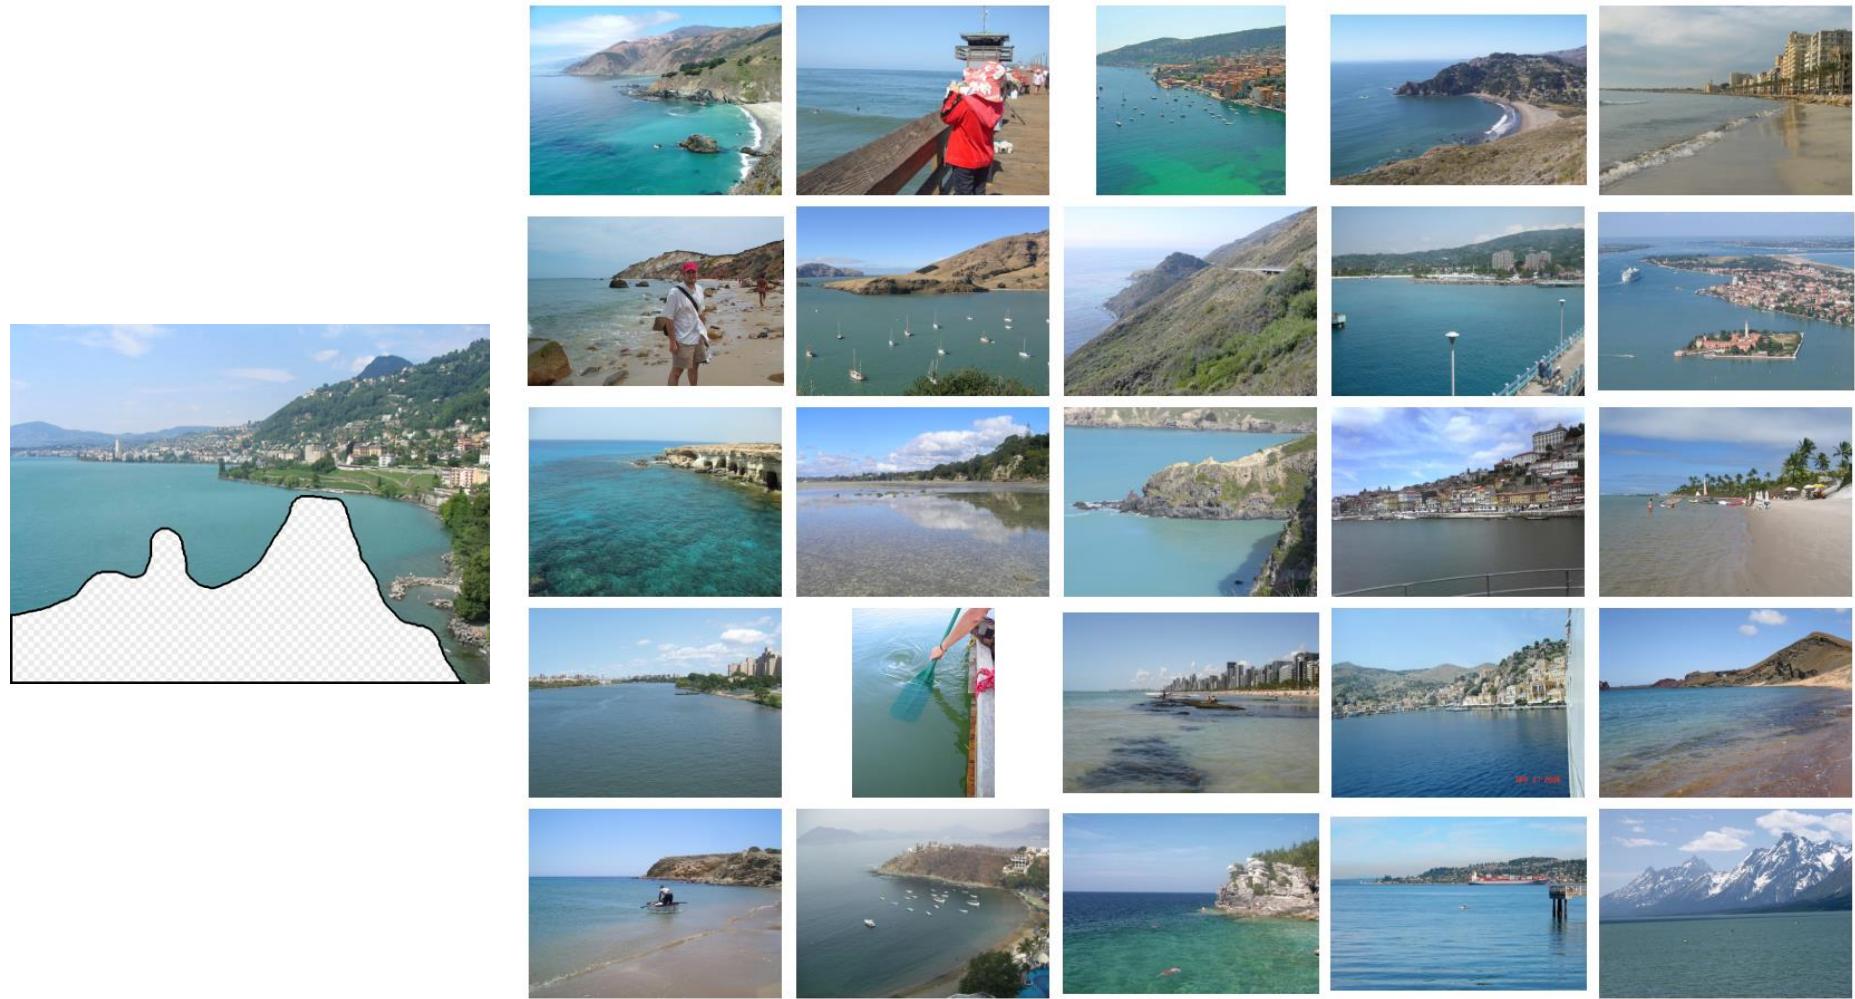
# Scene Descriptor



Scene Gist Descriptor  
(Oliva and Torralba 2001)

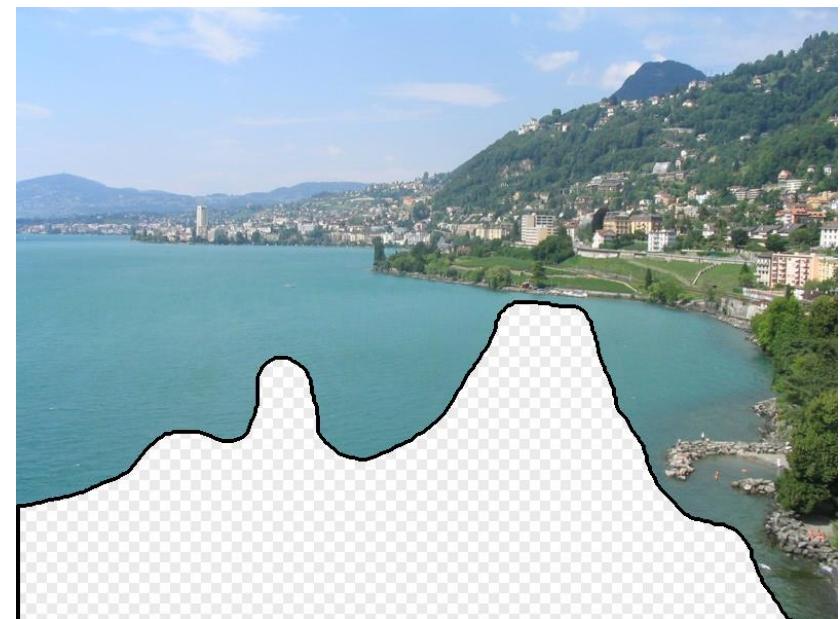
# 2 Million Flickr Images

The background of the image is a dense, uniform grid composed of numerous small, square thumbnail images. These thumbnails represent a vast collection of photographs from Flickr, showing a wide variety of subjects and colors. The overall effect is a visual representation of the scale and diversity of user-generated content.



... 200 total

# Context Matching

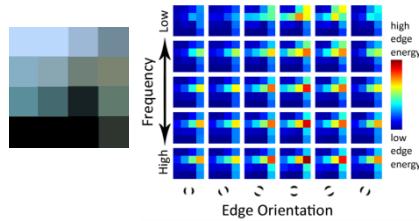




Graph cut + Poisson blending

# Result Ranking

We assign each of the 200 results a score which is the sum of:



The scene matching distance



The context matching distance  
(color + texture)



The graph cut cost

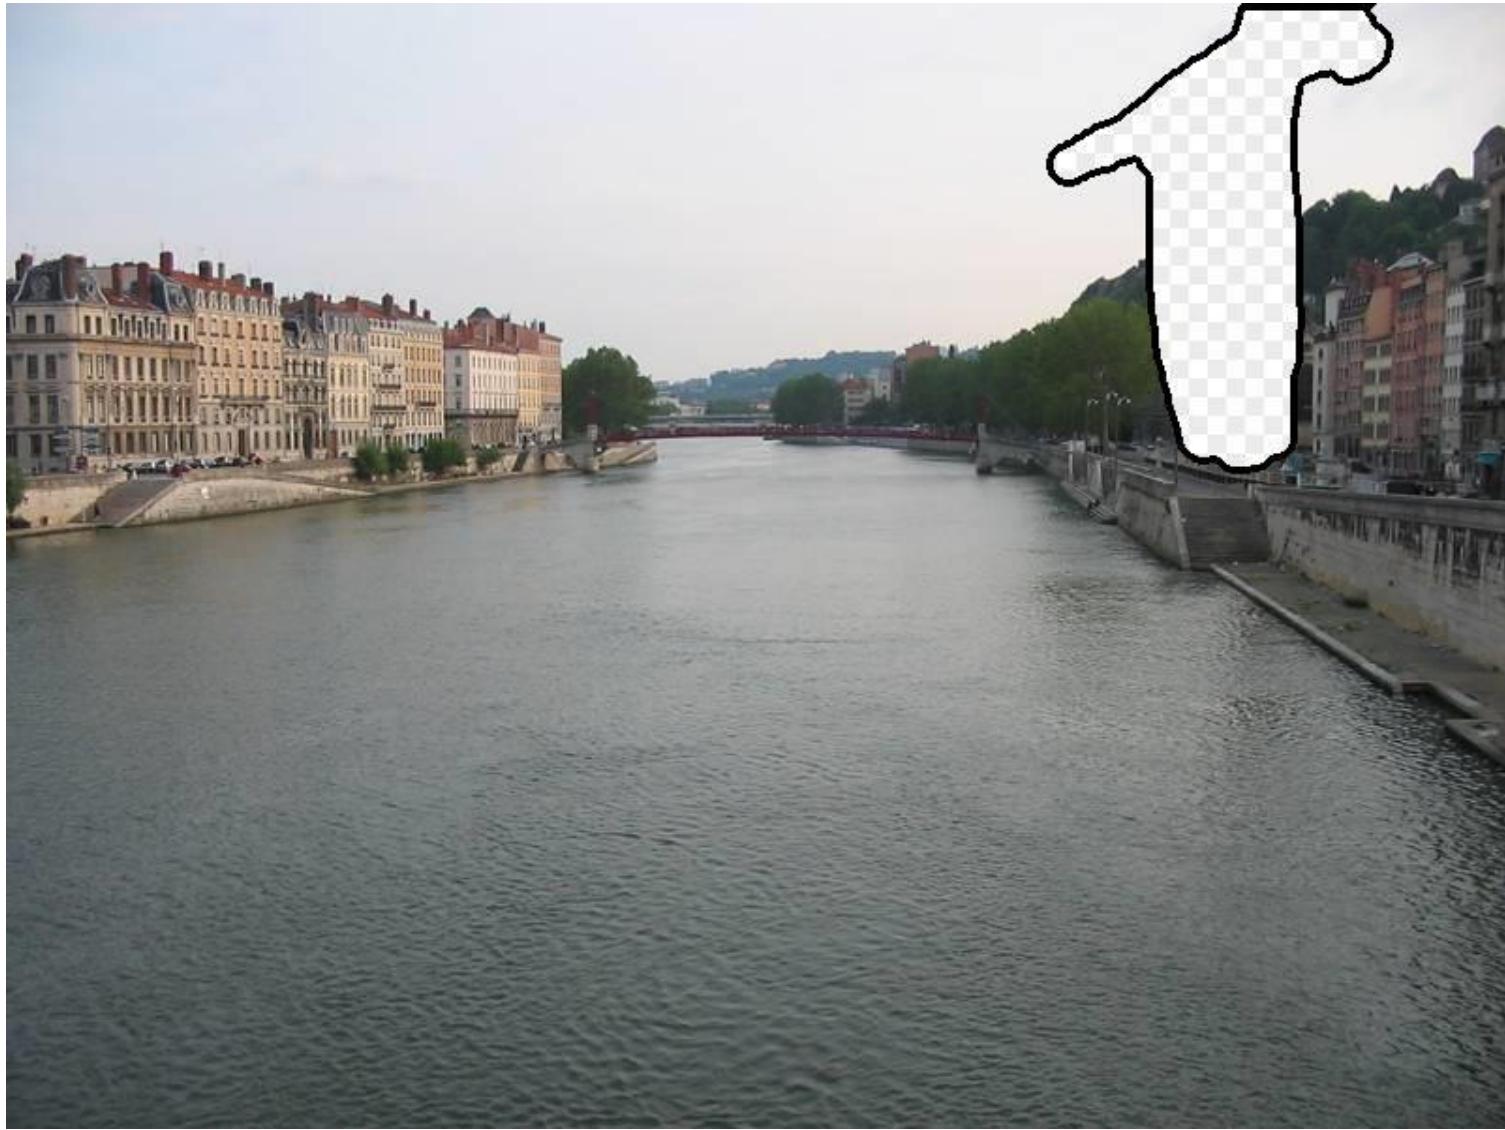




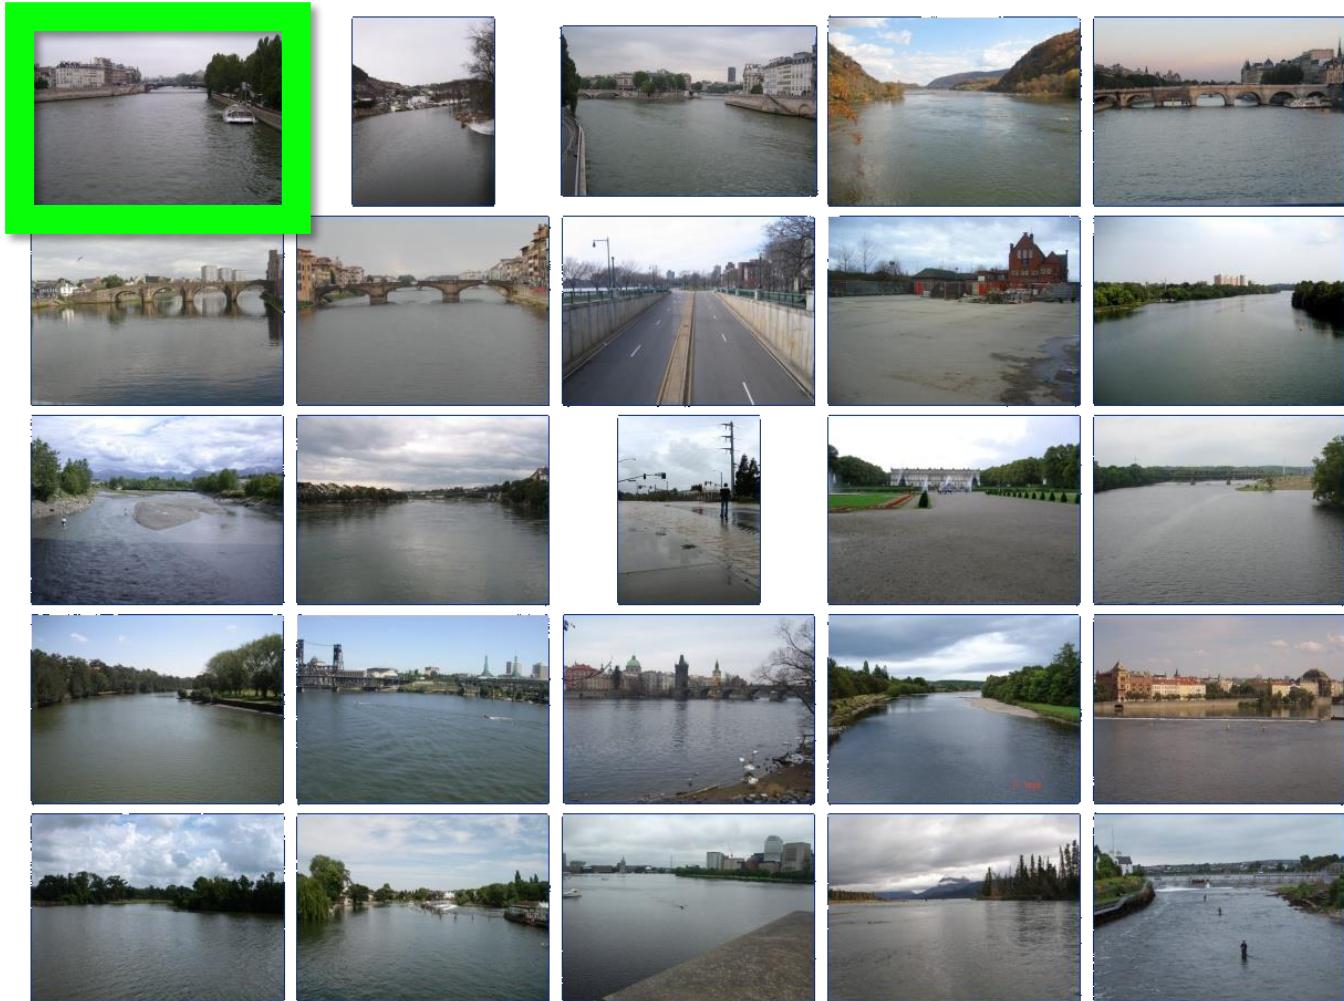








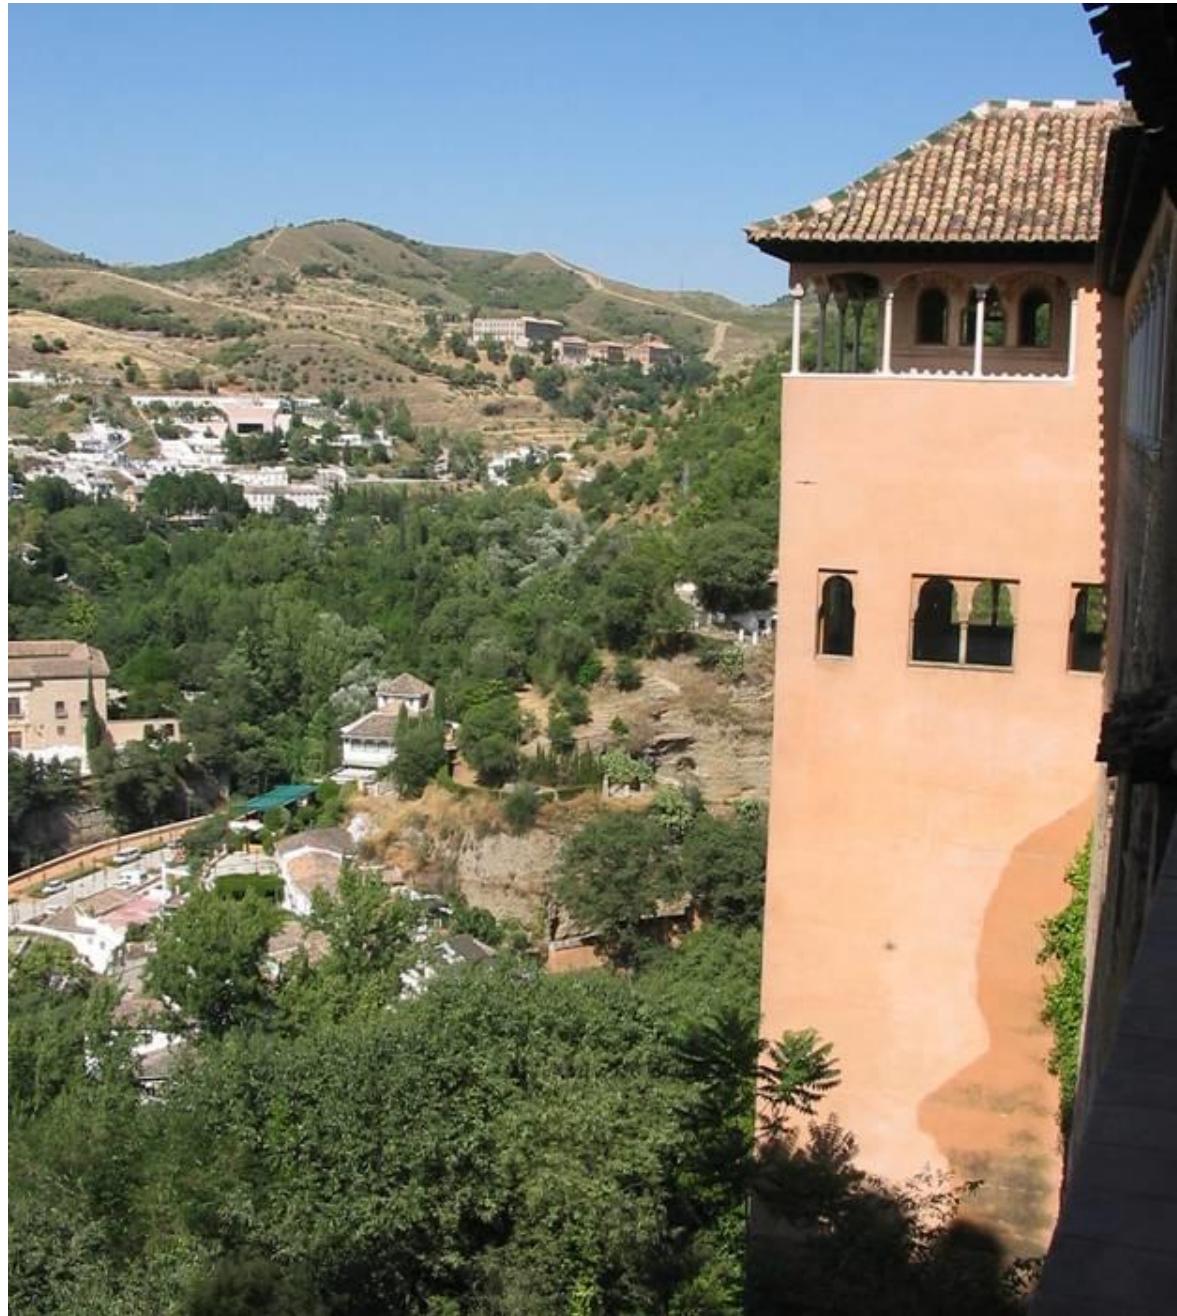


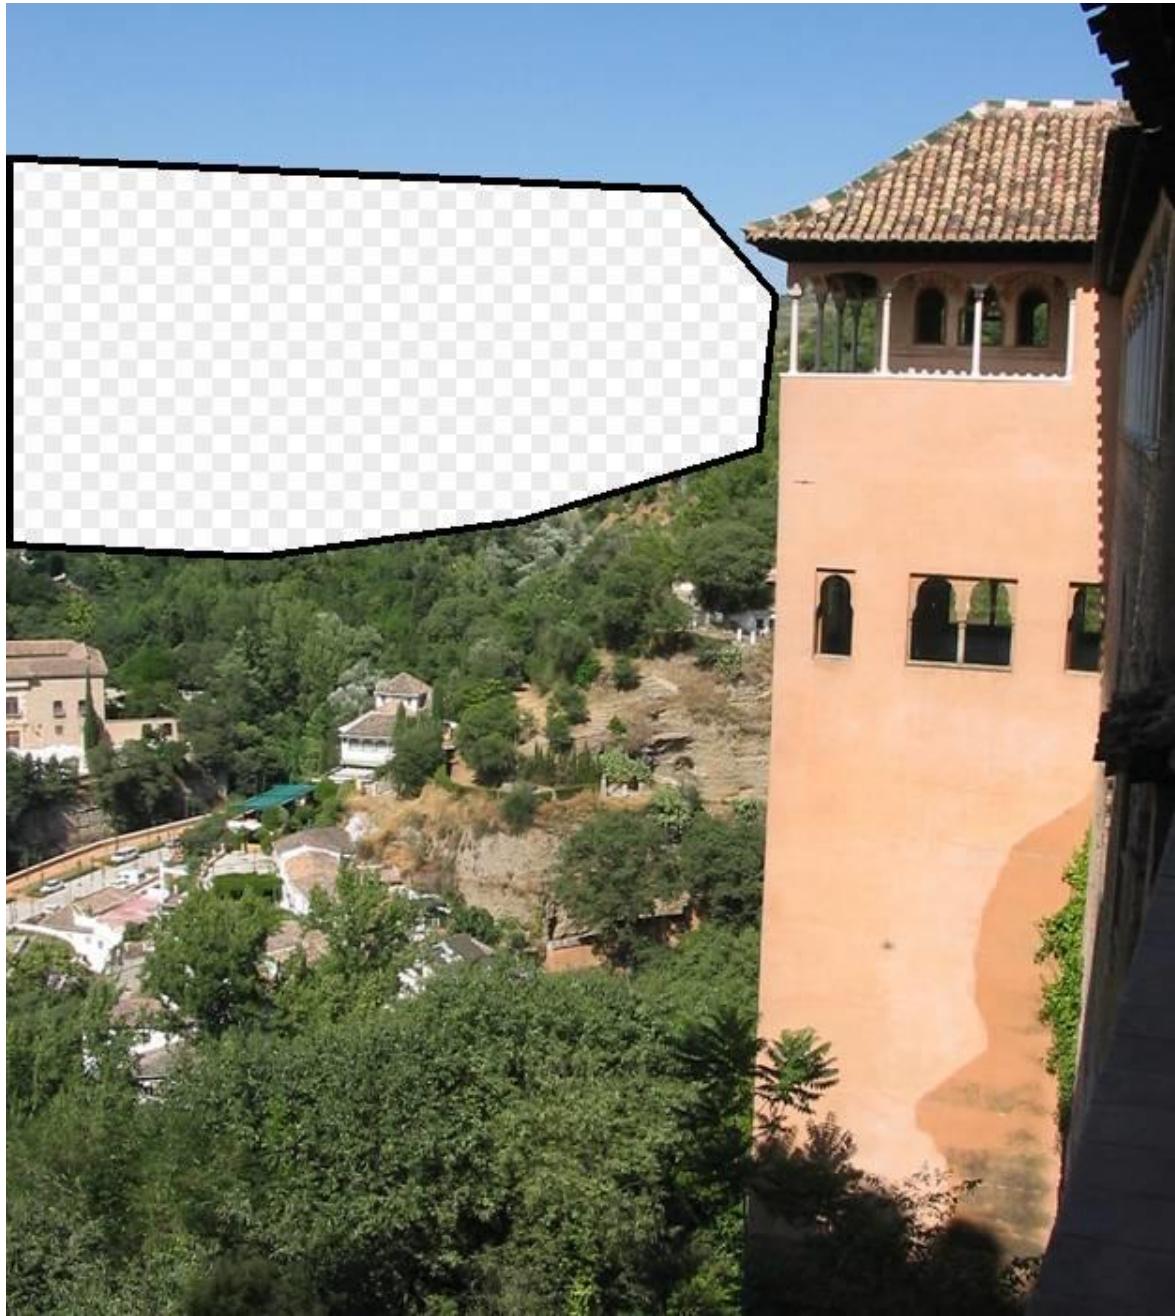


... 200 scene matches











# Which is the original?





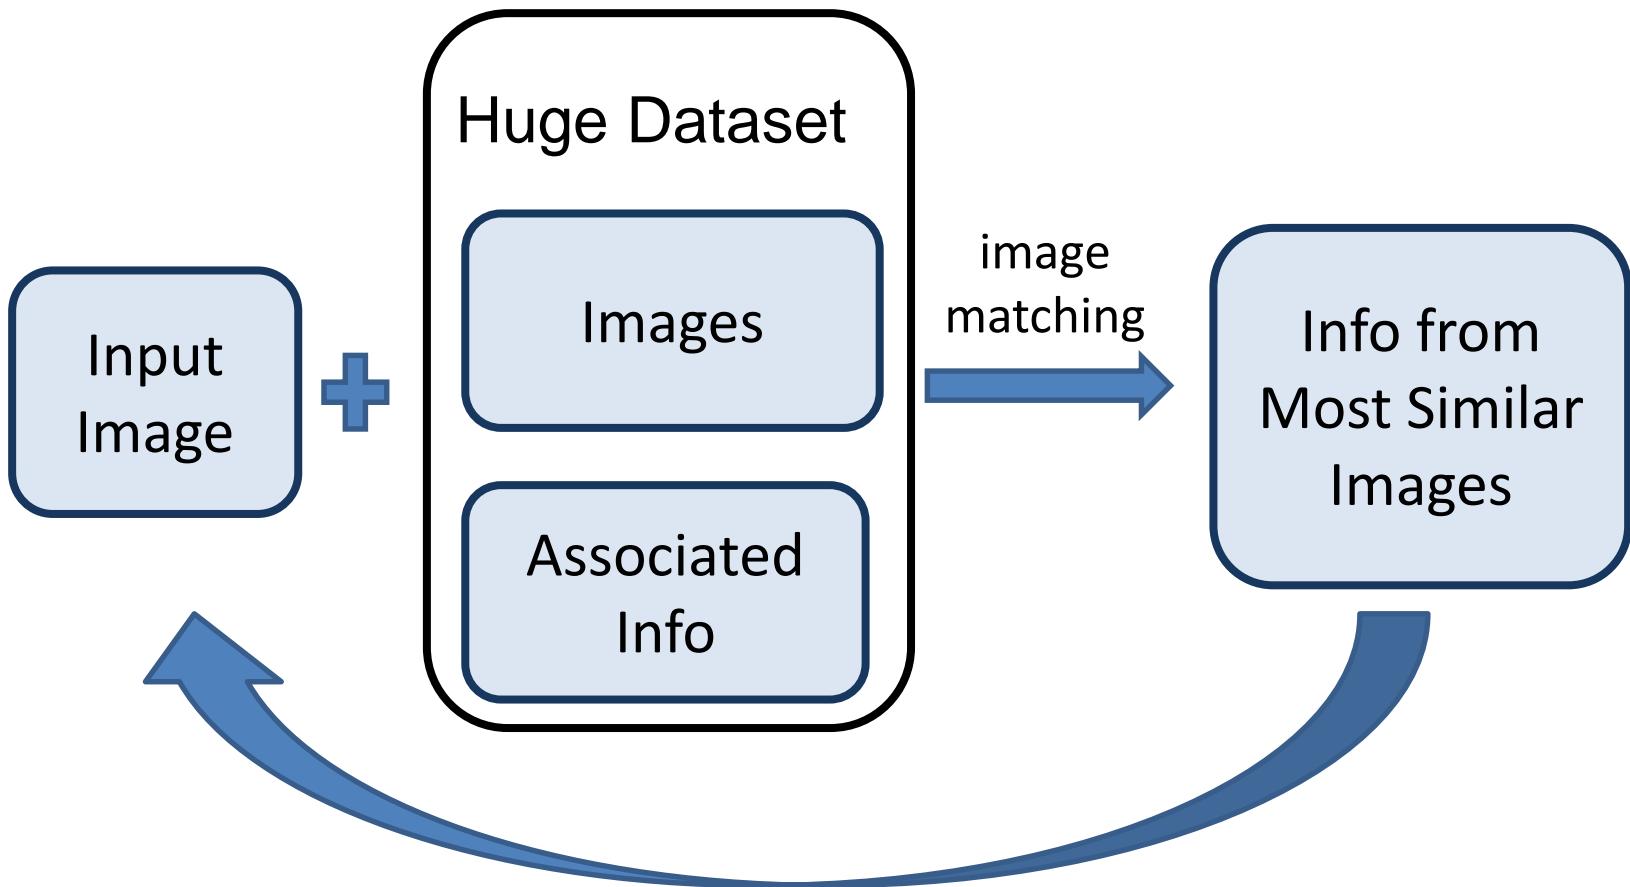


# Outline

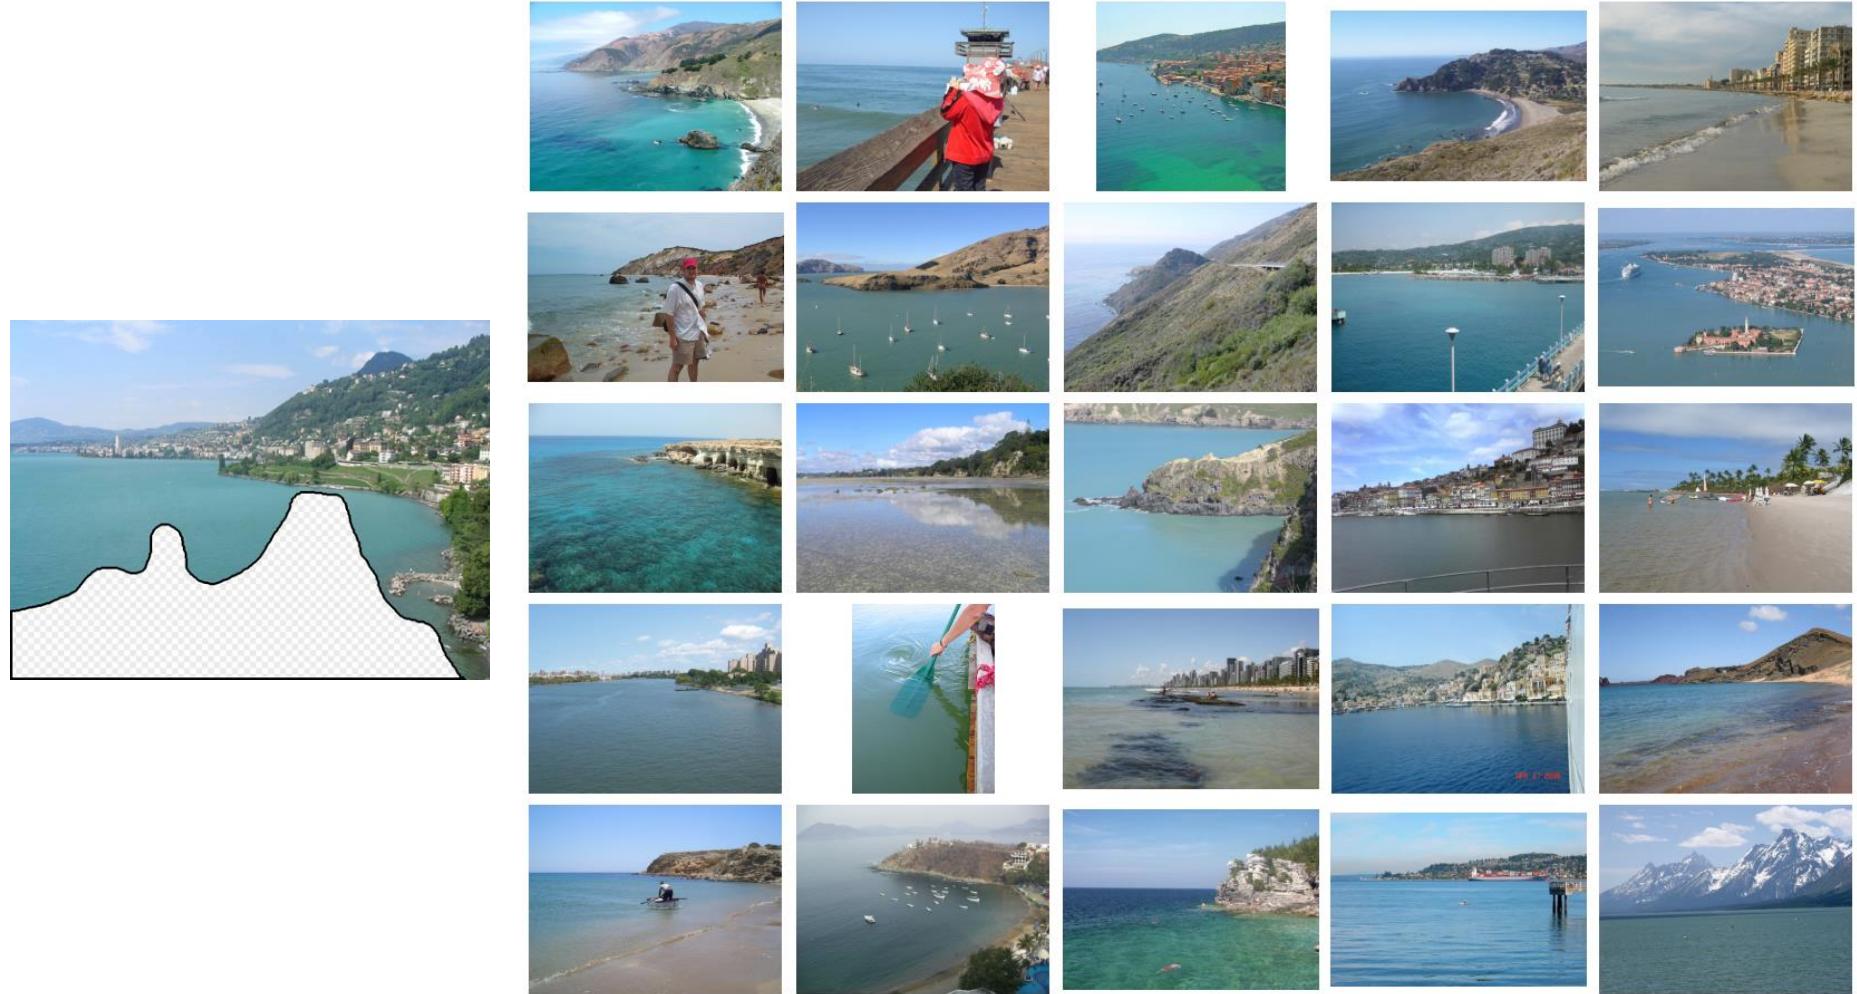
## Opportunities of Scale: Data-driven methods

- The Unreasonable Effectiveness of Data
- Scene Completion
- Im2gps
- Recognition via Tiny Images
- Project 5 Intro

# General Principal



Hopefully, If you have enough images, the dataset will contain very similar images that you can find with simple matching methods.



... 200 total



Graph cut + Poisson blending

# im2gps (Hays & Efros, CVPR 2008)

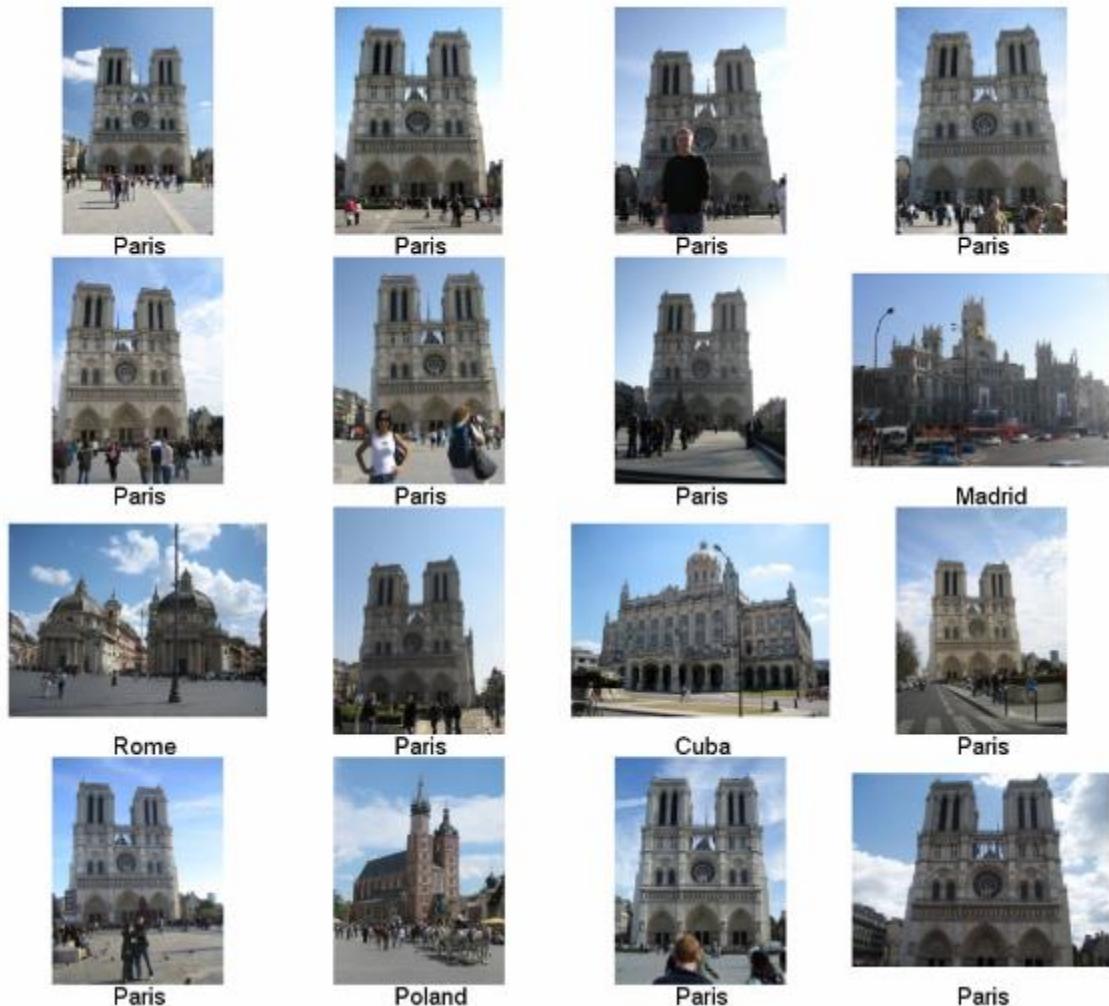


6 million geo-tagged Flickr images

<http://graphics.cs.cmu.edu/projects/im2gps/>

# How much can an image tell about its geographic location?





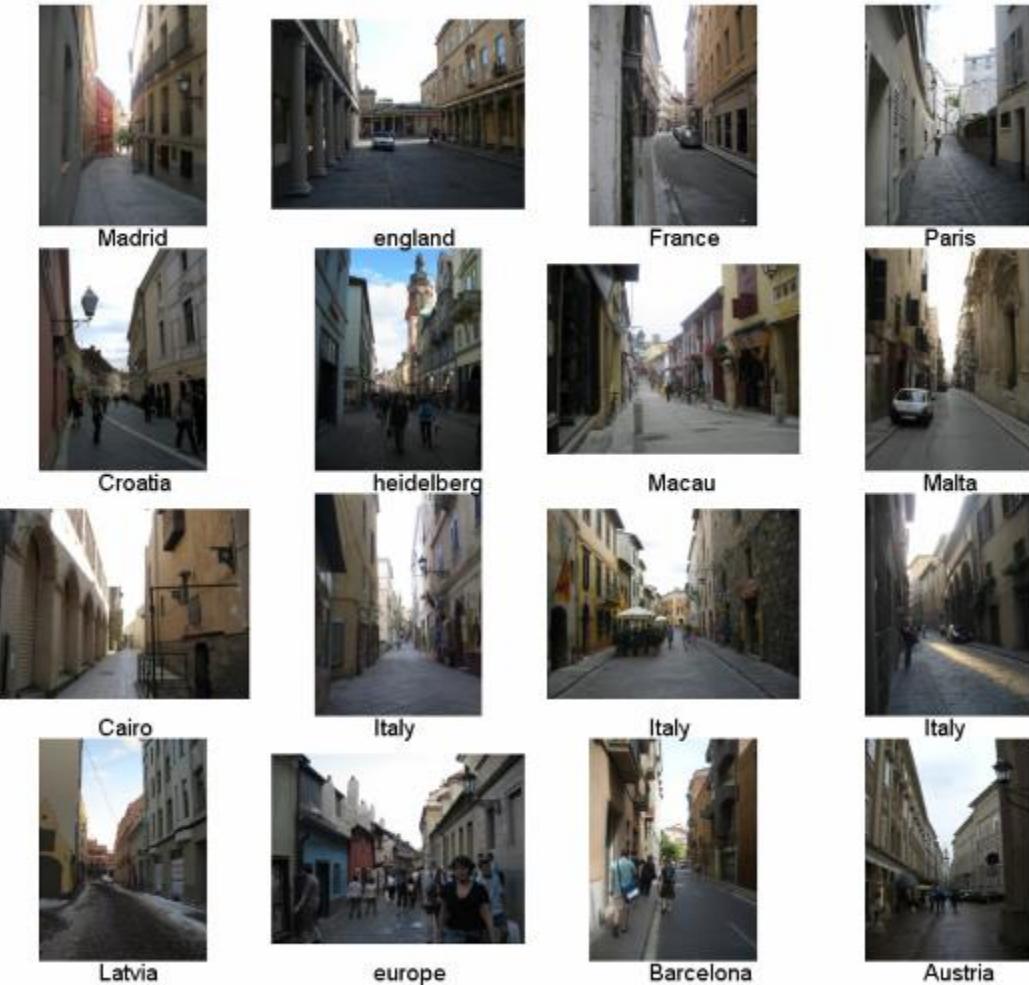
Nearest Neighbors according to gist + bag of SIFT + color histogram + a few others



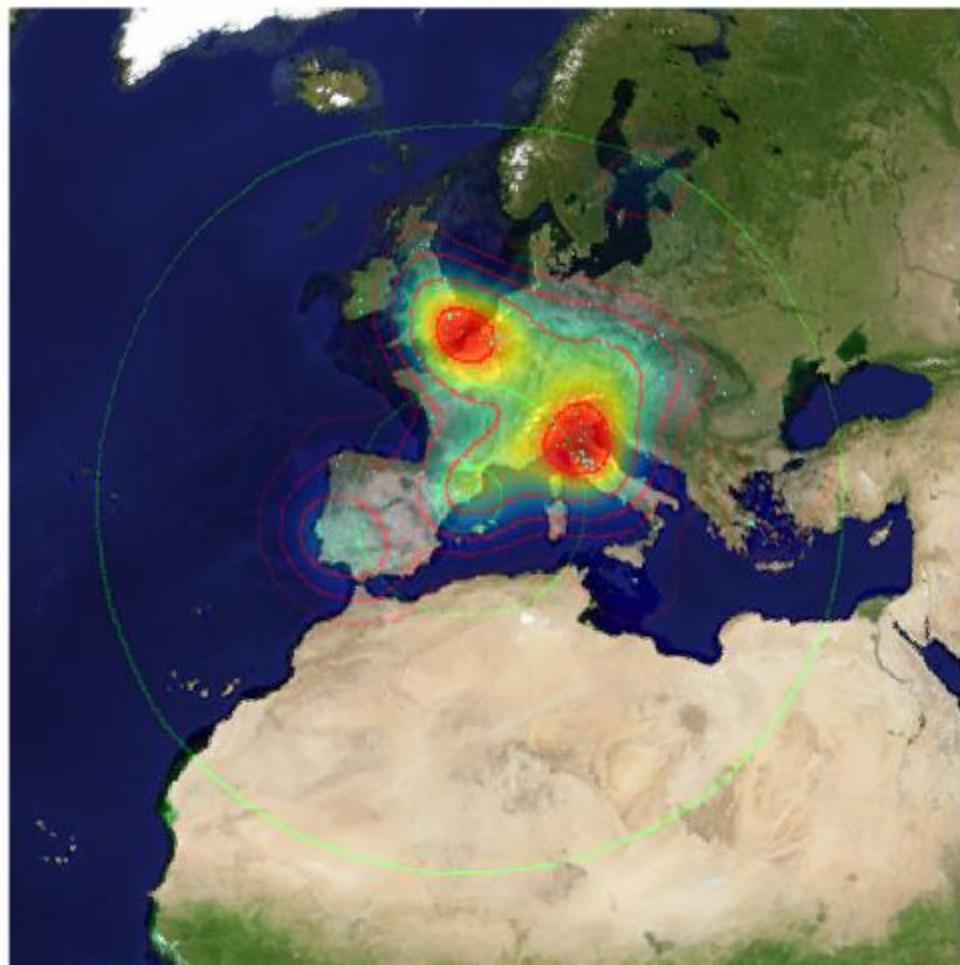
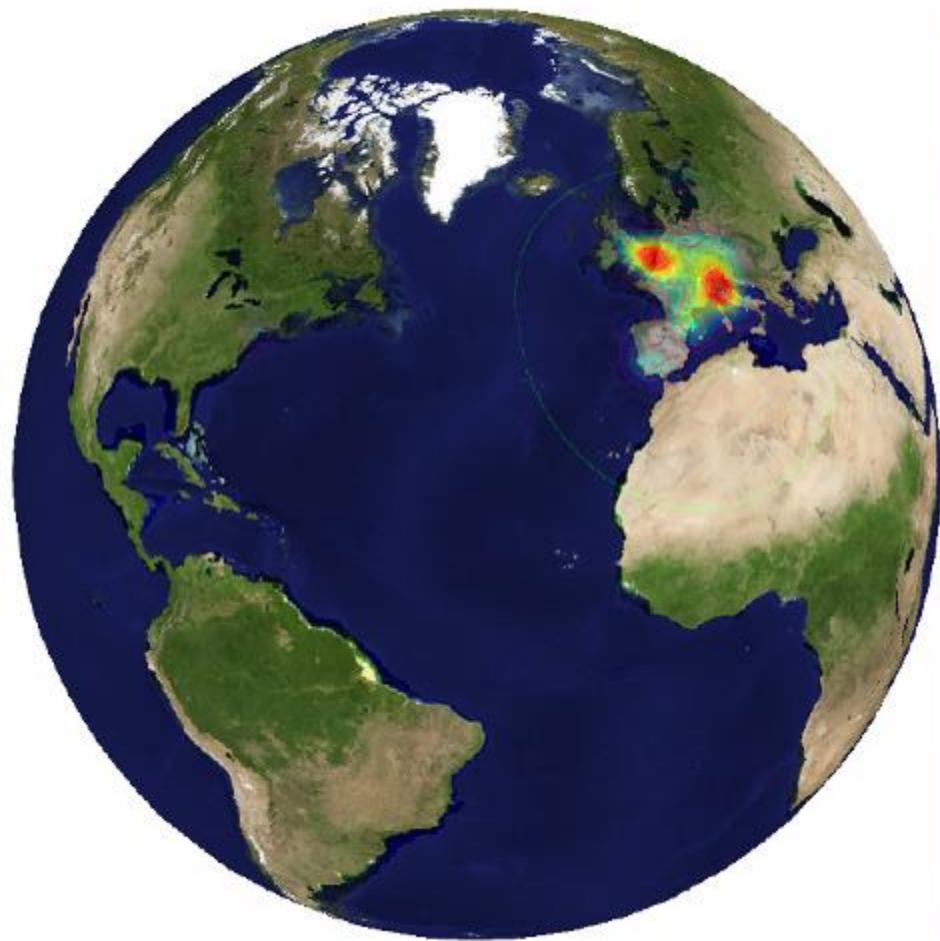
Im2gps



# Example Scene Matches

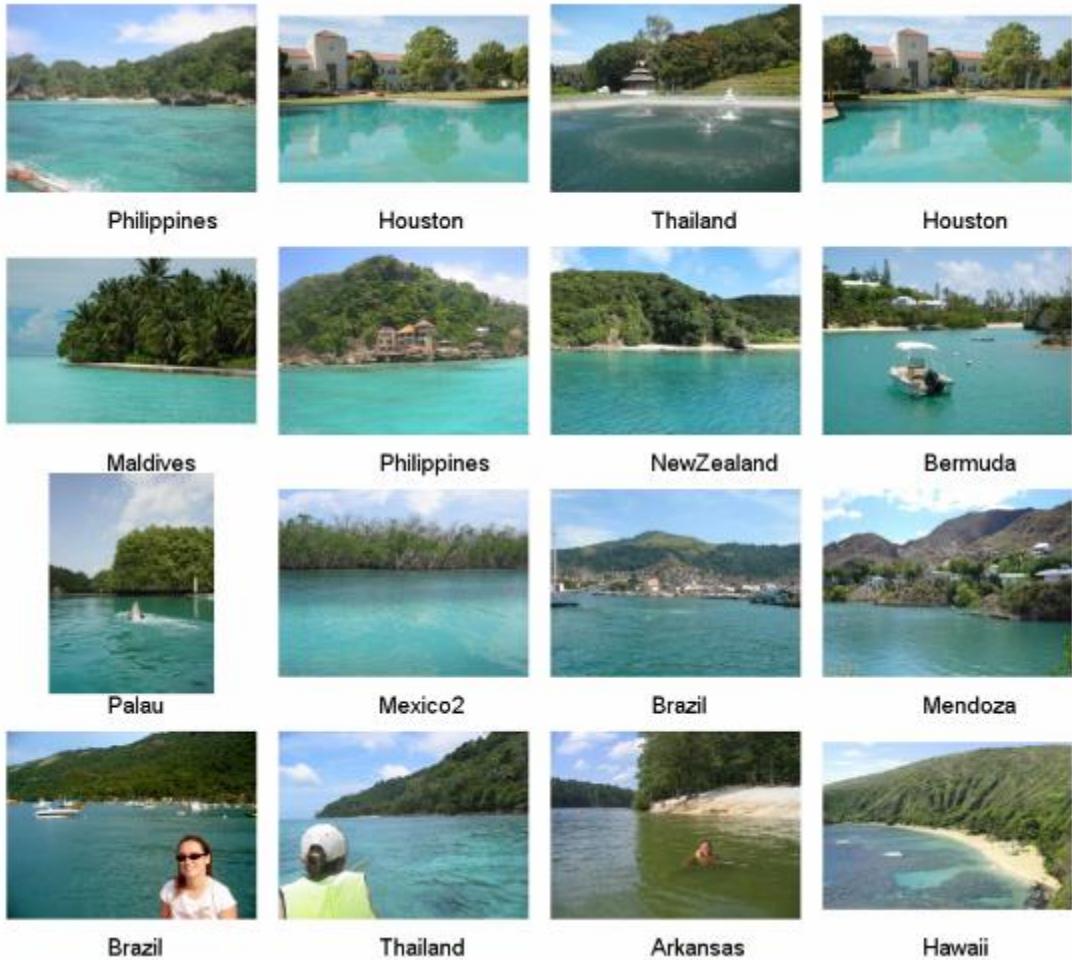


# Voting Scheme



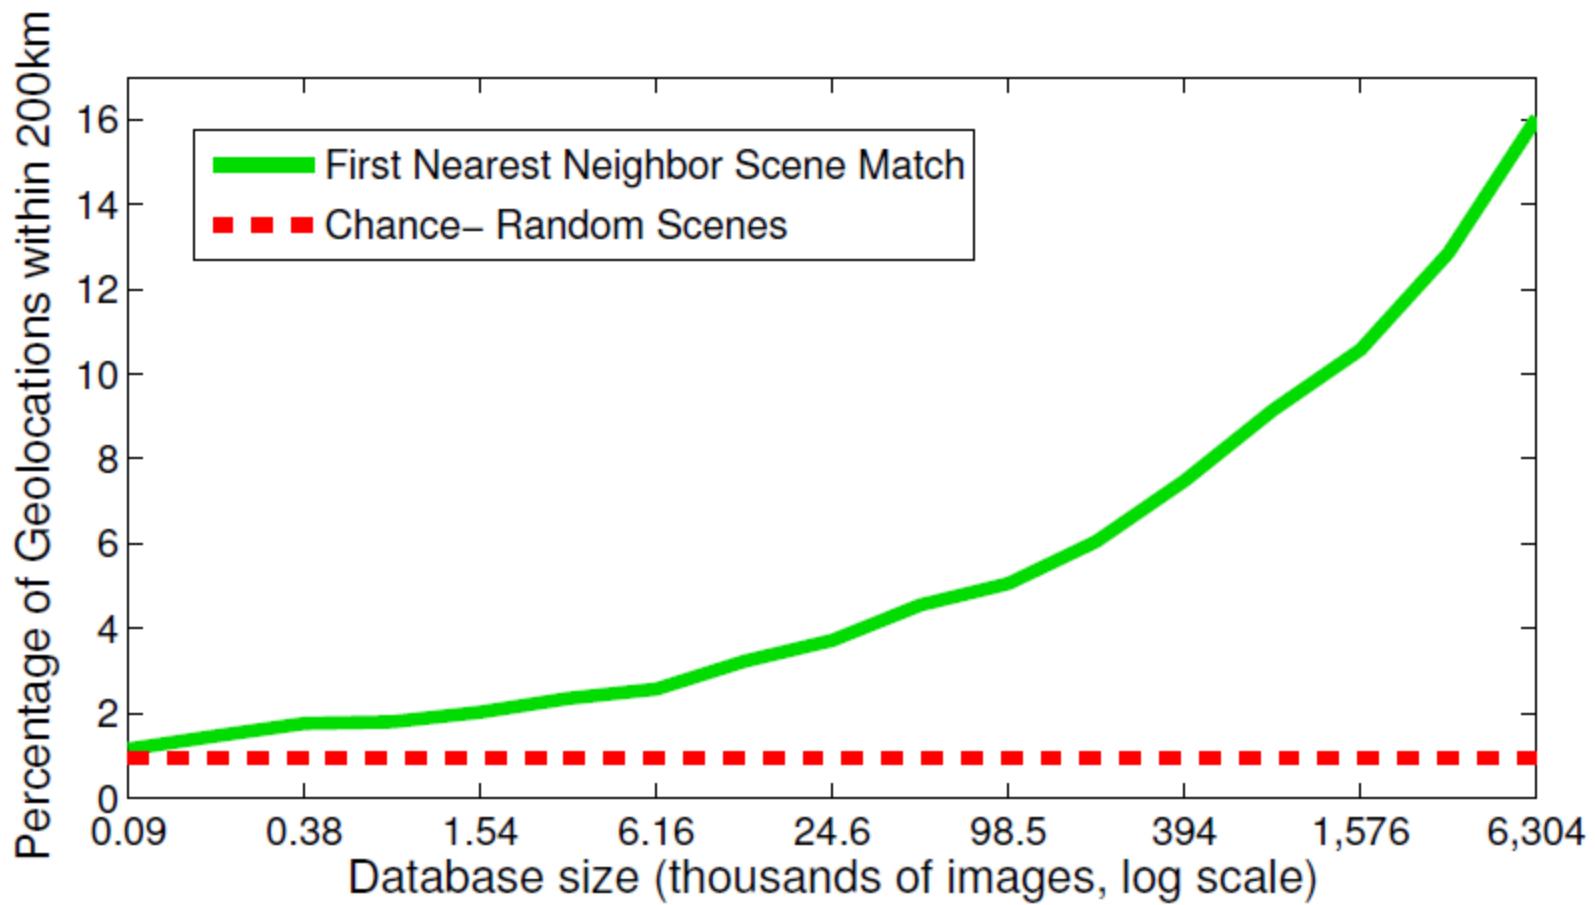
im2gps







# Effect of Dataset Size



# Population density ranking

High Predicted Density



...



Low Predicted Density

# Where is This?



[Olga Vesselova, Vangelis Kalogerakis, Aaron Hertzmann, James Hays, Alexei A. Efros. Image Sequence Geolocation. ICCV'09]

# Where is This?



# Where are These?



15:14,  
June 18<sup>th</sup>, 2006



16:31,  
June 18<sup>th</sup>, 2006

# Where are These?



15:14,  
June 18<sup>th</sup>, 2006



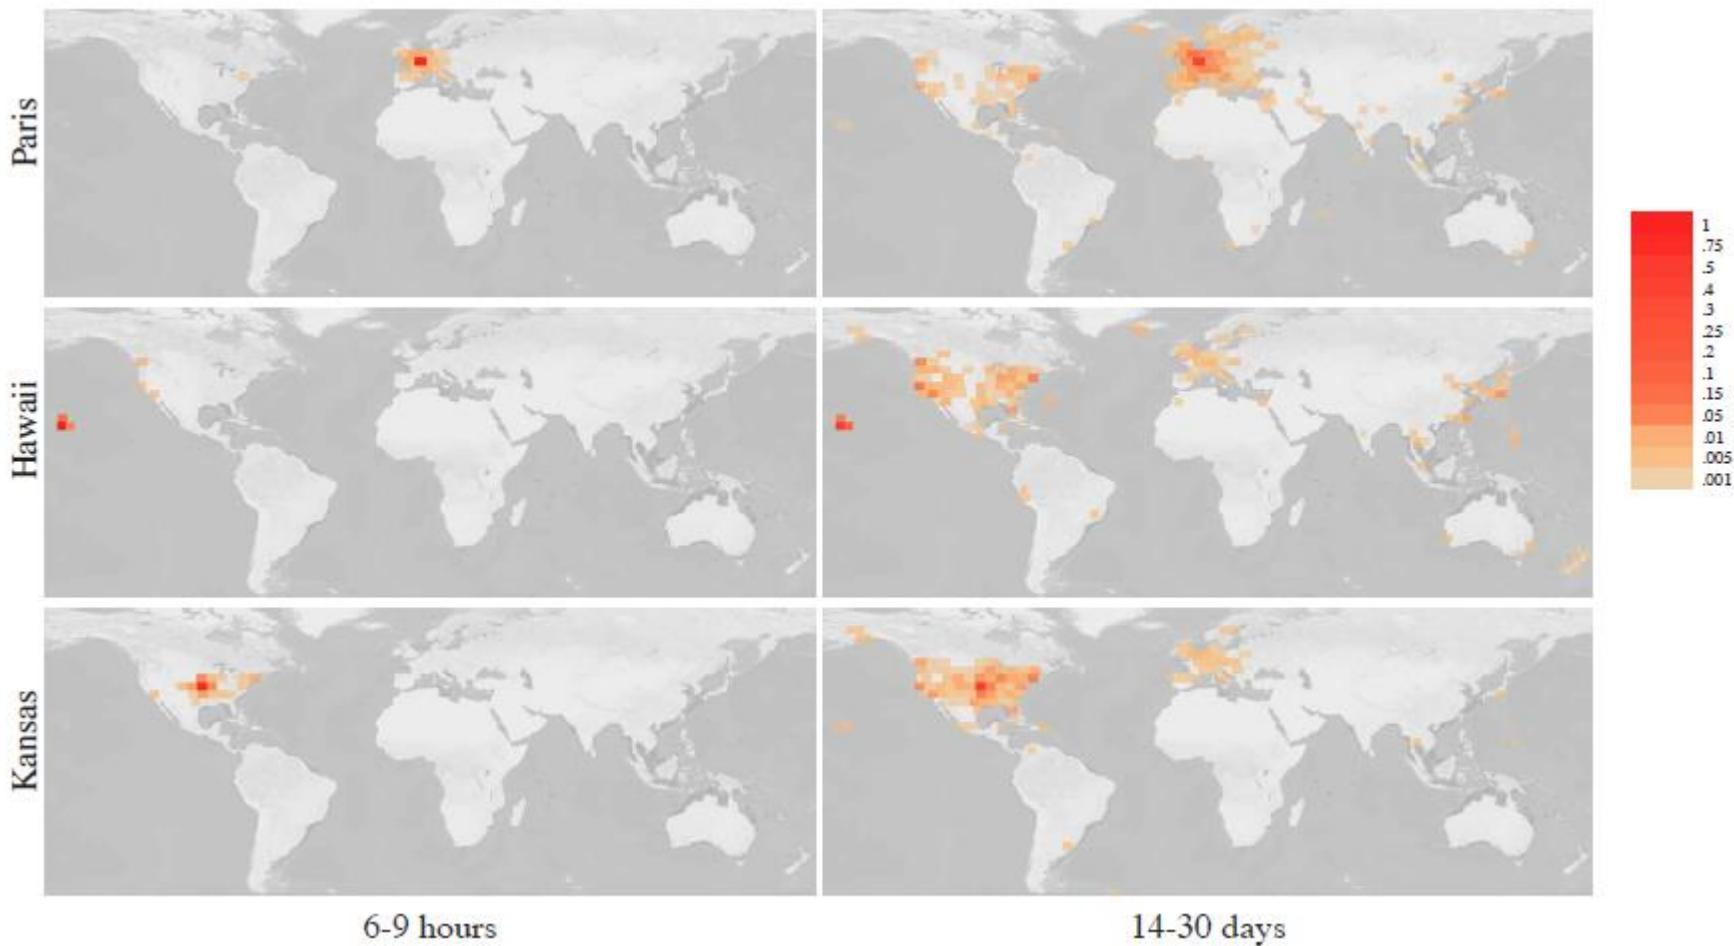
16:31,  
June 18<sup>th</sup>, 2006



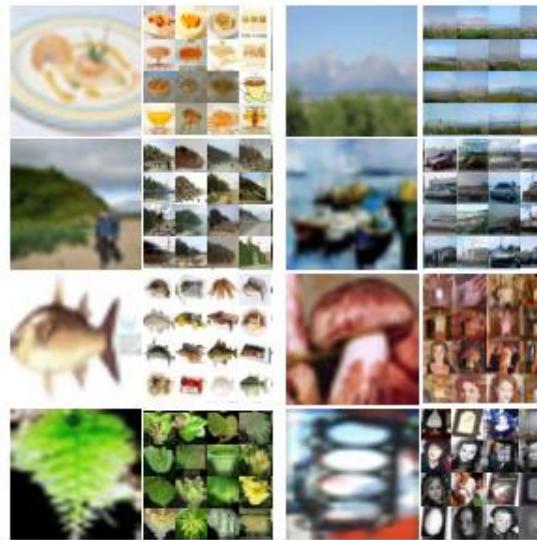
17:24,  
June 19<sup>th</sup>, 2006

# Results

- im2gps – 10% (geo-loc within 400 km)
- temporal im2gps – 56%



# Tiny Images



80 million tiny images: a large dataset for non-parametric object and scene recognition  
Antonio Torralba, Rob Fergus and William T. Freeman. PAMI 2008.

<http://groups.csail.mit.edu/vision/TinyImages/>

256x256



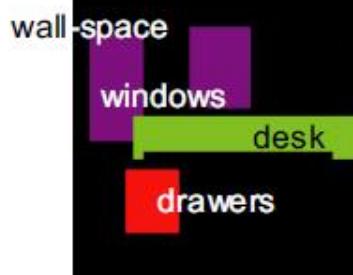
256x256



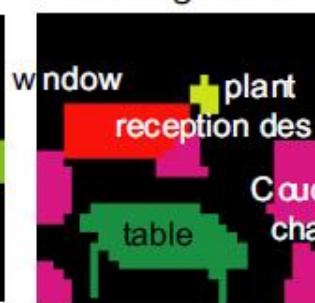
32x32



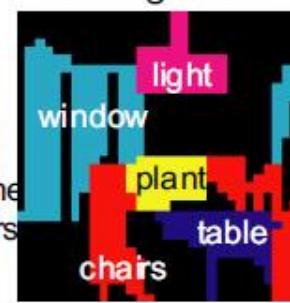
office



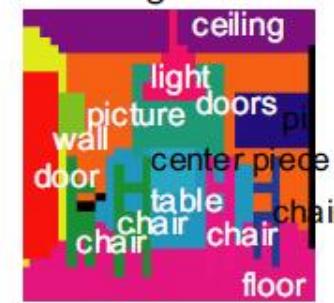
waiting area



dining room



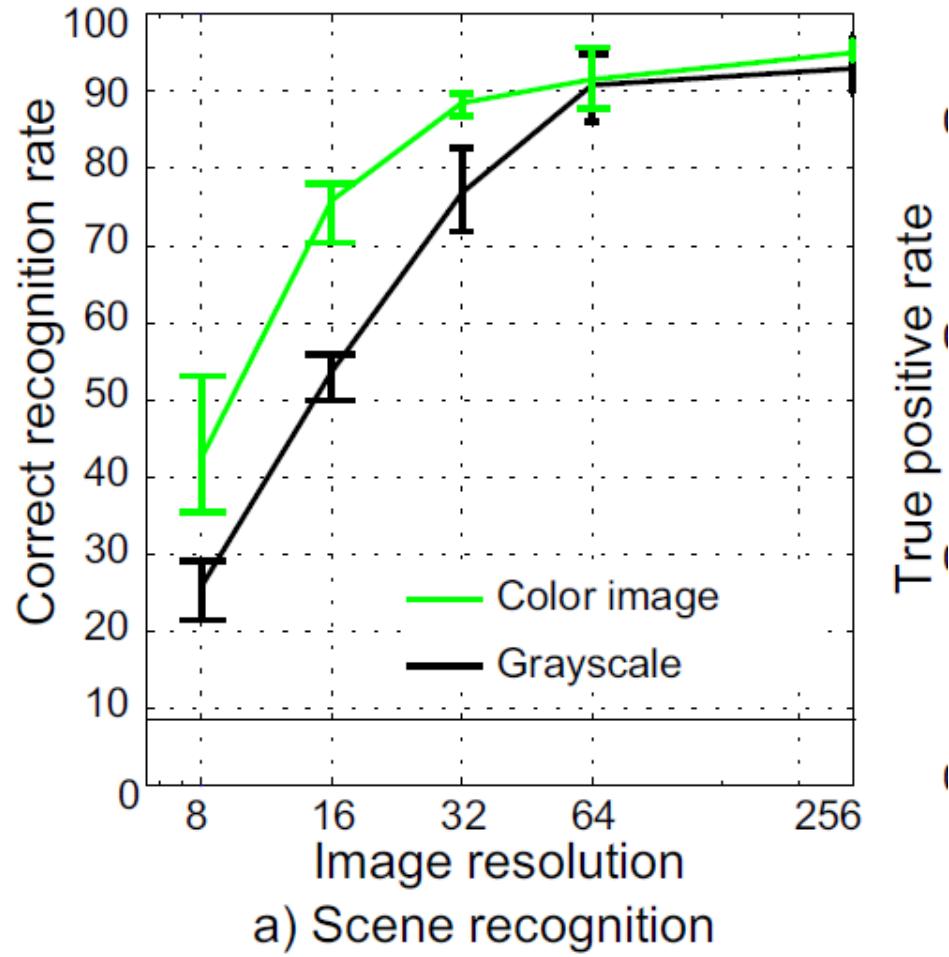
dining room



### c) Segmentation of 32x32 images



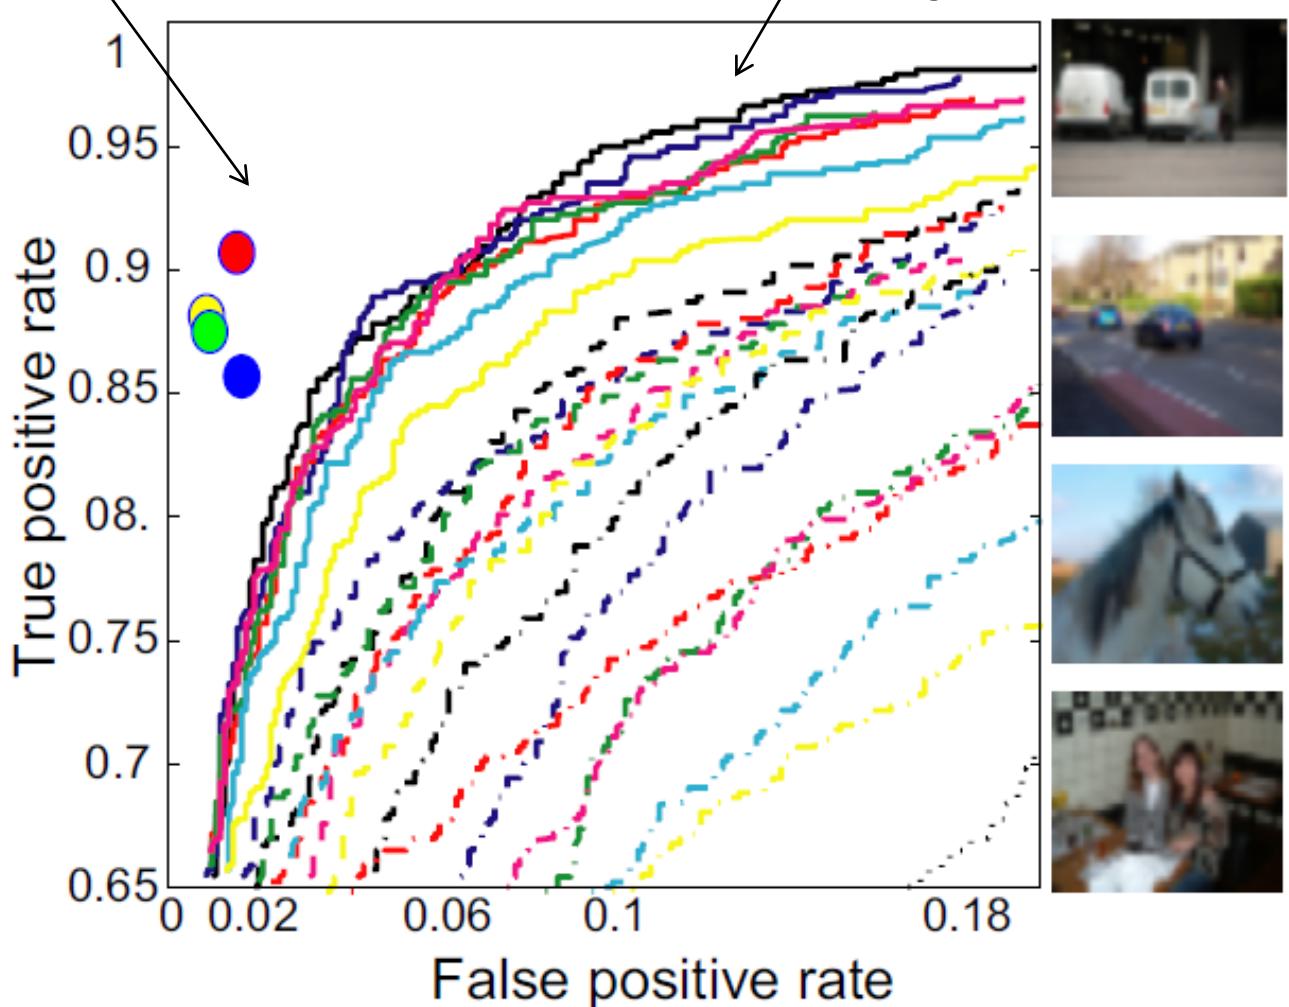
# Human Scene Recognition



# Humans vs. Computers: Car-Image Classification

Humans for 32 pixel tall images

Various computer vision  
algorithms for full resolution  
images



# Powers of 10

Number of images on my hard drive:  $10^4$



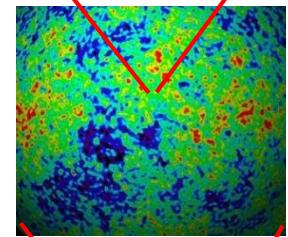
Number of images seen during my first 10 years:  $10^8$   
(3 images/second \* 60 \* 60 \* 16 \* 365 \* 10 = 630720000)



Number of images seen by all humanity:  $10^{20}$   
106,456,367,669 humans<sup>1</sup> \* 60 years \* 3 images/second \* 60 \* 60 \* 16 \* 365 =  
1 from <http://www.prb.org/Articles/2002/HowManyPeopleHaveEverLivedonEarth.aspx>



Number of photons in the universe:  $10^{88}$



Number of all 32x32 images:  $10^{7373}$   
 $256^{32 \times 32 \times 3} \sim 10^{7373}$



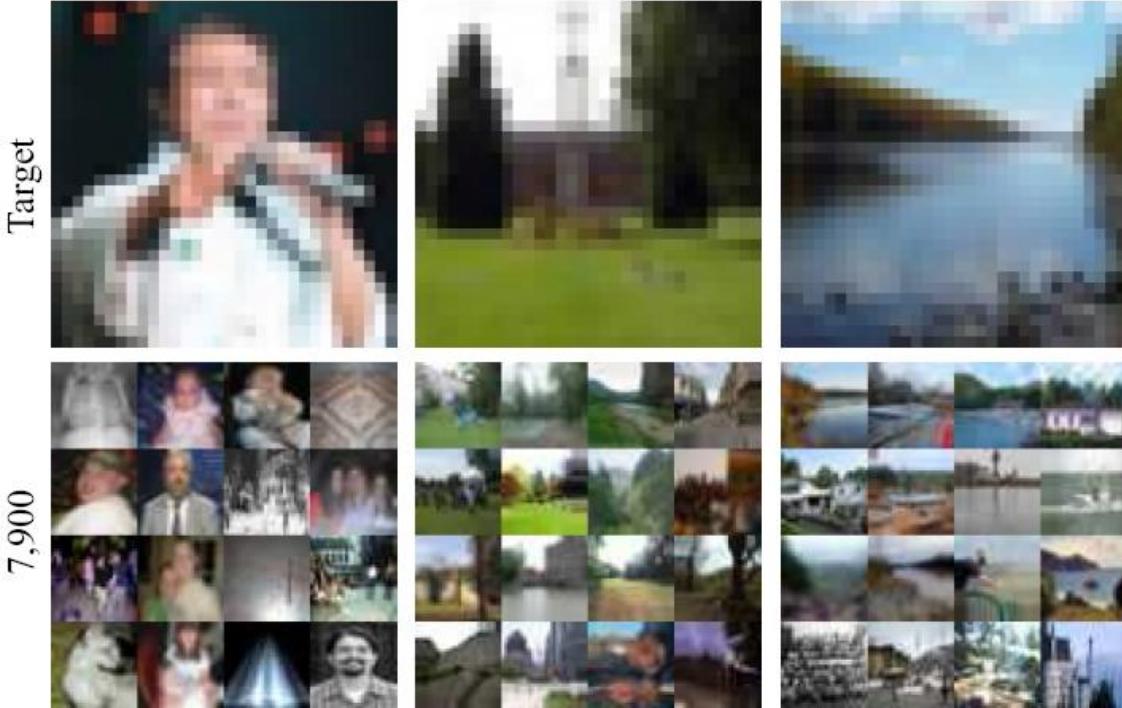
# Scenes are unique



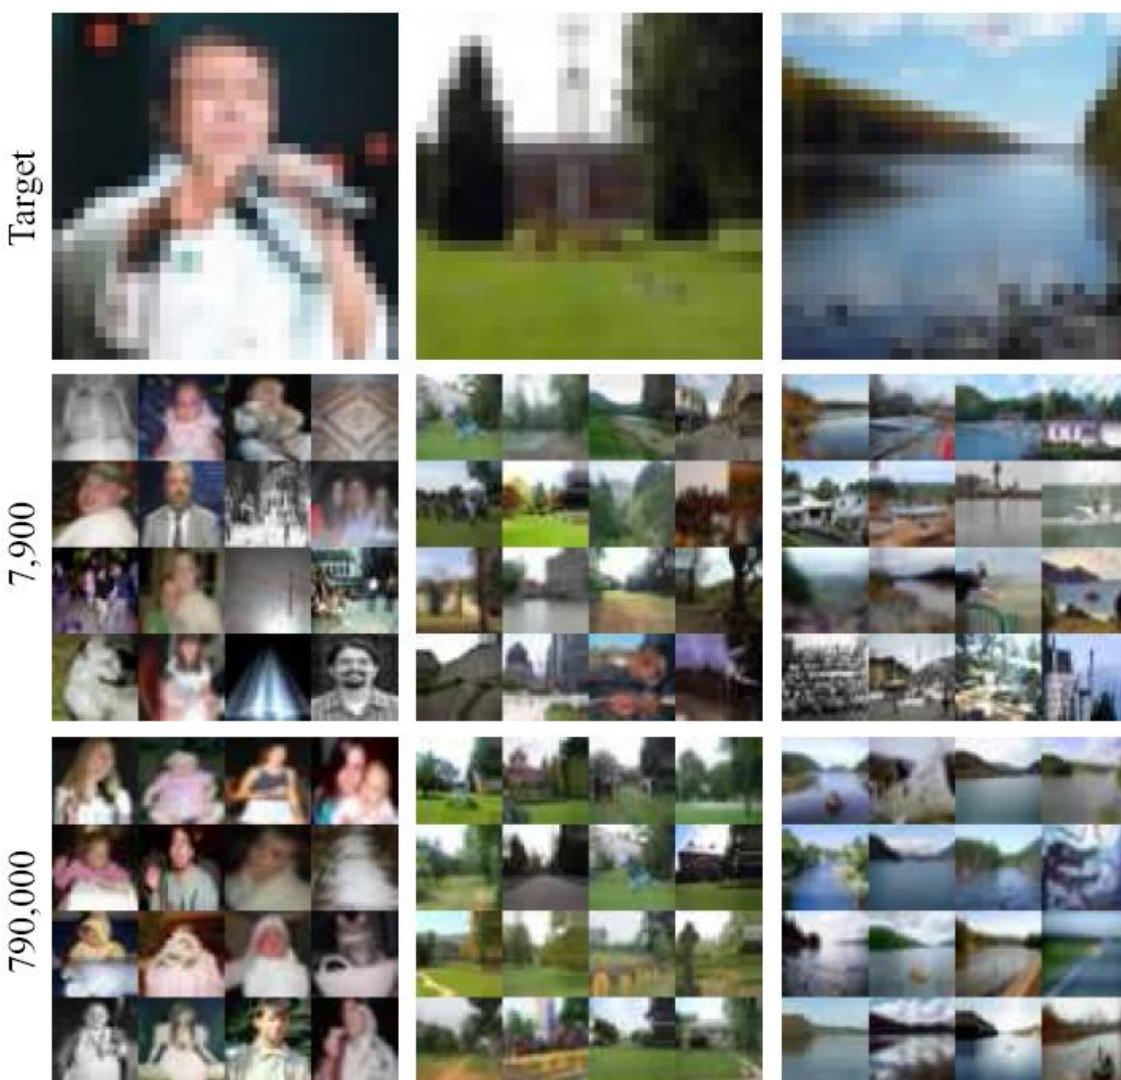
# But not all scenes are so original



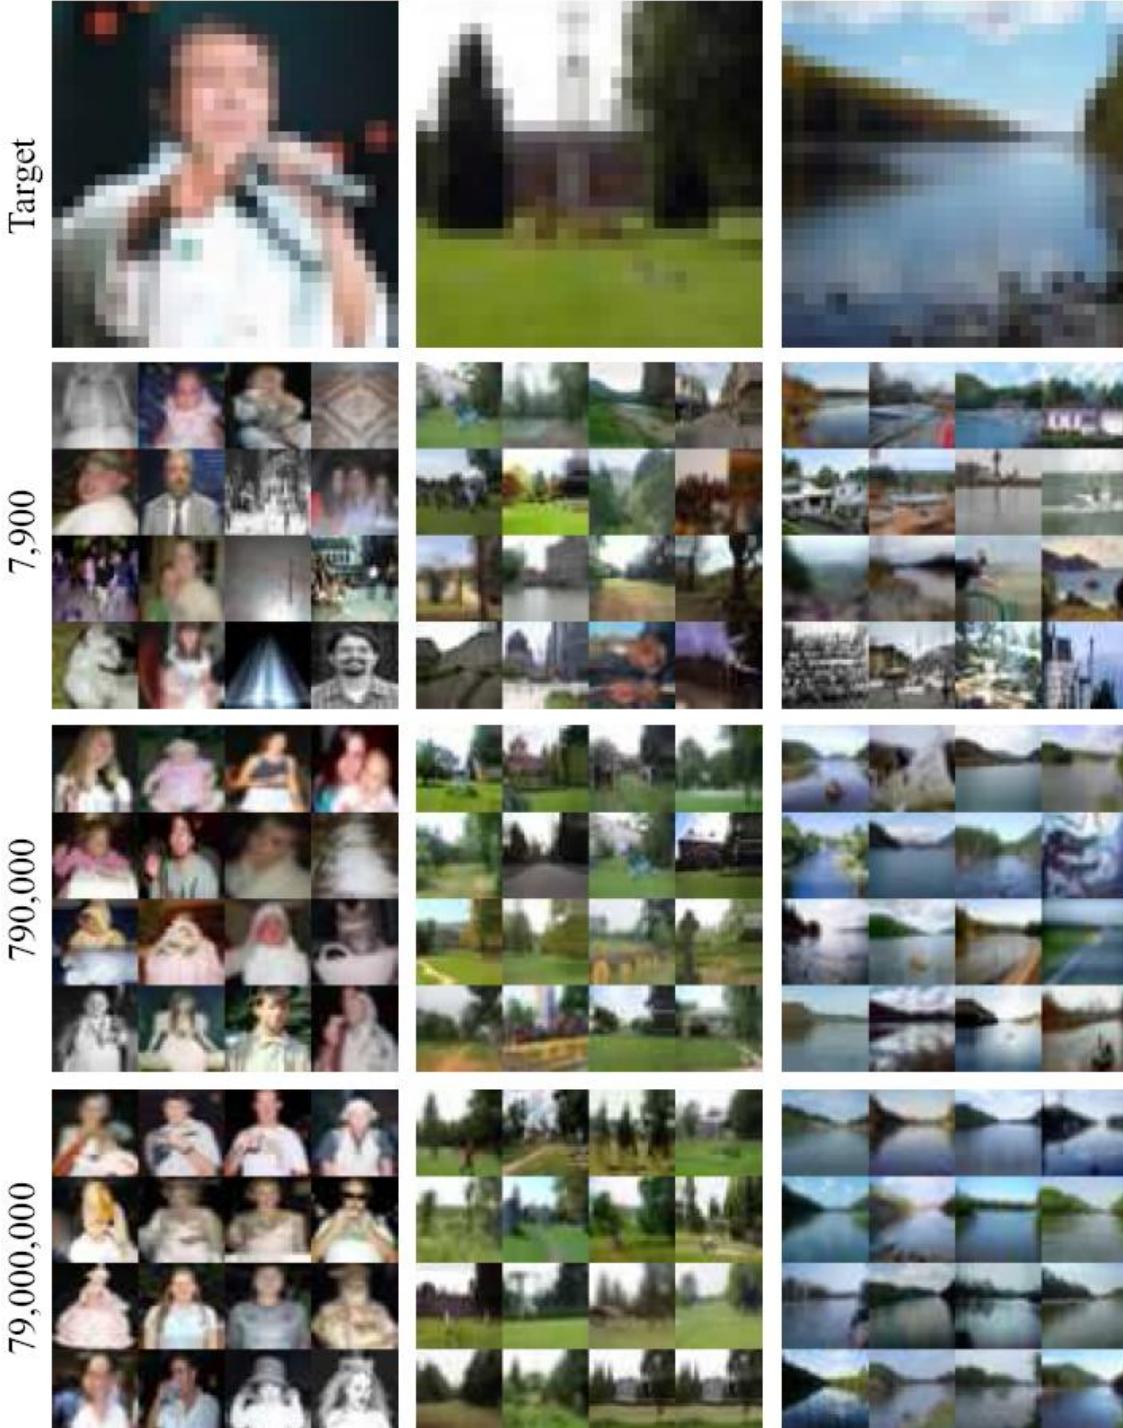
# Lots Of Images



# Lots Of Images



# Lots Of Images



# Application: Automatic Colorization



Input



Color Transfer



Color Transfer



Matches (gray)



Matches (w/ color)



Avg Color of Match

# Application: Automatic Colorization



Input



Color Transfer



Color Transfer



Matches (gray)

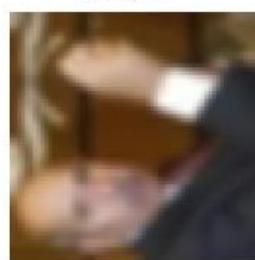
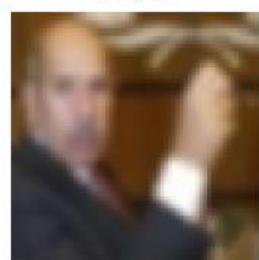


Matches (w/ color)



Avg Color of Match

# Automatic Orientation Examples



# Summary

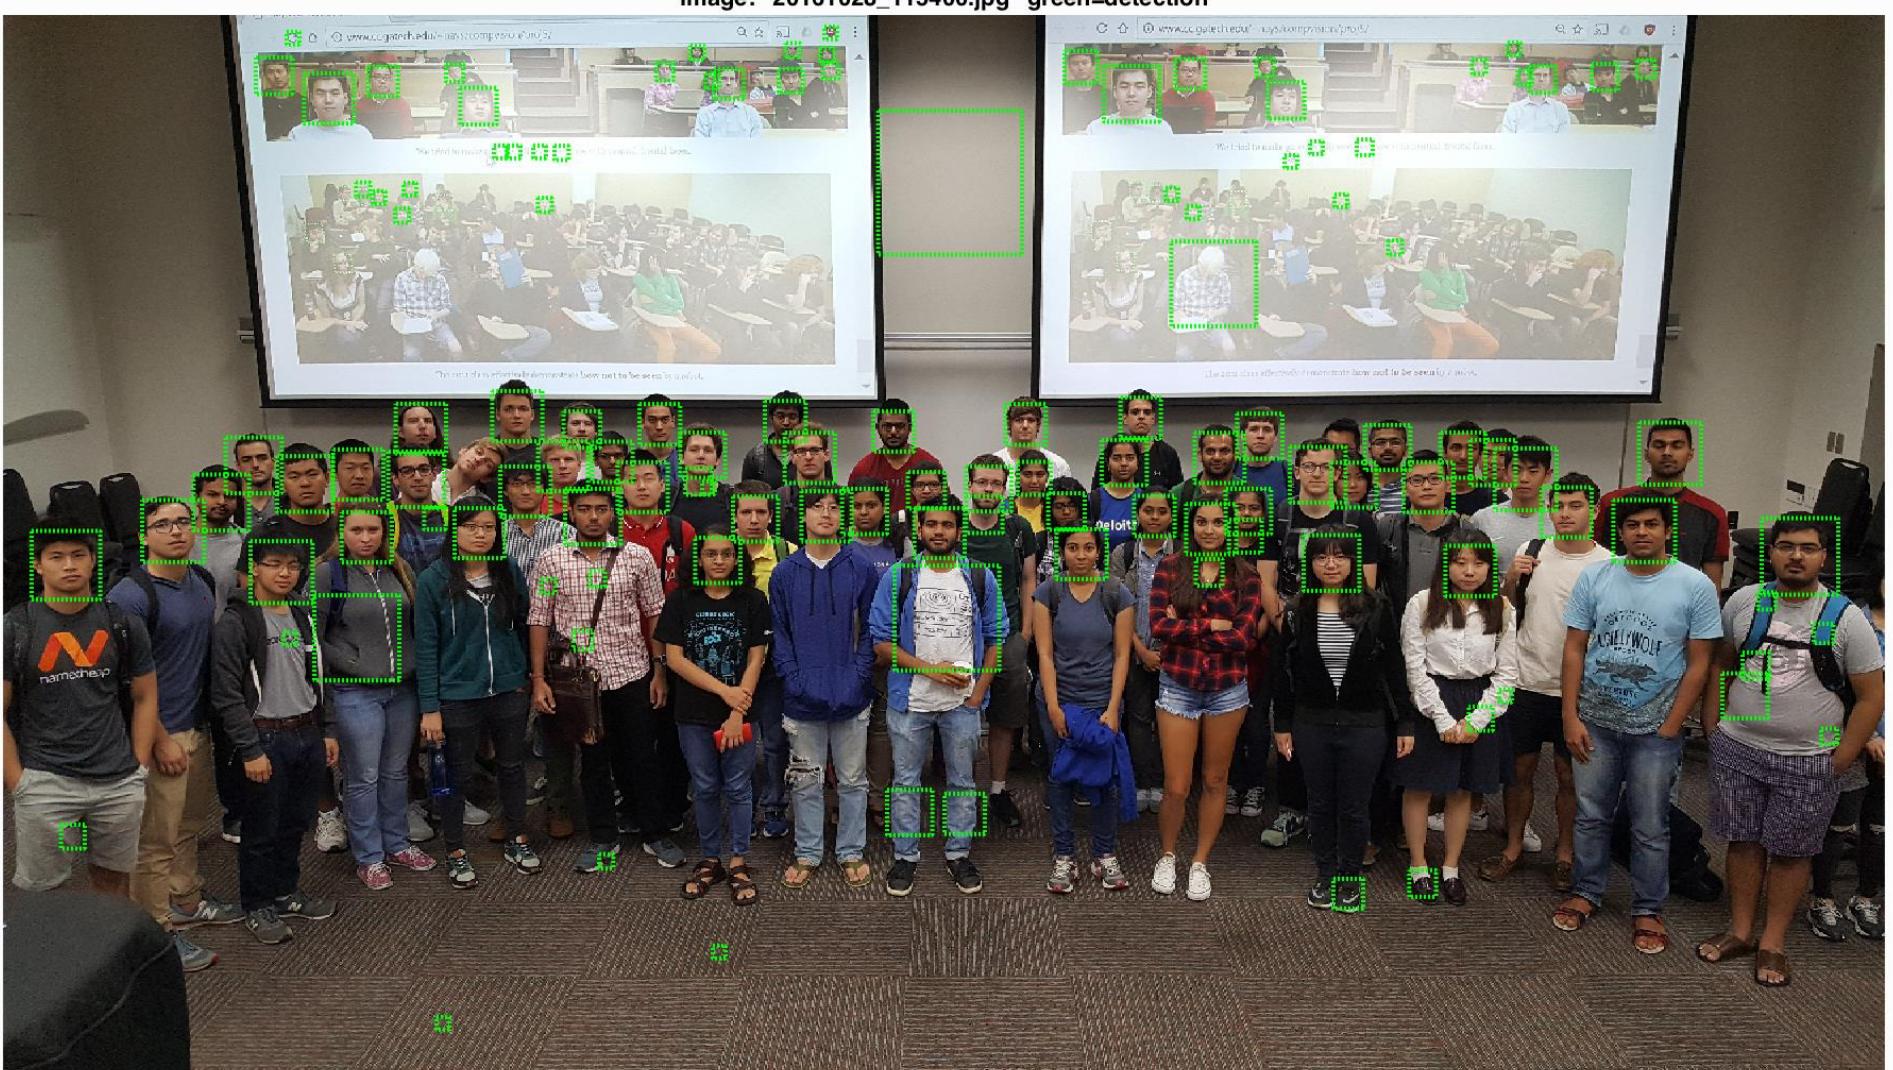
- With billions of images on the web, it's often possible to find a close nearest neighbor
- In such cases, we can shortcut hard problems by “looking up” the answer, stealing the labels from our nearest neighbor
- For example, simple (or learned) associations can be used to synthesize background regions, colorize, or recognize objects



# Project 5

- <http://www.cc.gatech.edu/~hays/compvision/proj5/>





Wenqi Xian